

Minimal residual certificates for the Erdős–Straus conjecture: a verified census to 10^{12}

J. M. Hopkins*

August 2026

Abstract

The Erdős–Straus conjecture asserts that $4/n = 1/a + 1/b + 1/c$ is solvable in positive integers for every $n \geq 2$; the open cases concentrate on primes in six residue classes mod 840 (*hard primes*). Every solution at a hard prime p has first denominator $a = (p + R)/4$ for an *admissible residual* $R \equiv 3 \pmod{4}$, and solvability at that R is the divisor-class condition $\exists k \mid m^2$, $m = pa$, with $k \equiv -m \pmod{R}$ — a criterion due to Bradford. The conjecture at p is therefore exactly the finiteness of the *minimal working residual* $R_{\min}(p)$, a structural invariant that verification records do not report.

We compute and independently verify $R_{\min}$ for all 1,175,215,396 hard primes below 10^{12} . The maximum is $R_{\min} = 111$, first attained at $p = 119,945,383,009$; it had stood at 107 across the preceding five orders of magnitude, and a heuristic model calibrated at 10^9 predicted in advance both the filling of a then-empty band $87 \leq R \leq 103$ and the decade in which the record would move. Verification is graded: exhaustive exact re-checking below 10^9 , complete deep tails and large systematic samples above, every certificate re-verified as $4abc = p(bc + ac + ab)$ in exact integer arithmetic, and all datasets published with checksums.

Around the census we assemble the apparatus it exercises. There are exact solvability criteria at $R = 3, 7, 11$ and at the first composite residual $R = 15$; those at $R = 3, 7$ are Yamamoto’s, restated in residual form, while the exact criteria at $R = 11$ and $R = 15$ are new. A support theorem — the critical-number theorem of Diderrich–Mann, transported to this setting — shows that failure at R forbids at least half the unit classes mod R on the shifted form $(p + R)/4$, uniformly in R , and its classical exceptions fall exactly at the residuals where an exact criterion is available. The uniform consequence bounds $R_{\min}$ itself: all but $O(x \exp(-c(\log \log x)^2))$ hard primes $p \leq x$ satisfy $R_{\min}(p) \leq \varepsilon \log \log x$, with explicit certificates — a threshold exponentially smaller than the one implicit in Vaughan’s covering congruences at comparable exceptional-set quality. The elementary layer and the finite enumerations behind the criteria are machine-checked in Lean 4/mathlib.

2020 Mathematics Subject Classification. Primary 11D68; Secondary 11N35, 11Y50, 68V20.

Keywords. Erdős–Straus conjecture, Egyptian fractions, unit fractions, sieve methods, formal verification.

*Independent researcher. jeff.m.hopkins@gmail.com. Computations and computer-assisted case analyses were carried out with the code repository accompanying this paper: <https://github.com/jeffmhopkins/Erd-s-Straus-attack>.

1 Introduction

Erdős and Straus conjectured (see [5]) that for every integer $n \geq 2$ the equation

$$\frac{4}{n} = \frac{1}{a} + \frac{1}{b} + \frac{1}{c} \tag{1.1}$$

has a solution in positive integers a, b, c . By multiplicativity it suffices to treat prime n , and classical polynomial identities settle every prime outside the six residue classes

$$\mathcal{H} \bmod 840 = \{1, 121, 169, 289, 361, 529\},$$

the squares of units modulo 840 (Mordell [23, Ch. 30]; the quadratic-residue nature of the obstruction goes back to Yamamoto [31]). We call primes in these classes *hard primes*. Direct verification has confirmed the conjecture far beyond any range considered here: to 10^{14} by Swett [26], to 10^{17} by Salez [24], and to 10^{18} by Mihnea and Bogdan [22]; the present paper is logically independent of those bounds. The structure of the solution count was analyzed by Elsholtz and Tao [4]. For calibration on the density axis: the first density bounds on the exceptional set are due to Webb [30], sieve bounds were advanced by Sander [25], and Vaughan’s classical mean-value bound [29] gives an exceptional set $\ll x \exp(-c(\log x)^{2/3})$ for (1.1) — asymptotically stronger than any fixed power of $\log x$, and still the record on that axis.

This paper is about a different quantity, and it is primarily a computation. A verification to 10^{18} answers “is (1.1) solvable at n ?” and records a yes; it does not record *how hard* the instance was, and the published records do not report the parameter in which the difficulty is measured. By Proposition 1.1 — Bradford’s criterion [9], restated — every solution at a hard prime p is a certificate at some admissible residual $R = 4a - p \equiv 3 \pmod{4}$, and the conjecture at p is exactly the finiteness of the least such R , written $R_{\min}(p)$. Our main object is a census of $R_{\min}$: its value for every one of the 1,175,215,396 hard primes below 10^{12} , verified, archived, and reproducible. What comes out of it is a distribution with sharp structure (Section 4), a broken record (the maximum moves from 107 to 111 in the last decade computed), and a heuristic model calibrated at 10^9 whose advance predictions about the tail were borne out at 10^{11} and 10^{12} .

The rest of the paper is the apparatus the census exercises and the structure it exhibits: exact per-residual solvability criteria (Section 3), which are what make the head of the distribution computable in closed form; a support theorem valid at every residual, which converts failure into a sieve condition; and, as its capstone, an almost-all bound on $R_{\min}$ itself (Theorem 5.5). Section 5 states these and sketches the sieve consequences; complete proofs of the elementary layer, together with the finite enumerations behind the criteria, are formalized in Lean 4 and inventoried in Appendix A.

New contributions

The census.

- $R_{\min}(p)$ computed and independently verified for every one of the 1,175,215,396 hard primes $p < 10^{12}$ (Section 4), with the full distribution (Table 3), a graded verification protocol (Appendix A), and archived datasets carrying checksums.
- The maximum of $R_{\min}$ moves for the first time in five orders of magnitude: it is 107 throughout $10^7 \leq x \leq 10^{11}$ and 111 at 10^{12} , first attained at $p = 119,945,383,009$. Section 4

exhibits the certificate and the exhaustive refutation of every admissible $R \leq 107$ at that prime.

- A calibrated independence model, fitted to the complete 10^9 solvability masks, that predicted in advance both the population of the band $87 \leq R_{\min} \leq 103$ — empty when the model was fitted — and the decade in which the record would move. Both predictions were tested by the 10^{11} and 10^{12} runs and held (Section 4).
- A reproducibility layer: exact re-verification of every certificate as $4abc = p(bc + ac + ab)$, a Lean 4 [3] / mathlib [21] formalization of the elementary layer and of the finite enumerations behind the exact criteria, and a published run kit. Appendix A gives the inventory and the exact trust base.

The apparatus, and what in it is new.

- *Exact* solvability criteria at $R = 11$ (Theorem 1.5) and at the first composite residual $R = 15$ (Theorem 1.6: a Jacobi-character dichotomy with no budget cases), together with the finite-reduction lemma (Theorem 1.3) that licenses their enumeration. The criteria at $R = 3$ and $R = 7$ (Theorems 1.2, 1.4) are *not* new — they are Yamamoto’s, in a different notation; see Remark 1.7, which also delimits what is new at $R = 11$ and $R = 15$.
- A reciprocity identity (Theorem 1.9) collapsing the per-residual conditions into one global statement about p , which we use to describe the record prime mechanically.
- The support theorem (Theorem 1.11): failure at R confines the factor classes of $(p + R)/4$ to at most half the unit classes mod R , for every admissible R at once. As an additive statement this is the critical-number theorem of Diderrich–Mann [13]; the contribution here is the transport to residuals (bounded exponent budgets, a single prescribed target) and the observation that its classical exceptional groups occur exactly at $R = 3, 7, 15$, where Theorems 1.2, 1.4 and 1.6 supply an exact criterion — so the sieve dimension $\frac{1}{2}$ is available at every admissible residual with no gaps.
- The sieve consequences (Section 5): certificate-producing exceptional-set bounds for every finite admissible list (Theorem 1.12, Proposition 1.13, Theorem 1.14) and, as the capstone, the *uniform chain* (Theorem 5.5): all but $O(xe^{-c(\log \log x)^2})$ hard primes $p \leq x$ have $R_{\min}(p) \leq \varepsilon \log \log x$, with an explicit certificate for each. Remark 1.15 states precisely what this does and does not improve on.
- An open problem carved out of the ceiling of that argument: Conjecture 5.10, an aperiodic branch-classification statement, verified exhaustively in the accessible range, whose resolution would raise the bound to $R_{\min}(p) \leq (\log x)^{c_0}$ (Theorem 5.12).

Read as a whole, the paper has three layers. (1) An exact reformulation of solvability — every solution of (1.1) at a hard prime is a residual certificate (Proposition 1.1), which is Bradford’s criterion. (2) Finite, machine-checkable inputs at individual residuals: the exact criteria (Theorems 1.2–1.6), the reciprocity identity (Theorem 1.9), and the support bounds (Lemma 1.10, Theorem 1.11). (3) The census of Section 4, which is what layers (1)–(2) are for: they make $R_{\min}$ computable at scale, and the exact criteria determine the head of its distribution in closed form. The sieve consequences of Section 5 run in the other direction, predicting the shape of the tail; the computation calibrates them and tests what they predict.

1.1 The residual method

Fix a hard prime p . All six classes satisfy

$$(H1) \ p \equiv 1 \pmod{8}, \quad (H2) \ p \equiv 1 \pmod{3}, \quad (H3) \ p \bmod 7 \in \{1, 2, 4\}.$$

For an *admissible residual* $R \equiv 3 \pmod{4}$ put

$$a = \frac{p+R}{4} \in \mathbb{Z}, \quad m = pa.$$

Substituting a into (1.1) and clearing denominators shows that a solution with first denominator a exists if and only if

$$\exists k \mid m^2 \quad \text{with} \quad k \equiv -m \pmod{R}, \quad (1.2)$$

in which case $b = (k+m)/R$ and $c = (m^2/k+m)/R$ are positive integers solving (1.1). We write $R_{\min}(p)$ for the least admissible R for which (1.2) holds. The formulation loses no solutions:

Proposition 1.1 (completeness). *Every solution of (1.1) at a hard prime p , ordered $a \leq b \leq c$, has $p/4 < a \leq 3p/4$; hence $R = 4a - p$ is an admissible residual in $(0, 2p]$ and the solution is a residual certificate at R . Consequently $R_{\min}(p) \leq 2p$ if and only if the conjecture holds at p .*

Proof. With $a \leq b \leq c$, the equation gives $1/a < 4/p \leq 3/a$, so $p/4 < a \leq 3p/4$ and $R = 4a - p \in (0, 2p]$; $R \equiv 3 \pmod{4}$ is automatic from $p \equiv 1 \pmod{4}$. By the displayed equivalence, the solution yields a divisor k as in (1.2), so $R_{\min}(p) \leq R \leq 2p$. Conversely (1.2) at any admissible R produces a solution. $\square$

A word on provenance. Proposition 1.1 is not new. The bounded search space together with the divisor-congruence criterion modulo $4a - p$ is due to Bradford [9], who proves that a solution exists at p if and only if some a in $[p/4] \leq a \leq [p/2]$ admits a divisor of a^2 in a prescribed class modulo $4a - p$ — his range is in fact slightly sharper than the $a \leq 3p/4$ recorded above, and the correspondence with solutions is likewise bijective. These are by now called the Bradford conditions [22]. Criteria of this general shape go back further. Yamamoto [31] constructs a complete system Σ of solvable coverings and proves that $4/p$ is solvable if and only if p lies in their union; his simplified subsystem Σ_0 is defined on precisely the six hard classes, and the individual coverings $\{-s/q\} = \{n : q \mid n + s\}$ are, for odd q , exactly the condition “ q divides $(p + s)/4$ ”. So Yamamoto’s system is a residual-indexed criterion in different notation, and several of the per-residual statements below are his; Remark 1.7 makes the dictionary explicit and says which are. Lopez [15] states the same shift-indexed divisor condition in general form for one solution shape. Mordell’s Chapter 30 [23, Ch. 30] collects the classical identities and his own non-existence theorem; the systematic divisor-counting form is Elsholtz–Tao’s [4]; and Salez [24] proves a completeness theorem for seven modular equations at the level of degree-one prime polynomials.

What this paper does with the criterion is to *measure* it. The organizing move is to index by the residual R rather than by the denominator a , which makes $R_{\min}(p)$ a well-defined statistic of p and the object of a census: it is computable at scale because each fixed R admits a finite exact criterion (Theorem 1.3), determined in closed form at $R = 3, 7, 11, 15$, and because failure at every other R is constrained uniformly (Theorem 1.11). Bradford’s range gives $R_{\min}(p) \ll p$; the computation of Section 4 shows $R_{\min} \leq 111$ throughout $p < 10^{12}$, and Theorem 5.5 bounds it by $\varepsilon \log \log x$ for almost every hard prime.

Two elementary congruences drive everything below. From $4a = p + R$,

$$a \equiv 4^{-1}p, \quad m \equiv 4^{-1}p^2 \pmod{R}, \quad (1.3)$$

and since $\gcd(p, R) = 1$ (as p is prime and $R \neq p$, $R \equiv 3 \not\equiv p \pmod{4}$), also $\gcd(a, R) = \gcd(m, R) = 1$: the setting is always a unit computation modulo R . By (1.3), m is a *square* modulo every prime $r \mid R$; hence for every prime $r \mid R$ with $r \equiv 3 \pmod{4}$ (at least one exists, since $R \equiv 3 \pmod{4}$),

$$\left(\frac{-m}{r}\right) = \left(\frac{-1}{r}\right) = -1 \quad (1.4)$$

— the target class $-m$ is a quadratic non-residue mod r .

A worked example. Take the hard prime $p = 1009 \pmod{840}$ and the smallest admissible residual $R = 3$. Then $a = (1009 + 3)/4 = 253 = 11 \cdot 23$ and $m = pa = 255,277$. Since $11 \equiv 2 \pmod{3}$, Theorem 1.2 below promises a certificate: take $k = 11$, a divisor of m^2 with $k \equiv -m \equiv 2 \pmod{3}$. Then $b = (k + m)/3 = 85,096$ and $c = (m^2/k + m)/3 = 1,974,822,872$, and indeed

$$\frac{4}{1009} = \frac{1}{253} + \frac{1}{85,096} + \frac{1}{1,974,822,872}.$$

Had every prime factor of a been $\equiv 1 \pmod{3}$, every divisor of m^2 would be $\equiv 1 \pmod{3}$ and the required class 2 would be unreachable — the prototype of all the obstructions studied below.

The shape of what follows. Everything in this paper grows out of the picture just seen, read at three levels. *Locally* (one residual): whether the target class $-m$ is attainable from the factor classes of a is a question about a finite multiplicative state — the multiset of classes with bounded exponent budgets — so each fixed R admits an exact, machine-enumerable solvability criterion (Theorems 1.2–1.5 below). This is what makes the census of Section 4 feasible: the head of the distribution of $R_{\min}$ is decided by a character test rather than a divisor search. *Globally* (all residuals at once): the conditions entering these criteria are quadratic characters, and reciprocity folds the condition at every residual into signs $\left(\frac{p}{q}\right)$ indexed by the primes q dividing the shifted values $(p + R)/4$ — failures at different residuals correlate because shifted values share prime factors (Theorem 1.9). *Quantitatively*: each criterion says that failure forces $(p + R)/4$ to avoid prime factors in many residue classes, a sifting condition; stacking residuals multiplies the savings (Theorem 1.12), though it provably cannot reach emptiness (Remark 5.3) — the honest endpoint of the method, discussed in Section 6.

1.2 Main results

Theorem 1.2 ($R = 3$; exact). *Let $p \equiv 1 \pmod{4}$ be prime with $3 \nmid p$, and $a = (p + 3)/4$. Then (1.2) holds for $R = 3$ if and only if a has a prime factor $q \equiv 2 \pmod{3}$. When it does, $k = q$ is a certificate.*

Modulo 3, character equals class and the criterion is elementary. For every other fixed residual we need to know that the search space is finite before we can enumerate it, which is the content of the next statement. We label it a theorem for reference, but it is bookkeeping on top of Bradford’s criterion, and finite-generation statements of this kind have a long history in the subject: Terzi [27] generated 198 solvable residue classes modulo 120120 algorithmically, Yamamoto’s [31, §3] whole apparatus is a finite generation procedure for solvable coverings, and

Salez [24] proves a completeness theorem for polynomial families. What is specific here is the finite datum: the capped multiset of factor classes of $(p + R)/4$ together with $p \bmod R$.

Theorem 1.3 (finite reduction). *For every fixed prime $R \equiv 3 \pmod{4}$, the validity of (1.2) at a hard prime p depends only on (i) the multiset of classes mod R of the prime factors of $a = (p + R)/4$, with multiplicities capped at $\lceil (R - 1)/2 \rceil$, and (ii) $p \bmod R$. Consequently an exact solvability criterion exists for every R and is computable by finite enumeration.*

Theorem 1.4 ($R = 7$, hard primes; exact). *Let p be a hard prime and $a = (p + 7)/4$. Then (1.2) holds for $R = 7$ if and only if a has a prime factor that is a quadratic non-residue modulo 7.*

Theorem 1.5 ($R = 11$, hard primes; exact). *Let p be a hard prime and $a = (p + 11)/4$; note $3 \mid a$, $2 \nmid a$, $7 \nmid a$ automatically. Then (1.2) fails for $R = 11$ if and only if either*

- (a) *every prime factor of a is a quadratic residue mod 11; or*
- (b) *$v_3(a) = 1$, every other prime factor of a is $\equiv 1 \pmod{11}$ except a non-residue part — its factors in the non-residue classes mod 11 — of exactly one of the shapes*

$$\{q \equiv 2\}, \quad \{q \equiv 6\}, \quad \{q \equiv 2\} \cup \{q' \equiv 6\} \pmod{11},$$

each with multiplicity one, paired respectively with $p \equiv 2, 6, 1 \pmod{11}$.

For composite R the certificate condition couples the prime-power components of R through the Chinese Remainder Theorem, and the divisor-class computation lives in the (generally non-cyclic) unit group $(\mathbb{Z}/R)^*$. Theorem 1.3 holds verbatim with $(\mathbb{Z}/R)^*$ in place of the cyclic group (multiplicities capped once exponent ranges cover each class's cyclic span; see Remark 3.1). The first composite case turns out to be *simpler* than $R = 11$: a pure character dichotomy, with the quadratic character modulo a prime replaced by the Jacobi character modulo 15, and no budget cases at all.

Theorem 1.6 ($R = 15$, hard primes; exact). *Let p be a hard prime and $a = (p + 15)/4$; note $p \bmod 15 \in \{1, 4\}$ and $2 \mid a$ automatically. Then (1.2) holds for $R = 15$ if and only if a has a prime factor q with $\left(\frac{q}{15}\right) = -1$, i.e. $q \equiv 7, 11, 13, 14 \pmod{15}$.*

Remark 1.7 (attribution at $R = 3, 7, 11, 15$). Theorems 1.2 and 1.4 are not new. In Yamamoto's notation [31, §3] the covering $\{-s/q\}$ — the set of n with $q \mid n + s$, equivalently, for odd q , $q \mid (n + s)/4$ — lies in his system Σ_1 of solvable coverings exactly for $q \equiv -1 \pmod{3}$ when $s = 3$, and for $q \equiv -1, -2, -4 \pmod{7}$ when $s = 7$; he records the common form explicitly, noting that “all the coverings $\{-s/q\}$ with the Jacobi symbol $\left(\frac{-s}{q}\right) = -1$ are contained in Σ_1 , for $s = 1, 2, 3, 4, 7$ and 8 ” [31, p. 45], and Σ_1 is defined on precisely the six hard classes. Since $\left(\frac{-s}{q}\right) = \left(\frac{q}{s}\right)$ for $s \equiv 3 \pmod{4}$, and since membership of a non-residue class is inherited by some prime factor, this is the sufficiency half of Theorems 1.2 and 1.4 verbatim; the necessity half is the character obstruction (Proposition 2.1), which is Yamamoto's Lemma 4 specialised. What the statements below add at these two residuals is only that the certificate can be taken at residual exactly R , and a machine-checked proof.

At $R = 11$ and $R = 15$ the situation is different, and it is worth being exact about the boundary. Yamamoto's Table 3 [31, p. 45] gives, for $s = 11$, the classes $q \equiv -1, -3, -4 \equiv 7, 8, 10 \pmod{11}$: precisely the three non-residue classes in which a single prime factor already forces

solvability. Theorem 1.5 supplies the complement — that $q \equiv 2$ and $q \equiv 6 \pmod{11}$ suffice as well, outside the explicitly described budget family of case (b) — and the exactness. Over the 1,005 hard primes below 4×10^5 , Yamamoto’s classes certify 413 and Theorem 1.5 certifies 628. Yamamoto does not treat $s = 15$ (his tables list prime s only), but his constructions (i)–(ii) apply verbatim there and yield the three classes $q \equiv 7, 13, 14 \pmod{15}$; Theorem 1.6 completes the Jacobi dichotomy with the fourth class $q \equiv 11 \pmod{15}$ and proves it exact (315 versus 355 primes certified in the same range). For $R \geq 19$ the divergence is no longer marginal: Yamamoto’s rules give 1, 3 or 9 classes according to $R \bmod 24$, against the $(R-1)/2$ non-residue classes reached by the criteria and support bounds below — 3 versus 53 at $R = 107$.

Remark 1.8 (computer-assisted components). The exhaustive verifications behind Theorems 1.4, 1.5, 1.6 and Lemma 1.10 below are *parts of the respective proofs*, on the same footing as a hand case analysis. Each is a deterministic enumeration of a finite configuration space whose finiteness and sufficiency are proved in advance (Theorem 1.3; for Lemma 1.10, monotonicity of target-reachability in support and multiplicities, checked at minimal exponents). The checks are short (seconds each, up to 331 s for the largest, $R = 107$), involve only exact integer arithmetic, and the code is available in the accompanying repository; Table 5 records the exact configuration counts and entry points. The enumerations at $R = 7, 11$ and the $R = 19, 23, 31$ cases of Lemma 1.10 are additionally formalized in Lean 4, with the reductions they rely on proved as theorems rather than trusted (the $R = 15$ enumeration and the remaining cases of Lemma 1.10 run in the Python layer only); Appendix A gives the inventory and the exact trust base.

The obstruction framework of Section 2 yields a reciprocity identity that converts the per-residual conditions into a single global statement about p — a structural interpretation of joint failure across residuals:

Theorem 1.9 (reciprocity structure). *Let p be a hard prime, $R \equiv 3 \pmod{4}$ prime, and q an odd prime factor of $a = (p + R)/4$ (necessarily $q \neq R$). Then*

$$\left(\frac{q}{R}\right) = \left(\frac{p}{q}\right).$$

The sieve consequences below need one input per residual: a statement that failure at R forces the shifted form $(p + R)/4$ to avoid prime factors in a positive-density set of residue classes. For $R = 3, 7, 11, 15$ the exact criteria provide this; for every other residual it is the support bound, which we state first in the per-residual machine-verified form in which it was originally established (Lemma 1.10) and then prove unconditionally for every residual at once (Theorem 1.11; both proved in Section 5.1).

Lemma 1.10 (support bound; verified for every prime $R \equiv 3 \pmod{4}$, $19 \leq R \leq 107$). *Let R be one of the twelve primes 19, 23, 31, 43, 47, 59, 67, 71, 79, 83, 103, 107, and call the set of unit classes mod R occupied by the prime factors of $(p + R)/4$ its support. Every configuration for which residual R fails at a hard prime has support of size at most $(R-1)/2$, at most $(R-3)/2$ of whose classes differ from the identity class (which never affects solvability). Consequently failure at R forbids at least half the classes, through finitely many maximal-support branches.*

Theorem 1.11 (unconditional support bound). (i) *Let $d \geq 4$ be even and $S \subseteq \mathbb{Z}/d \setminus \{0\}$. If the bounded subset-sum set*

$$M(S) = \left\{ \sum_{v \in S} e_v v : e_v \in \{0, 1, 2\} \right\}$$

is not all of $\mathbb{Z}/d$, then $|S| \leq d/2 - 1$, and the bound is attained ($S =$ the nonzero even classes). (ii) Consequently, for every prime $R \equiv 3 \pmod{4}$ and every p : if the prime factors of $(p+R)/4$ occupy at least $(R-1)/2$ nonzero unit classes mod R , then every unit class is a divisor class of m^2 and (1.2) holds; failure forces the conclusion of Lemma 1.10, with no upper limit on R . (iii) In a finite abelian group G of even order g with t involutions (multiplicative form, $M(S) = \prod_{v \in S} \{1, v, v^2\}$): if $M(S) \neq G$ then $|S| \leq \max(\lfloor (g+t-2)/2 \rfloor, g/2 - 1)$. For $G = (\mathbb{Z}/R)^*$ with $\omega(R) \leq 2$ this is $\leq g/2$, so failure at any admissible composite R forbids at least $\varphi(R)/2 - 1$ unit classes on $(p+R)/4$ — sieve dimension $\geq \frac{1}{2} - 1/\varphi(R)$. For the admissible composites $R \leq 107$ an exhaustive enumeration (`kneser_support_general`, a proof component of this clause) sharpens the maximal failing support to $g/2 - 1$, so those residuals forbid at least $\varphi(R)/2$ classes and join the chain with full dimension $\geq \frac{1}{2}$ each. For arbitrary $\omega(R)$ the same bound reads $|S| \leq \max(\lfloor (g+t-2)/2 \rfloor, g/2 - 1)$ with $t = 2^{\omega(R)} - 1$, so failure forbids at least $(g-t)/2 - 1 \geq \varphi(R)/2 - 2^{\omega(R)-1} - 1$ unit classes — dimension $\geq \frac{1}{2} - (2^{\omega(R)-1} + 1)/\varphi(R)$, and since $2^{\omega(R)} = R^{o(1)}$ this is ≥ 0.44 for every admissible R beyond an absolute bound: every residual, with no restriction on ω , is a usable sieve input.

Theorem 1.12 (the chain). *The number of hard primes $p \leq x$ failing every residual in $\{3, 7\}$, in $\{3, 7, 11\}$, in $\{3, 7, 11, 15\}$, and in $\{3, 7, 11, 15, 19\}$ is, respectively,*

$$O\left(\frac{x}{(\log x)^2}\right), \quad O\left(\frac{x}{(\log x)^{5/2}}\right), \quad O\left(\frac{x}{(\log x)^3}\right), \quad O\left(\frac{x}{(\log x)^{7/2}}\right).$$

More generally, by Theorem 1.11, for every finite set $\mathcal{P}$ of admissible residuals $R \equiv 3 \pmod{4}$ — prime, or composite with $\omega(R) \leq 2$ — the number of hard primes failing every residual in $\mathcal{P}$ is $O_{\mathcal{P}}\left(x/(\log x)^{1+\kappa(\mathcal{P})}\right)$, where $\kappa(\mathcal{P}) = \sum_{R \in \mathcal{P}} \delta_R$ with $\delta_R = \frac{1}{2}$ for R prime or composite ≤ 107 and $\delta_R = \frac{1}{2} - 1/\varphi(R)$ for composite $R > 107$ (the two tiers of Theorem 1.11(iii)); in particular $\kappa(\mathcal{P}) = |\mathcal{P}|/2$ for every list drawn from the residuals up to 107. For the fifteen primes $\equiv 3 \pmod{4}$ up to 107 the exponent is $17/2$, and for the full 27-element admissible list $\{3, 7, 11, \dots, 107\}$ it is $29/2$. All but a proportion $O((\log x)^{-2})$ of hard primes are solved within $\{3, 7, 11, 15\}$ — one power of $\log x$ less than the counts above, since there are $\asymp x/\log x$ hard primes up to x .

Proposition 1.13 (aggregate identity families). *Let $p \equiv 1 \pmod{4}$ be prime and $R \equiv 3 \pmod{4}$. If $R \mid p+1$, then $k = ap^2$ is a certificate for residual R ; if $R \mid p+4$, then $k = a^2p$ is. Hence any prime factor $\equiv 3 \pmod{4}$ of $p+1$ or of $p+4$ yields an explicit solution, and no divisor shape $k = a^j p^i$ yields a further p -only condition. Adding the two corresponding linear forms to Theorem 1.12: the number of hard primes $p \leq x$ failing every prime residual up to 107 and both families is $O(x/(\log x)^{19/2})$; with the full 27-element admissible list the exponent is $31/2$.*

Theorem 1.14 (density reduction). *For every $A > 0$ there is a finite admissible set $S(A)$ such that the number of hard primes $p \leq x$ for which every residual in $S(A)$ fails is $O_A\left(x/(\log x)^{1+A}\right)$.*

Remark 1.15 (relation to Vaughan's bound). As exceptional-set estimates, Theorem 1.14, the chain of Theorem 1.12, and Theorem 5.5 are all weaker than Vaughan's bound $\ll x \exp(-c(\log x)^{2/3})$ [29], which already exceeds every fixed power of $\log x$. They are not offered as improvements on it, and the honest comparison needs one further step, because Vaughan's argument is *constructive*: it proceeds by covering congruences, and every n it certifies comes

with an explicit solution and hence with an explicit admissible residual, bounded by a power of his sifting parameter. Read that way it already gives an almost-all bound on $R_{\min}$ — roughly $R_{\min}(n) \leq \exp(O((\log x)^{1/3}))$ outside a set of size $x \exp(-c(\log x)^{2/3})$, and $R_{\min}(n) \leq (\log x)^\theta$ outside $x \exp(-c(\log \log x)^2)$ when the threshold is set to match Theorem 5.5’s exceptional set. So Theorem 5.5 is not the first almost-all bound on $R_{\min}$. Its content is that at comparable exceptional-set quality it lowers the threshold from a power of $\log x$ to a constant multiple of $\log \log x$ — an exponential gain in the certified residual, which is the quantity Proposition 1.1 shows *is* the conjecture at p . Two further differences of kind: the exceptional sets here are cut out by finitely many *explicit* multiplicative conditions on shifted forms, each failure constructively certified by an exhibited divisor k , with the underlying lemmas exact or machine-verified; and the certificates are of a single controlled shape, indexed by R , rather than assembled from a covering system. A mean-value argument, by contrast, counts the primes that fail and produces no residual at all for the ones that succeed.

Organization. Section 2 develops the obstruction framework (Propositions 2.1, 2.2) and proves Theorem 1.9, whose Corollary 2.3 depends on the character/II dichotomy introduced there. Section 3 proves Theorem 1.2, then Theorem 1.3, then Theorems 1.4, 1.5 and 1.6, which rest on it. Section 4 is the census: data, verification protocol, the record, and the calibrated model. Section 5 is the structural context — Section 5.1 proves Theorem 1.12 (base case first, then, after the monotonicity reduction (Lemma 5.1), Lemma 1.10 and its unconditional upgrade Theorem 1.11, the higher exponents), followed by Proposition 1.13 and Theorem 1.14; Section 5.2 proves the uniform chain (Theorem 5.5); Section 5.3 localizes what stands between it and a polylogarithmic bound; and Section 5.4 sketches, in compressed form, the one conditional program that would pass the ceiling.

Table 1 records, for each main result, exactly what its proof consumes and its formal-verification status, so that each layer can be audited independently.

Three printed results are retained as stepping stones although later results subsume them, and the reader may want the map stated once: Theorem 1.14 is superseded by Theorem 5.5 (which gives $(\log x)^{-A}$ for every A at once), but its fixed-list proof is simpler and is the template the uniform chain scales; Lemma 1.10 is subsumed by Theorem 1.11, and is retained as the machine-verified confirmation — it is the form carried by the Lean development; and Theorem 5.16 is dominated unconditionally by Theorem 5.5 (the fixed chain beats the p -adapted ladder unconditionally), its role being structural: it opens the conditional program of Section 5.4, the only known frame whose measured inputs would pass the ceiling of Remark 5.7.

The census of Section 4 also calibrates the base-case constant and the decay exponents of the chain, and it is where the per-residual criteria of Section 3 are exercised at scale.

2 The obstruction framework

Proposition 2.1 (character obstruction). *Let $r \mid R$ be prime with $r \equiv 3 \pmod{4}$. If every prime factor q of $m = pa$ satisfies $\left(\frac{q}{r}\right) = +1$, then (1.2) fails for R .*

Proof. Every divisor k of m^2 is a product of prime factors of m , so $\left(\frac{k}{r}\right) = +1$. A certificate requires $k \equiv -m \pmod{r}$, whence $\left(\frac{k}{r}\right) = \left(\frac{-m}{r}\right) = -1$ by (1.4) — impossible. $\square$

We call failures produced by Proposition 2.1 *character failures*; failures where m possesses a non-residue prime factor modulo every relevant r , yet no divisor of m^2 attains the exact class

$-m \pmod R$, are *budget failures*. (We avoid “character failure/II”, which Elsholtz and Tao [4] use for the two principal *solution* parametrizations.)

Proposition 2.2 (guaranteed-success classes). *Let $t \equiv -4^{-1}p^2 \pmod R$. If a has a prime factor q with $q \equiv tp^{-i} \pmod R$ for some $i \in \{0, 1, 2\}$, then $k = qp^i$ is a certificate. In particular, since $tp^{-2} \equiv -4^{-1} \pmod R$, any prime factor of a in the fixed class $-4^{-1} \pmod R$ certifies success regardless of p .*

Proof. $k = qp^i$ divides m^2 (as $q \mid a$ and $i \leq 2$) and lies in class $t \equiv -m$. Integrality of b, c is automatic from $k \equiv -m$ and $m^2/k \equiv m^2(-m)^{-1} \equiv -m \pmod R$. $\square$

The contrapositive of Proposition 2.2 is the sieve input for Theorem 1.14: failure at R forces $(p + R)/4$ to avoid prime factors in up to three explicit classes modulo R .

The next identity converts the per-residual character conditions into a single global statement about p , collapsing joint failure across residuals into one set of quadratic-character conditions.

Proof of Theorem 1.9. From $q \mid p + R$ we get $R \equiv -p \pmod q$. By quadratic reciprocity with $R \equiv 3 \pmod 4$, $\left(\frac{q}{R}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{R}{q}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{-1}{q}\right) \left(\frac{p}{q}\right) = \left(\frac{p}{q}\right)$. $\square$

Corollary 2.3 (forced primes and correlated failure). *(i) The hard classes are the squares mod 840, so p is automatically a quadratic residue mod 3, 5, 7, and $p \equiv 1 \pmod 8$ handles $q = 2$ (which divides a only for $R \equiv 7 \pmod 8$, where $\left(\frac{2}{R}\right) = +1$): no prime forced into a shifted value $(p + R)/4$ by the congruences (H1)–(H3) carries a negative character. (ii) p fails character at every prime residual $R \leq B$ simultaneously if and only if $\left(\frac{p}{R}\right) = +1$ for each such R and $\left(\frac{p}{q}\right) = +1$ for every unforced odd prime q (one not forced by the congruences of (H1)–(H3)) dividing any of the shifted values $(p + R)/4$, $R \leq B$ — one condition per distinct prime, shared across every residual it divides.*

Proof. Immediate from Theorem 1.9 together with (H1)–(H3) and the quadratic characters of the forced small primes. $\square$

In words, part (i) says that the small primes forced into the shifted values by the Mordell structure never break a character-failure — a conceptual reading of why the hard classes are exactly the square classes. By part (ii), failures at distinct residuals correlate through shared prime factors: shifted values with few distinct prime factors leave few independent sign conditions. Section 4 quantifies this for the record prime. Both parts are verified at scale (135,588 instances of Theorem 1.9, 61,438 forced-prime instances, zero violations of either; Appendix A).

3 Exact criteria

The individual criteria proved in this section are instances; the conceptual point is the framework behind them — every residual, prime or composite, admits a finite computable exact criterion (Theorem 1.3), and every criterion doubles as a sieve input. The four residuals worked out here are the ones the computations show carry 95% of the hard primes.

3.1 Proof of Theorem 1.2

Proof. Modulo 3 the unit group has order 2, so quadratic character equals class. Since $p^2 \equiv 1 \pmod{3}$, (1.3) gives $m \equiv 1$ and the target is $-m \equiv 2 \pmod{3}$; also $3 \nmid a$. *Necessity:* if every prime factor of a is $\equiv 1 \pmod{3}$, then $a \equiv 1 \pmod{3}$, hence $p = 4a - 3 \equiv 1 \pmod{3}$ as well, so every prime factor of $m = pa$ is $\equiv 1 \pmod{3}$ and every divisor of m^2 is $\equiv 1 \pmod{3}$: class 2 is never attained. *Sufficiency:* $k = q \equiv 2 \pmod{3}$ divides m^2 , and the recovery formulae give integers as in Proposition 2.2. $\square$

3.2 Proof of Theorem 1.3

Proof. A certificate is a divisor $k \mid m^2$ in class $-m \bmod R$. The set of classes of divisors of m^2 is determined by the classes of the prime factors of $m = pa$ together with their multiplicities, each factor being usable with exponent at most twice its multiplicity; exponents are only meaningful modulo $\text{ord}_R(\text{class}) \leq R - 1$, so multiplicities beyond $\lceil (R - 1)/2 \rceil$ are equivalent to that cap (Remark 3.1 illustrates why). The target class and the consistency relation $\prod(\text{classes}) \equiv a \equiv 4^{-1}p \pmod{R}$ are functions of $p \bmod R$ by (1.3). Hence solvability is a function of the capped class multiset and $p \bmod R$, a finite amount of data. $\square$

Remark 3.1 (why the cap works, and the size of the space). The one step above that is not pure bookkeeping is the multiplicity cap, and a small example shows both why it is harmless and what it buys. Take $R = 7$ and a factor of a in a class x of order 6 (say $x = 3$): a factor of multiplicity μ may enter a divisor of m^2 with any exponent $0, 1, \dots, 2\mu$, and once $2\mu \geq 6$ these powers already sweep out the full cyclic group $\langle x \rangle$ — raising μ past $\lceil (R - 1)/2 \rceil = 3$ enlarges nothing that a divisor subproduct can reach. For classes of smaller order the cap is generous rather than tight (a class of order 2 stabilizes at $\mu = 1$), which is fine: the enumeration ranges over a finite over-approximation. What finiteness buys is concrete: at $R = 7$ the capped space has 4^6 multisets before the structural constraints (i)–(iii) in the proof of Theorem 1.4 below and 1,536 consistent configurations after them; at $R = 11$ it compresses to 25 equivalence states. The enumeration is decidable, not merely heuristic, precisely because the cap and the consistency relation are proved before any code runs.

3.3 Proof of Theorem 1.4

Proof. $(\Rightarrow)$ is Proposition 2.1 with $r = R = 7$. $(\Leftarrow)$ reduces to a finite verification via Theorem 1.3. Write $(\mathbb{Z}/7)^* \cong \mathbb{Z}/6$ in discrete-log coordinates. Three structural constraints hold at hard primes:

- (i) $p \bmod 7 \in \{1, 2, 4\}$, by (H3);
- (ii) $2 \mid a$: by (H1), $p + 7 \equiv 0 \pmod{8}$, so $a = (p + 7)/4$ is even — the class of 2 (a residue mod 7) is always present;
- (iii) the factor classes multiply to $a \equiv 4^{-1}p \pmod{7}$.

Enumerating all class multisets with multiplicities capped at 3 subject to (i)–(iii) yields 1,536 consistent configurations; for each, one checks whether the target class is attained by a bounded-exponent subproduct. The verification (code: `theory.verify_R7_finite`) confirms: the target is attainable precisely when some factor class is a non-residue. $\square$

Remark 3.2. The hard-class hypotheses are essential: for $p = 701 \equiv 1 \pmod{4}$ (not hard), $a = 177 = 3 \cdot 59$ has the non-residue factor 3, yet $R = 7$ fails. Constraint (ii) — the forced even factor — is exactly what rescues every would-be budget configuration at hard primes. Empirically (Section 4), among 41,421 sampled failures of $R = 7$ at hard primes, *all* are character failure.

3.4 Proof of Theorem 1.5

Proof. By Theorem 1.3 with $R = 11$ (the unit group has order 10): multiplicities cap at 5, and at hard primes $3 \mid a$ (from (H2), $11 \equiv 2 \pmod{3}$), while $2 \nmid a$ (from (H1), $p + 11 \equiv 4 \pmod{8}$) and $7 \nmid a$ (from (H3), $-11 \pmod{7} = 3 \notin \{1, 2, 4\}$). A dynamic-programming enumeration over the finite configuration space (code: `theory.finite_criterion_dp`) compresses it to 25 equivalence states: 16 succeed, 6 fail with all factor classes quadratic residues (case (a)), and exactly 3 fail with a non-residue present (case (b)); the three failing states decode to the three listed patterns. Two structural facts emerge from the enumeration. First, no state fails by a *proper-subgroup* obstruction: the only subgroup of $\mathbb{Z}/10$ containing an odd element properly is $\{0, 5\}$, and the consistency relation forces p 's class to extend it to all of $\mathbb{Z}/10$. Second, the case-(b) failures are pure *budget* failures — the exponent bounds, not the character obstruction, block the target: any of $v_3(a) \geq 2$, a factor in one of the classes 4, 5, 9, or a second non-residue factor restores solvability. $\square$

3.5 Proof of Theorem 1.6

Proof. ($\Rightarrow$, the Jacobi obstruction.) By (1.3), $m \equiv 4^{-1}p^2 = (2^{-1}p)^2 \pmod{15}$ is the square of a unit, so the Jacobi symbol satisfies $\left(\frac{m}{15}\right) = +1$; and $\left(\frac{-1}{15}\right) = \left(\frac{-1}{3}\right)\left(\frac{-1}{5}\right) = -1$. Hence $\left(\frac{-m}{15}\right) = -1$. The classes with $\left(\frac{q}{15}\right) = +1$ are exactly $\{1, 2, 4, 8\} = \langle 2 \rangle$, a subgroup of $(\mathbb{Z}/15)^*$ containing p 's class ($p \pmod{15} \in \{1, 4\}$, since $p \equiv 1 \pmod{3}$ by (H2) and $p \equiv \pm 1 \pmod{5}$ for all six hard classes). If every prime factor of a lies in $\langle 2 \rangle$, then every divisor of m^2 does, and the target $-m \equiv 11 \pmod{15}$ (an equality valid for both p -classes) is unreachable. This composite-modulus form of Proposition 2.1 is strictly stronger than the proposition applied to the single prime $r = 3 \mid 15$.

($\Leftarrow$.) By the composite form of Theorem 1.3 over $(\mathbb{Z}/15)^* \cong \mathbb{Z}/2 \times \mathbb{Z}/4$ (Carmichael exponent $\lambda(15) = 4$, so reachability saturates at multiplicity 2; multiplicities 0..5 capture every pair of saturation state and class product). The structural constraints at hard primes are: $p \pmod{15} \in \{1, 4\}$; $2 \mid a$ (from (H1), $8 \mid p + 15$), so the class of 2 always appears; and consistency $\prod(\text{classes}) \equiv a \equiv 4p \pmod{15}$. Enumerating the full space (code: `theory.verify_R15_finite`) yields 349,920 consistent configurations: 349,380 succeed and 540 fail — and *every* failing configuration has all factor classes in $\langle 2 \rangle$. In contrast to $R = 11$ there are no budget failures at all: the consistency relation forces $\prod\left(\frac{q}{15}\right)^{\mu_q} = \left(\frac{4p}{15}\right) = +1$, so non-residue factors occur with even total multiplicity, and the enumeration confirms that any two non-residue units (or one squared), together with the forced powers of 2, always reach the target. $\square$

Remark 3.3. The general composite- R engine (`theory.solvable_exact_general`), which evaluates the criterion of the composite finite reduction directly in $(\mathbb{Z}/R)^*$, agrees with ground-truth divisor search on *every* composite admissible residual $R \in \{15, 27, 35, \dots, 99\}$ (7,938 sampled hard primes per residual below 10^9 , zero disagreements).

4 The census to 10^{12}

This section is the paper’s main object: the value of $R_{\min}(p)$ for every hard prime below 10^{12} . For each p the admissible residuals are tried in increasing order; at each R the criterion is the exact one of Section 3 where available, and otherwise a direct integer factorization of $a = (p + R)/4$ (trial division; the factorization of $m^2 = p^2 a^2$ follows) with a scan of the divisor classes for (1.2). The first R that succeeds is $R_{\min}(p)$, and the certificate it yields is re-verified in exact integer arithmetic as $4abc = p(bc + ac + ab)$. All code and data are in the accompanying repository; the 10^{12} run is a single 10.5 h job on 32 cores, reproduced by `scripts/run_1e12.sh`.

4.1 Verification protocol

The claims of this section rest on the datasets of Table 2, all archived in the repository under `data/` with sha256 manifests. A certificate (R, a, b, c) for p is *accepted* only if $4abc = p(bc + ac + ab)$ holds in exact integer arithmetic and $R = 4a - p$; for the compact minimal- R maps, the verifier *reconstructs* the triple from (p, R) by the divisor search of (1.2) and then applies the same exact check, so a map entry is verified as a genuine certificate, not taken on faith.

The masks dataset of Table 2 is the basis for the covering-set statements (the 18-residual cover, the lower bound 12, the four critical primes), the failure taxonomy, and the calibration below; the minimal- R maps are the basis for Table 3, the record, and the gap analysis.

4.2 Data

Every one of the 1,175,215,396 hard primes $p < 10^{12}$ possesses a residual certificate with $R \leq 111$. The verification behind this claim is graded by dataset: exhaustive reconstruction and exact re-checking of every entry below 10^9 ; systematic samples (143,000, 257,343 and 2,350,431 entries) at 10^{10} , 10^{11} and 10^{12} , together with the *complete* tails (842 entries with $R \geq 51$, 20,151 and 118,210 with $R \geq 43$) re-checked with full minimality — Table 2 gives the per-dataset detail. The minimal residuals distribute as in Table 3 and Figure 1.

Three features of the distribution deserve comment.

- **The maximum moves at 10^{12} .** Throughout $10^7 \leq x \leq 10^{11}$ — five successive orders of magnitude — the maximal minimal residual over hard primes $p \leq x$ stood at $R = 107$, attained by the single prime $8,803,369 < 10^7$. At 10^{12} it is $R_{\max} = 111$, attained by three primes, the least being

$$p = 119,945,383,009 \equiv 529 \pmod{840},$$

for which every admissible $R \leq 107$ fails and $R = 111$ certifies with $a = (p + 111)/4 = 29,986,345,780$. The old record 107 also ceased to be unique, acquiring three further primes in the same decade. The prime 8,803,369 is extremal for a closely related statistic as well: in independent contemporaneous work, Bello-Hernández, Benito and Fernández [10] report it as the largest “survivor” below 10^7 in a $k \equiv 3 \pmod{4}$ divisor sieve, requiring $k = 107$ there too. In the language of Corollary 2.3, 8,803,369’s 27 shifted values contain only 43 distinct odd primes (typical: 45–52), and 81% of them satisfy $\left(\frac{p}{q}\right) = +1$ (typical: 44–60%) — an extreme value of exactly the statistic Theorem 1.9 identifies as governing joint failure. We attach no significance level to this: the prime was selected *because* it was the record, so its extremeness in a correlated statistic is expected rather than surprising. The




Figure 1: The data of Table 3 on a logarithmic count scale. The head is dominated by $R = 3, 7, 11$; the tail thins geometrically; the shaded band 87–111 was empty below 10^{10} , and its population at 10^{11} and 10^{12} followed the calibrated prediction (see the prediction bullet below), including the move of the maximum from $R = 107$ to $R = 111$.

observation is descriptive — it says what a record prime looks like mechanically, not that its existence is improbable.

- A prediction, and what it was worth.** The *calibrated independence model* treats the per-residual failure events as independent across R , with per-residual failure rates fitted from the complete 10^9 solvability masks and a single deep-tail correlation factor calibrated on the $R > 83$ counts; it is a heuristic, not a theorem, and its role is to make falsifiable predictions. Below 10^{10} no prime had minimal $R \in \{87, 91, 95, 99, 103\}$; the model attributed this gap to small-number statistics (expected counts ≤ 0.12 per value at 10^9) and predicted that deep-tail primes (failing every $R \leq 83$) land at $87, 91, 95, 99, 103, \geq 107$ with probabilities $0.53, 0.11, 0.22, 0.03, 0.05, 0.06$ (the last being the total mass at or beyond 107). Between 10^{10} and 10^{11} the five values acquired 8, 3, 5, 1, 1 primes respectively (all in $(1.3 \times 10^{10}, 10^{11})$), and none passed 103; that distribution was *consistent* with the prediction ($\chi^2 \approx 1.3$) but, on 19 points with several expected cells below 5, had little power to discriminate. The model's sharper commitment was about *when the record would move*: the mass at or beyond 107 was put at 6% per deep-tail prime, which with the deep-tail growth rate places the first record-breaking prime in the decade 10^{11} – 10^{12} . It broke there, at $R = 111$. The deep tail now holds 73 primes with $R_{\min} \geq 87$ against 19 a decade earlier, distributed 40, 5, 12, 4, 5, 4, 3 over $87, \dots, 111$. A model fitted at 10^9 thus predicted, before the computation, both the filling of a band that was empty at the time and the decade in which the maximum would move for the first time in five orders of magnitude.

- Among all 1,587,581 hard primes below 10^9 , exactly *four* admit a unique working residual

≤ 107 (8,803,369 with 107; 142,361,209 with 59; 287,567,281 with 83; 794,037,841 with 63)
— the entire obstruction to a shorter covering list. A greedy cover of the full solvability table uses 18 residuals; a packing argument gives a lower bound of 12.

4.3 Calibration of the base case of Theorem 1.12

The proportion of hard primes failing both $R = 3$ and $R = 7$, multiplied by $\log x$, is constant to within 1% across three decades of the 10^9 data (Table 4).

That is, the joint exceptional density is $\approx 5.17/\log x$: the decay matches the order predicted by the base case of Theorem 1.12, and 5.17 estimates the (conjectural) Selberg–Delange constant. The exceptional set is density zero but decays slowly: 26.5% of hard primes below 10^9 fail both residuals; all are settled by deeper residuals of the list. The fitted per-residual decay exponents are consistent with the sieve prediction $\kappa = \frac{1}{2}$ where exact criteria exist: $\hat{\kappa} = 0.539$ for $R = 7$, and $\hat{\kappa} = 0.387$ for $R = 3$, where the short fitting range and Selberg–Delange lower-order terms bias the exponent downward.

4.4 Failure taxonomy

Sampling the complete solvability table at 10^9 : failures at $R = 7$ are 100% character failures (as Theorem 1.4 requires); at $R = 11$, 89.7% character failure and 10.3% the budget pattern of Theorem 1.5(b); the character-failure share broadly decreases with $\varphi(R)$ (e.g. 51% at $R = 103$, though not monotonically), as hitting one specific class among $\varphi(R)$ grows harder while the character dichotomy stays binary. Taken together, the data of this section matches the framework of Sections 2–5.1: each distributional feature — the head, the geometric tail, the once-empty band, the static record — has a structural or heuristic explanation (the exact criteria for the head; the calibrated model, built on the sieve dimensions, for the tail and the band; Theorem 1.9 for the record).

5 The sieve chain and its ceiling

5.1 The chain

We now prove the chain of Theorem 1.12, beginning with the base case, and then Lemma 1.10, its unconditional upgrade Theorem 1.11, the aggregate identities, and the density reduction.

Proof of Theorem 1.12, base case $\{3, 7\}$. Step 1 (exact criteria). Let p be a hard prime and set $n = (p + 3)/4$, so that $(p + 7)/4 = n + 1$: the two shifted forms are consecutive integers. By Theorems 1.2 and 1.4, p is exceptional (both residuals fail) precisely when

$$n \text{ has no prime factor } \equiv 2 \pmod{3} \quad \text{and} \quad n + 1 \text{ has no prime factor } \equiv 3, 5, 6 \pmod{7},$$

with $p = 4n - 3$ prime.

Step 2 (congruence bookkeeping). Fix one of the six hard classes $h \pmod{840}$; then $n \equiv (h + 3)/4 \pmod{210}$ is fixed, and the conditions at the primes 2, 3, 5, 7 are determined by the class (in particular $3 \nmid n$ and $7 \nmid n + 1$ automatically). It suffices to bound each class and sum.

Step 3 (containment in a sifted set). Let $z = x^{1/10}$. After discarding the $O(z)$ primes $p \leq z$, every exceptional p gives $n \leq (x + 3)/4$, $n \equiv n_0 \pmod{210}$, such that for every prime q with $7 < q \leq z$:

- (a) $q \equiv 2 \pmod{3} \implies n \not\equiv 0 \pmod{q}$;
- (b) $q \equiv 3, 5, 6 \pmod{7} \implies n \not\equiv -1 \pmod{q}$;
- (c) $q \nmid 4n - 3$ (since $4n - 3 = p$ is prime, $p > z$).

Conditions (a), (b) only use “no forbidden factor up to z ”, which is implied by the full criteria — an upper bound needs no more.

Step 4 (sieve of dimension 2). This is a standard upper-bound sieve for the three linear forms n , $n + 1$, $4n - 3$, with sifting function

$$\omega(q) = \mathbf{1}_{q \equiv 2(3)} + \mathbf{1}_{q \text{ NQR}(7)} + 1 \quad (q > 7),$$

whose forbidden classes $0, -1, 3 \cdot 4^{-1} \pmod{q}$ are pairwise distinct for $q \nmid 21$. By Dirichlet, ω has average $\kappa = \frac{1}{2} + \frac{1}{2} + 1 = 2$ in the Halberstam–Richert sense $\sum_{q \leq w} \omega(q) \log q / q = \kappa \log w + O(1)$. The Fundamental Lemma of sieve theory (see [8, Thm. 2.5] and [6, Ch. 6]) with $z = x^{1/10}$ bounds the exceptional set $E(x)$ of hard primes $p \leq x$ failing both residuals:

$$\#E(x) \ll x \prod_{7 < q \leq z} \left(1 - \frac{\omega(q)}{q}\right) \asymp x (\log z)^{-1} (\log z)^{-1/2} (\log z)^{-1/2} \asymp \frac{x}{(\log x)^2},$$

by Mertens’ theorem and its arithmetic-progression form. Summing over the six classes preserves the bound. The higher exponents of the chain are deduced after Lemma 1.10 below. $\square$

The verification behind Lemma 1.10 reduces to finitely many checks through the following elementary but load-bearing observation, which we isolate for scrutiny.

Lemma 5.1 (monotonicity reduction). *For a class multiset C (classes v with multiplicities μ_v) let $D(C) = \{\sum_v e_v v : 0 \leq e_v \leq 2\mu_v\}$ denote the set of attainable divisor logs (before the p -shifts). If $\mu_v \leq \mu'_v$ for all v , then $D(C) \subseteq D(C')$, where C' carries multiplicities μ'_v on the same classes; likewise enlarging the support enlarges D . Consequently a support S carries some failing configuration (any multiplicities ≥ 1 on S) if and only if the all-multiplicity-one configuration on S fails.*

Proof. Any admissible exponent vector for C is admissible for C' , giving the inclusions. If a configuration on S with multiplicities $\mu_v \geq 1$ fails, the target lies outside $D(C)$, hence outside the subset $D(C_1)$ attained at multiplicities one — so C_1 fails too; the converse inclusion is trivial. The same argument applies verbatim after adjoining the p -shifts. $\square$

Proof of Lemma 1.10. By Lemma 5.1 it suffices to check that for every support of size $(R - 1)/2$ over the nonzero discrete logs and every class of $p \pmod{R}$, the target is reachable already at minimal multiplicities (exponent budget 2 per class, p -budget 2); success there kills every configuration on every superset support at any multiplicities. The check deliberately ignores the consistency relation $\prod(\text{classes}) \equiv 4^{-1}p$: it therefore examines a *superset* of the realizable configurations, and can only overstate the failing supports — the lemma is conservative. For $R = 19, 23$ a direct enumeration of all $(R - 1) \binom{R-2}{(R-1)/2}$ (support, class of p) pairs suffices (437,580 resp. 7,759,752 pairs; code: `theory.verify_support_bound`). For the larger residuals the binomial count explodes, and the check is instead performed by a subset dynamic program (`theory.verify_support_bound_dp`): fold all supports into a table mapping each reachability mask to the maximal support size achieving it — valid because, by the same monotonicity,

failure at any multiplicities implies failure at multiplicity one. The realized masks are heavily structured, so the state count stays tractable (3,001 states at $R = 31$ against 77.6 million supports; 19,530,971 states at $R = 107$, 331 seconds). All twelve residuals verify with zero violations; exact state counts are tabulated in Table 5. The bound is tight: the all-residue (character failure) configurations occupy exactly $(R - 3)/2$ nonzero classes. $\square$

Theorem 1.11 is a statement of the critical-number literature, and we should say so before proving it. The *critical number* $\text{cr}(G)$ of a finite abelian group is the least t such that every $A \subseteq G \setminus \{0\}$ with $|A| \geq t$ has $\Sigma(A) = G$, where $\Sigma(A)$ is the set of subset sums. Diderrich and Mann [13] proved that $\text{cr}(G) = |G|/2$ for every abelian G of even order, with exactly four exceptions — $\mathbb{Z}/4$, $\mathbb{Z}/6$, $\mathbb{Z}/8$ and $\mathbb{Z}/2 \times \mathbb{Z}/4$ — and the determination for all finite abelian groups was completed by Freeze, Gao and Geroldinger [14]; see also Gao–Hamidoune–Lladó–Serra [7] for the covering form. Since $M(S) \supseteq \Sigma(S)$, parts (i) and (iii) of Theorem 1.11 follow from those theorems wherever they apply, and the extremal family (the nonzero even classes) is theirs. More precisely $M(S) = \Sigma(S \uplus S)$, the set of subsequence sums of the sequence in which each element of S occurs twice, so everything in this subsection and the next lives inside the zero-sum / Davenport-constant circle of ideas (see [12, Ch. 6, 16]).

Three things justify giving the proof anyway. First, the exceptions have to be handled, and the exponent budget removes them. In the residual setting $G = (\mathbb{Z}/R)^*$, and among admissible $R \equiv 3 \pmod{4}$ the exceptional groups occur at exactly $R = 7$ (giving $\mathbb{Z}/6$) and $R = 15$ (giving $\mathbb{Z}/2 \times \mathbb{Z}/4$); no admissible R has $\varphi(R) = 4$. At those two residuals the Diderrich–Mann bound $|S| \leq g/2 - 1$ is false for $\Sigma(S)$ — supports of size $g/2$ with $\Sigma(S) \neq G$ exist — but true for $M(S)$, because allowing each class twice is exactly enough to close the gap. The proof below gives $|S| \leq g/2 - 1$ for M at every even-order G with no exceptions, so the sieve dimension $\frac{1}{2}$ is available at *every* admissible residual. (At $R = 7$ and $R = 15$ this is belt and braces: Theorems 1.4 and 1.6 give an exact criterion there in any case.) Second, the summands here are not arbitrary: each factor class contributes the bounded-exponent set $\{1, v, v^2\}$, and the derivation below yields the quantitative form we need — a single prescribed target, bounds uniform in R , and the involution bookkeeping of part (iii) — directly. Third, it is short, and it is what the Lean development formalizes.

Proof of Theorem 1.11. (i) Write $M = A_1 + \cdots + A_k$ with $A_i = \{0, v_i, 2v_i\}$, $k = |S|$, and let $H = \text{Stab}(M) = \{h : M + h = M\}$, a subgroup of $\mathbb{Z}/d$. Kneser’s addition theorem [20] (in the many-summand form obtained by iterating [28, Thm. 5.5]) gives

$$|M| \geq \sum_{i=1}^k |A_i + H| - (k - 1) |H|.$$

Case $H = \{0\}$. Every $|A_i| = 3$, except that the (at most one) index i with $v_i = d/2$ has $|A_i| = 2$ (there $2v_i = 0$); so $|M| \geq (3k - 1) - (k - 1) = 2k$. Since $M \neq \mathbb{Z}/d$, $2k \leq d - 1$, and d even forces $k \leq d/2 - 1$. *Case $|H| = h > 1$.* M is a union of H -cosets and $M \neq \mathbb{Z}/d$, so $|M| \leq d - h$. Split $k_0 = |S \cap H| \leq h - 1$ (distinct nonzero elements of H) and $k_1 = k - k_0$. For $v \in H$, $|A_v + H| = h$; for $v \notin H$, $A_v + H \supseteq H \cup (v + H)$, so $|A_v + H| \geq 2h$. Kneser’s bound gives $d - h \geq h(k_0 + 2k_1 - k + 1) = h(k_1 + 1)$, so $k_1 \leq d/h - 2$ and $k \leq (h - 1) + (d/h - 2) = h + d/h - 3$. For every divisor h with $2 \leq h \leq d/2$, the quadratic $h^2 - (\frac{d}{2} + 2)h + d \leq 0$ (its roots are 2 and $d/2$), i.e. $h + d/h - 3 \leq d/2 - 1$. Tightness: for S the nonzero even classes, $M(S)$ is the even subgroup and $|S| = d/2 - 1$.

(ii) In discrete-log coordinates on $(\mathbb{Z}/R)^* \cong \mathbb{Z}/d$, $d = R - 1$, the divisor classes of m^2 contain the set $M(S)$ over the nonzero support S of $(p+R)/4$ at exponent budget 2 per class (Lemma 5.1: multiplicities ≥ 1 only enlarge budgets). If $|S| \geq d/2$, part (i) forces $M(S) = \mathbb{Z}/d$: every class — in particular the target $-m$ — is a divisor class of m^2 , without invoking the p -powers, and (1.2) holds. Contrapositively, failure at R forces $|S| \leq d/2 - 1 = (R - 3)/2$, which is the conclusion of Lemma 1.10 for arbitrary R .

(iii) The argument in (i) never used cyclicity. In a finite abelian G of even order g : in the H -trivial case each involution $v \in S$ has $|\{1, v, v^2\}| = 2$, so $|M| \geq (3k - t) - (k - 1) = 2k - t + 1$ and $M \neq G$ gives $k \leq \lfloor (g+t-2)/2 \rfloor$; the $|H| = h > 1$ case is verbatim, giving $k \leq h+g/h-3 \leq g/2-1$ over proper subgroups. For $G = (\mathbb{Z}/R)^*$, $t = 2^{\omega(R)} - 1$; with $\omega(R) \leq 2$, $t \leq 3$ and the bound is $\leq g/2$, so failure forbids at least $(g-1) - g/2 = \varphi(R)/2 - 1$ non-identity unit classes — dimension $\geq \frac{1}{2} - 1/\varphi(R)$. The general- t clause is the same computation without the specialization: failure forbids at least $(g-1) - (g+t-2)/2 = (g-t)/2$ non-identity classes (minus one when the max is attained by $g/2 - 1$), and $2^{\omega(R)} \leq R^{O(1/\log \log R)} = R^{o(1)}$ gives the stated 0.44 threshold beyond an absolute bound. The upgrade to $\varphi(R)/2$ forbidden classes for $R \leq 107$ is the enumeration cited in the statement: the maximum support of a proper product-set is exactly $g/2 - 1$ for every composite admissible $R \leq 107$ (`theory.kneser_support_general`; e.g. $R = 95$: $g = 72$, maximum 35 over 583,455 realizable product-sets). That enumeration is a proof component of the $R \leq 107$ clause, on the same footing as the checks of Remark 1.8; the $\frac{1}{2} - 1/\varphi(R)$ tier needs Kneser alone. $\square$

Remark 5.2. Theorem 1.11(ii) is strictly stronger than what the sieve needs: at support $\geq (R - 1)/2$ every class is reachable, for every target and every p . The strong form is also confirmed numerically (`theory.verify_support_bound_strong`): the maximum support of a non-full reachability mask equals $(R - 3)/2$ exactly at every prime residual checked. Lemma 1.10 and its dynamic-program verifications stand as independent machine confirmations of special cases (and remain the form carried by the Lean development).

Proof of Theorem 1.12, higher exponents. Write $n = (p + 3)/4$; the relevant forms are n , $n + 1$, $n + 2 = (p + 11)/4$, $n + 3 = (p + 15)/4$, $n + 4 = (p + 19)/4$, and $4n - 3 = p$. Joint failure of $\{3, 7, 11\}$ implies, by Theorems 1.2, 1.4, 1.5: n free of primes $\equiv 2 \pmod{3}$; $n + 1$ free of non-residues mod 7; and, splitting by the two cases of Theorem 1.5: *branch (a)* — $n + 2$ free of the five non-residue classes mod 11 (density $\frac{1}{2}$); *branch (b)* — $n + 2$ free of the six classes $\{4, 5, 7, 8, 9, 10\}$ mod 11 (density $\frac{3}{5}$). Each branch is a four-form sieve problem of dimension $1 + \frac{1}{2} + \frac{1}{2} + \frac{1}{2} = \frac{5}{2}$, resp. $1 + \frac{1}{2} + \frac{1}{2} + \frac{3}{5} = \frac{13}{5}$. Since $\frac{13}{5} > \frac{5}{2}$, the budget branch contributes $O(x/(\log x)^{13/5}) = o(x/(\log x)^{5/2})$: the bound is the maximum of the branch exponents, attained by the character branch (a). The Fundamental Lemma applied to each branch gives $O(x/(\log x)^{5/2})$ in total.

Adjoining $R = 15$: by Theorem 1.6, failure at 15 is the single condition “ $n + 3$ free of the four classes $\{7, 11, 13, 14\} \pmod{15}$ ” — one branch, density $\frac{4}{8} = \frac{1}{2}$ over the $\varphi(15) = 8$ unit classes — raising each combined branch to dimension ≥ 3 : joint failure of $\{3, 7, 11, 15\}$ is $O(x/(\log x)^3)$.

For $\{3, 7, 11, 15, 19\}$, Theorem 1.11 (equivalently, at this residual, Lemma 1.10) splits failure at 19 into finitely many maximal-support branches, each forbidding an explicit set of at least 9 of the 18 unit classes mod 19 on $n + 4$ — dimension $\geq \frac{1}{2}$. Crossing with the branches at 11 leaves finitely many combined branches, each a six-form problem of dimension $\geq \frac{7}{2}$ (those involving the budget case at 11 have dimension $\geq \frac{31}{10} + \frac{1}{2} > \frac{7}{2}$ and are again lower-order); the forbidden residues of the six forms are pairwise distinct for $q > 19$, smaller primes being absorbed into the congruence class of n . Summing the finitely many branches gives $O(x/(\log x)^{7/2})$.

The general case is identical: for a finite admissible list $\mathcal{P}$ (prime, or composite with $\omega(R) \leq 2$), the forms $n + (R - 3)/4$, $R \in \mathcal{P}$, are distinct integer shifts of n ; Theorems 1.2, 1.4, 1.5, 1.6 handle $R = 3, 7, 11, 15$ and Theorem 1.11 handles every other residual (part (ii) for primes, part (iii) for composites), each contributing dimension δ_R through finitely many maximal-support branches — $\delta_R = \frac{1}{2}$ for primes and for composites ≤ 107 , $\delta_R = \frac{1}{2} - 1/\varphi(R)$ for larger composites (the two tiers of part (iii)) — plus 1 for primality: total dimension $\geq 1 + \kappa(\mathcal{P})$, with forbidden residues pairwise distinct for $q > \max \mathcal{P}$ and smaller primes absorbed into the congruence class of n . For the fifteen primes up to 107 this gives exponent $\frac{17}{2}$; for the full 27-element admissible list, $\frac{29}{2}$. $\square$

Proof of Proposition 1.13. $k = ap^2$ divides $m^2 = p^2a^2$ and $k \equiv -m$ if and only if $ap^2 \equiv -pa$, i.e. $p \equiv -1 \pmod{R}$; $k = a^2p$ divides m^2 and $a^2p \equiv -pa$ if and only if $a \equiv -1$, i.e. $p \equiv -4 \pmod{R}$ by (1.3). Integrality of b, c is automatic. Scanning all shapes $k = a^j p^i$ ($j, i \leq 2$) yields no further p -only condition. The failure of both families states that the two forms $(p+1)/2$ and $p+4$ (in n : $2n-1$ and $4n+1$) are free of primes $\equiv 3 \pmod{4}$ — two further sifting conditions of dimension $\frac{1}{2}$ each on forms distinct from all previous ones. $\square$

Empirically the two families alone cover 74% of hard primes below 10^9 — and none of the four primes of Section 4 that admit a unique working residual ≤ 107 .

Remark 5.3. By Theorem 1.11, every residual — prime or composite — adds $\frac{1}{2}$ to the exponent, so the chain's exponent is bounded only by the number of residuals used; the constants $O_{\mathcal{P}}$, however, grow with $|\mathcal{P}|$ (branch counts and the remainder term of the Fundamental Lemma). The chain is a concrete instantiation of Theorem 1.14 with explicit constants. One caveat frames its scope, and it is fundamental: no fixed finite list can be pushed to *emptiness* by upper-bound sieving: the count bound $x/(\log x)^\kappa$ diverges for every fixed κ , and extrapolating the calibration of Section 4 to any fixed list yields a positive density constant, i.e. infinitely many expected exceptional primes with the record minimal residual growing slowly without bound. The correctly posed remaining problem is a slowly growing bound $R_{\min}(p) \ll f(p)$, which requires existence (lower-bound) technology beyond upper-bound sieves.

Proof of Theorem 1.14. Fix $A > 0$. Choose $B = B(A)$ so large that

$$\sum_{\substack{3 \leq R \leq B \\ R \equiv 3 \pmod{4}}} \frac{1}{\varphi(R)} \geq \sum_{\substack{3 \leq R \leq B \\ R \equiv 3 \pmod{4}}} \frac{1}{R} > A; \quad (5.1)$$

this is possible since the reciprocal sum over the progression $R \equiv 3 \pmod{4}$ diverges, and $\varphi(R) \leq R$. Let $S(A) = \{R \equiv 3 \pmod{4} : 3 \leq R \leq B\}$ and $M = 840 \cdot \prod_{R \in S(A)} R$.

Write $n = (p+3)/4$ as before, so $(p+R)/4 = n + (R-3)/4$ for each $R \in S(A)$: a family of $|S(A)|$ integer linear forms in n with pairwise distinct shifts, together with the form $4n-3 = p$. Set $z = x^{1/10}$ and discard the $O(z)$ primes $p \leq z$. Fix a residue class $n \equiv n_0 \pmod{M}$ compatible with p hard; there are finitely many, and it suffices to bound each. Within such a class, for every $R \in S(A)$ the quantities $t \equiv -4^{-1}p^2$ and $p^{-1} \pmod{R}$ are constant, so the class set $S_R(p) = \{t, tp^{-1}, tp^{-2}\} \pmod{R}$ of Proposition 2.2 is a fixed nonempty set of $t_R = t_R(n_0) \in \{1, 2, 3\}$ classes. By the contrapositive of Proposition 2.2, if every residual in $S(A)$ fails at p , then for each $R \in S(A)$ the form $n + (R-3)/4$ has *no* prime factor in $S_R(p)$; and $4n-3$, being prime and greater than z , has no prime factor $\leq z$.

Sift the class $\{n \leq (x+3)/4 : n \equiv n_0 \pmod{M}\}$ by the primes $q \leq z$, $q \nmid M$. A prime q forbids: the class $n \equiv 3 \cdot 4^{-1} \pmod{q}$ (from the primality of $4n-3$), and, for each $R \in S(A)$ with $q \bmod R \in S_R(p)$, the single class $n \equiv -\frac{R-3}{4} \pmod{q}$ (a prime factor q of the form $n + (R-3)/4$ in a guaranteed-success class cannot occur). Thus

$$\omega(q) = 1 + \#\{R \in S(A) : q \bmod R \in S_R(p)\},$$

which is bounded by $|S(A)| + 1$; the forbidden residues of distinct forms are pairwise distinct for $q > B$, since the shifts $(R-3)/4$ differ by less than B and $q \nmid M$. For each R , the density of primes q with $q \bmod R \in S_R(p)$ is $t_R/\varphi(R)$ by Dirichlet's theorem, so ω has average

$$\kappa = 1 + \sum_{R \in S(A)} \frac{t_R}{\varphi(R)} \geq 1 + \sum_{R \in S(A)} \frac{1}{\varphi(R)} > 1 + A,$$

in the Halberstam–Richert sense, by (5.1). The Fundamental Lemma [8, Thm. 2.5] in dimension κ then bounds the sifted count by $O_A\left(x \prod_{q \leq z} (1 - \omega(q)/q)\right) = O_A(x(\log x)^{-\kappa}) = O_A\left(x/(\log x)^{1+A}\right)$, and summing over the finitely many classes $n_0 \bmod M$ preserves the bound. $\square$

Remark 5.4. Theorem 1.14 is the strongest statement the pure sieve can give: density $O((\log x)^{-A})$ for every A , but not emptiness. The conjecture for hard primes would follow from a *finite covering hypothesis*: some finite S admits, for every hard prime, a residual $R \in S$ with (1.2) (the converse is not automatic — the conjecture bounds $R_{\min}(p)$ by $2p$, not uniformly). The data to 10^{12} is covered by $S = \{3, 7, 11, \dots, 107\}$ — yet both the calibration (Remark 5.3) and the conditional construction of Section 6 predict that any fixed list eventually fails.

5.2 The uniform chain

Theorem 1.14 fixes the list $S(A)$ once A is fixed; letting the list *grow with* x is purely a question of uniformity in the constants of Theorem 1.12 — the branch count per residual, the class modulus, and the sieve constants in growing dimension. Auditing those constants shows the growth is affordable for lists of length a small multiple of $\log \log x$, which yields the strongest unconditional statement of this paper.

Theorem 5.5 (uniform chain). *There are absolute constants $\varepsilon, c > 0$ and an (ineffective) x_0 such that for all $x \geq x_0$,*

$$\#\{\text{hard primes } p \leq x : R_{\min}(p) > \varepsilon \log \log x\} \leq x \exp(-c(\log \log x)^2).$$

In particular, all but $O(xe^{-c(\log \log x)^2})$ hard primes $p \leq x$ possess an explicit, divisor-exhibited residual certificate at some admissible $R \leq \varepsilon \log \log x$ — unconditionally.

Proof. Write $L = \log \log x$, $B = \varepsilon L$, and let $\mathcal{R}$ be the admissible residuals $3 \leq R \leq B$ ($|\mathcal{R}| \sim B/4$). Discard the $O(\sqrt{x})$ primes $p \leq \sqrt{x}$. A hard prime with $R_{\min}(p) > B$ fails every $R \in \mathcal{R}$; we bound the count of such p by a single sieve application per *branch*, as follows.

Dimensions. Each $R \in \mathcal{R}$ contributes a sifting dimension δ_R : the exact criteria at $R = 3, 7, 11, 15$ and Theorem 1.11(ii) at primes give $\delta_R = \frac{1}{2}$; composites ≤ 107 give $\frac{1}{2}$ by the sharpened tier of Theorem 1.11(iii); and every remaining composite — with *no* restriction on $\omega(R)$ — gives

$\delta_R \geq \frac{1}{2} - (2^{\omega(R)-1} + 1)/\varphi(R) \geq 0.44$ beyond an absolute bound R_1 , by the general- t clause of Theorem 1.11(iii). Hence $\kappa := 1 + \sum_{R \in \mathcal{R}} \delta_R \geq 1 + 0.44(|\mathcal{R}| - R_1) \geq 1 + B/10$ for B large.

Branches. For each R , failure confines the support of $(p+R)/4$ to one of at most $2^{\varphi(R)}$ failing supports S_R (crudely, all subsets); fixing $(S_R)_{R \in \mathcal{R}}$ and the class n_0 of $n = (p+3)/4$ modulo $\widetilde{M} = \text{lcm}(840, \prod_{\ell \leq B} \ell) = e^{O(B)}$ costs a factor $\prod_R 2^{\varphi(R)} \cdot e^{O(B)} = e^{O(B^2)}$ branches.

Sieve. Within a branch, sift the integers $n \leq (x+3)/4$, $n \equiv n_0 \pmod{\widetilde{M}}$, by the primes $\ell \in (B, z]$, $z = x^{1/s}$: the forbidden classes mod ℓ are $n \equiv 3 \cdot 4^{-1}$ (so that $p = 4n - 3$ keeps no factor $\leq z$) and, for each R with $\ell \bmod R$ outside S_R , $n \equiv -(R-3)/4$ — pairwise distinct for $\ell > B$, so $\omega(\ell) \leq |\mathcal{R}| + 1 < \ell/2$. The density condition holds with the dimension κ above and (Ω) -constant $O(B \log B)$, by Mertens' theorem in the progressions modulo $R \leq B$ (Siegel–Walfisz range, whence the ineffective x_0). Apply Brun's sieve of order $r \asymp \kappa$ ([8, Ch. 2]; for constants uniform in the dimension, [6, Thm. 6.9]) at $s = 10r$, so that $z^r = x^{1/10}$: since the sifted sequence consists of *integers* in a progression — no primality is detected anywhere — the remainder is $\sum_{d \leq z^r} (B+2)^{\omega(d)} \ll x^{1/10} e^{O(BL)} = x^{1/10+o(1)}$ for a suitable absolute choice of the implied constants, and the main term is

$$\ll \frac{x}{\widetilde{M}} \prod_{B < \ell \leq z} \left(1 - \frac{\omega(\ell)}{\ell}\right) \leq x \exp(-\kappa(L - O(\log B)) + O(B \log B)).$$

Accounting. Multiplying by the branch count,

$$\begin{aligned} \#\{p \leq x : R_{\min}(p) > B\} &\leq e^{O(B^2)} x \exp(-\kappa L + O(B \log B) + O(\kappa \log B)) \\ &\leq x \exp(-BL/10 + O(B^2)), \end{aligned}$$

and with $B = \varepsilon L$ the exponent is $-\varepsilon L^2/10 + O(\varepsilon^2 L^2) \leq -\varepsilon L^2/20$ for ε below an absolute threshold. $\square$

Remark 5.6 (the constant, and the range). The constant c in Theorem 5.5 is not innocuous. It inherits the branch factor $e^{\Theta(B^2)}$ against the sieve gain $e^{-BL/10}$, so the ε below which the accounting closes is small, and $c \asymp \varepsilon$. The bound beats the trivial count $\asymp x/(32 \log x)$ of hard primes only once $c(\log \log x)^2 - \log \log x > \log 32$, which for small c places the crossover at values of x far beyond any computational range: at $c = 10^{-2}$ one needs $\log \log x \gtrsim 350$. The theorem is therefore asymptotic in the strict sense — nontrivial for all large x , ineffectively so (Siegel–Walfisz), and of no numerical content in any accessible range. Its interest is structural, and the numerical picture at accessible heights is supplied instead by Section 4.

Remark 5.7 (the ceiling, and the one lever). Theorem 5.5 sits at the boundary set by its *branch factor*: $e^{\Theta(B^2)}$ from the crude all-subsets bound on failing supports, against the sieve gain $e^{-BL/10}$, forces $B \ll L$ and the ceiling $\exp(-cL^2)$. A closer audit (Section 5.3) shows the branch factor is the *only* binding wall: the class modulus $\widetilde{M} = e^{O(B)}$ never multiplies the main term (the classes partition the integers), and the Mertens/Siegel–Walfisz inputs remain valid for B up to any fixed power of $\log x$ at the already-paid price of ineffectivity. Classifying the failing supports into polynomially many structured branches — the subject of Section 5.3 — would therefore raise Theorem 5.5 to an exceptional set $x \exp(-(\log x)^{c_0})$ with $R_{\min} \leq (\log x)^{c_0}$ (Theorem 5.12), with the family ceiling strictly below exponent $\frac{1}{2}$ (see the discussion there). Vaughan's bound $x \exp(-c(\log x)^{2/3})$ [29] remains stronger as a pure density statement at every tier of this family; what Theorem 5.5 adds is structural: an explicit certificate at a small residual for almost every hard prime, a statement mean-value methods do not address. The calibrations $J = (\log x)^\theta$ of Theorem 5.22 remain out of reach of every unconditional variant considered here.

5.3 The branch classification

Fix a prime residual R , $d = R - 1$, and work in the discrete-log model of Theorem 1.11(i): a support $S \subseteq \mathbb{Z}/d \setminus \{0\}$ at exponent budget 2 per class has reach $M(S) = \sum_{v \in S} \{0, v, 2v\}$, and S fails for a target t if $t \notin M(S)$ (the p -budget is absorbed as one more summand). By Lemma 5.1 every failing configuration lives on a failing support, and every failing support extends to a maximal one (S failing, $S \cup \{w\}$ succeeding for all w). Two facts calibrate the problem. First, maximality forces no size concentration: maximal failing supports exist at every size scale from $\Theta(\log d)$ (complete-partition constructions realizing $M(S) = \mathbb{Z}/d \setminus \{t\}$, a *singleton miss*) up to the extremal $d/2 - 1$. Second, *enumerating* them is hopeless:

Theorem 5.8 (branch count blow-up). *There is $c > 0$ such that for every large even d , the number of maximal failing supports for the target $t = -1$ in $\mathbb{Z}/d$ is at least $2^{c\sqrt{d}}$: the supports $[1, r] \cup T' \cup \{X\}$, $T' \subseteq [r + 1, 2r]$, $|T'| = \lceil r/2 \rceil$, $r = \lceil \sqrt{d/4} \rceil$, with X completing the sum to $d/2 - 1$, are pairwise distinct complete partitions (for d large, $X > 2r$, so X collides with no other part and T' is recoverable from the support; the chain condition at the part X reduces to $d/2 - 2 \leq 3r(r + 1)$, which holds since $r^2 \geq d/4$), and each has reach exactly $\mathbb{Z}/d \setminus \{-1\}$, hence is maximal failing.*

The exponential-in- $\sqrt{d}$ order is not a surprise: the objects constructed are complete partitions with distinct parts, whose total count is $\exp(\Theta(\sqrt{n}))$ classically (Andrews–Beck–Hopkins [11]). What Theorem 5.8 records is the version adapted to $\mathbb{Z}/d$ and to maximality, because that is what rules out enumeration. A useful classification must therefore produce *containers* — polynomially many structured sets covering all failing supports. The kernel (periodic) case is under complete control, by the standard stabilizer/quotient reduction that drives the critical-number literature (see [12, Ch. 6]); we state it in the form and with the size bookkeeping we need, and claim no novelty in the mechanism.

Theorem 5.9 (kernel branch; lossless reduction). *Let S fail for t and $H = \text{Stab}(M(S)) \neq \{0\}$, and write $\varphi_H: \mathbb{Z}/d \rightarrow \mathbb{Z}/\bar{d}$. Then $t \notin M(S) + H$, the projection $\bar{T} = \varphi_H(S \setminus H)$ is a failing support for $\bar{t} = \varphi_H(t) \neq 0$ at the same budget, and $S \subseteq (H \setminus \{0\}) \cup \varphi_H^{-1}(\bar{T})$; if S is maximal then moreover $H \setminus \{0\} \subseteq S$. Iterating: every failing support of $\mathbb{Z}/d$ is contained in $\psi^{-1}(\{0\} \cup T') \setminus \{0\}$ for a composed quotient $\psi: \mathbb{Z}/d \rightarrow \mathbb{Z}/d'$ and an aperiodic maximal failing support T' downstairs, and this container has at most $\lfloor d/2 \rfloor - 1$ elements. (The support bound extends to odd moduli, as needed for the quotients: for odd $\bar{d}$ no summand is degenerate, so in the aperiodic case Kneser's bound gives $|T| \leq (\bar{d} - 3)/2$; for nontrivial stabilizer of order h the count is $\leq h + \bar{d}/h - 3 \leq \bar{d}/3 \leq (\bar{d} - 3)/2$ for $\bar{d} \geq 9$, smaller odd $\bar{d}$ having no proper nontrivial subgroups.)*

Proof. $M + H = M$ gives $H \subseteq M$ and $t \notin M + H$ (else $t \in M$), so $\bar{t} \neq 0$. The projection $\varphi_H(M(S)) = \sum_v \{0, \bar{v}, 2\bar{v}\}$ contains $M(\bar{T})$ computed downstairs at budget 2 (collisions only enlarge budgets; Lemma 5.1), and H -saturation of $M(S)$ turns $t \notin M(S)$ into $\bar{t} \notin M(\bar{T})$: $\bar{T}$ fails. The containment is definitional. For maximality: any $w \in H \setminus (S \cup \{0\})$ has $A_w \subseteq H$, so $M(S \cup \{w\}) = M(S) + A_w = M(S)$ still fails — w must already lie in S . Iterate on the quotient (strong induction on d , extending $\bar{T}$ to a maximal support downstairs first); the size bookkeeping per step is $(h - 1) + h|T'| \leq h - 1 + h\lfloor (\bar{d} - 2)/2 \rfloor \leq d/2 - 1$, exact at even $\bar{d}$, with strict inequality at odd $\bar{d}$. $\square$

The open core is the aperiodic case. All the evidence points to one container shape — a coset-progression *window*: for $K \leq \mathbb{Z}/d$, a unit direction and a circular window W of quotient

classes, set $C(K, W) = \varphi_K^{-1}(\{0\} \cup W) \setminus \{0\}$. The family $\{C(K, W) : |C(K, W)| \leq \lfloor d/2 \rfloor - 1\}$ has fewer than d^3 members (a container is determined by the kernel size $\kappa \mid d$, a unit direction, a window start, and a window length $\leq \bar{d}/2$: at most $\sum_{\kappa \mid d} \bar{d}^3/2 \leq \zeta(3)d^3/2$ choices).

Conjecture 5.10 (aperiodic branch classification). *Every maximal failing support S with $\text{Stab}(M(S)) = \{0\}$ for a target t in $\mathbb{Z}/d$ satisfies $S \subseteq C(K, W)$ for some window container with $|C(K, W)| \leq \lfloor d/2 \rfloor - 1$.*

The evidence: exhaustive verification for *all* moduli $d \leq 60$ and all targets (24.25 million orbit-weighted maximal supports; zero violations; the bound $\lfloor d/2 \rfloor - 1$ is attained, hence optimal); exhaustive verification at $d = 72$, the smallest *non-Hajós* cyclic order — the one place aperiodic factorization pathologies could first appear — over all 11 target orbits and 14.99 million maximal supports, with extensive stratified sampling (including supports seeded from explicit aperiodic de Bruijn factorizations of $\mathbb{Z}/72$) at every non-Hajós order up to 240; independent verification in the residual model at $R = 19, 23, 31, 43, 47, 59$ (495,782 maximal supports at $R = 59$) and at the composite moduli $R = 15, 35, 39$ (entry point `erdos_straus.branch_enum`; archived); a proved subfamily (the signed complete partitions, windows with room to spare); and a rank-2 witness ($S = \{11, 17, 19, 29\} \subseteq \mathbb{Z}/30$, a singleton miss contained in no short one-dimensional window but in a $K = \langle 10 \rangle$ window container of size 14) showing the subgroup component of the shape is necessary. Combining Theorem 5.9 with Conjecture 5.10 in the quotients (windows pull back to windows along composed quotients):

Theorem 5.11 (container form; conditional on Conjecture 5.10). *Every failing support for (d, t) is contained in one of fewer than d^3 explicit window containers of size $\leq \lfloor d/2 \rfloor - 1$; hence failure of residual R confines the factor classes of $(p + R)/4$ to one of $O(R^3)$ branches, each forbidding at least half the nonzero classes — sieve dimension $\frac{1}{2}$ per residual, with the branch count reduced from $2^{\varphi(R)}$ to $O(R^3)$.*

Theorem 5.12 (conditional uniform chain, polylog tier). *Assume Conjecture 5.10 (for all moduli). Then there is an absolute $c_0 > 0$ such that all but $O(x \exp(-(\log x)^{c_0}))$ hard primes $p \leq x$ satisfy $R_{\min}(p) \leq (\log x)^{c_0}$.*

Proof sketch. Rerun the proof of Theorem 5.5 with list length $B = (\log x)^\theta$. The branch factor over the list drops from $e^{\Theta(B^2)}$ to $\prod_{R \leq B} O(R^3) = e^{(3+o(1)) \cdot \frac{B}{4} \log B}$; the class conditioning never multiplies the main term (the classes partition the sifted integers, so the per-class main terms resum); and Siegel–Walfisz covers the Mertens-in-progressions inputs for moduli up to any fixed power of $\log x$. The per-residual balance is then a sieve gain $\delta(1 - \theta)L$ against a branch-and-slop cost $(3 + c_M)\theta L$, where $L = \log \log x$, $\delta = \frac{1}{2}$, and c_M collects the per-residual slop — the Mertens-in-progressions constants, the discarded initial segment $\ell \leq B$ of sieving primes, and the (Ω) -condition constant, together of order 1. This is admissible for every $\theta < \delta/(3 + c_M + \delta)$ — with the constants proved here, any θ below roughly $1/9$. The family ceiling sits strictly below $\frac{1}{2}$: no container scheme of this shape has fewer than $\sim d$ members (distinct progression directions already require distinct containers), so the count exponent is at least 1 and the balance caps θ at $\delta/(1 + c_M + \delta) < \frac{1}{3}$ — consistent with, and sharper than, the $\exp(-(\log x)^{1/2+o(1)})$ outer ceiling of the wider method family, itself short of Vaughan’s $2/3$. A quasi-polynomial classification ($d^{O(\log d)}$ containers) would still give exceptional sets $x \exp(-\exp(c\sqrt{\log \log x}))$, past every $\exp(-C(\log \log x)^A)$. $\square$

Conjecture 5.10 is thus a single well-posed question of additive combinatorics — a container theorem of Kneser–Freiman type for bounded-exponent subset sums — standing between Theorem 5.5 and its polylog tier. Where its difficulty lives can now be pinned exactly. Near-critical Kneser analysis is *vacuous* here (aperiodic failing supports with $|S| \geq 3$ are never Kneser-tight — the governing mechanism is exact covering by complete sequences, not small doubling); the singleton-miss supports satisfy the symmetry $2 \sum_{v \in S} v = 2t$, which pins their windows (and at *odd* moduli singleton misses do not exist at all: the miss set always has even size, so the floor family there is $|F| = 2$, auto-maximal and window-shaped). The one mechanism previously feared — chaining of *punctured* (cofinite) intermediate reaches, where direction information degenerates and non-Hajós factorization pathologies could hide — is provably empty:

Proposition 5.13 (β -vacuity). *Let S fail for t in $\mathbb{Z}/d$, let $v \in S$, and suppose the deletion step at v is Kneser-critical ($|M(S)| = |M(S \setminus \{v\})| + 2$) of punctured type: the two new elements are isolated punctures of the complement. Then $F - v = F$ for the miss set $F = \mathbb{Z}/d \setminus M(S)$, so $v \in \text{Stab}(M(S))$. In particular, for an aperiodic maximal failing support every critical deletion step is of window (α) type.*

Proof. Write $A = M(S \setminus \{v\})$ and $\mathbb{Z}/d \setminus A = F \sqcup \{x, y\}$ with x, y the punctures. For $f \in F$: since $f \notin M(S) = A \cup (A + v) \cup (A + 2v)$, both $f - v$ and $f - 2v$ avoid A , so $f - v \in F \cup \{x, y\}$. If $f - v = x$, then $x - v = f - 2v \notin A$, contradicting the isolation of x ($x - v \in A$; the $v = d/2$ two-cycle case is the same contradiction). Likewise for y . So $F - v \subseteq F$, hence $F - v = F$ by finiteness. $\square$

What remains open in Conjecture 5.10 is therefore only the control of *non-critical* deletion steps — slack-consuming extensions, a question about exact-covering efficiency that is indifferent to the Hajós/non-Hajós divide (consistent with the exhaustive verification at $d = 72$ above). And the payoff is robust to the conjecture’s scope: by a Chen-type P_3 count there are $\gg B/(\log B)^2$ admissible residuals $R \leq B$ whose modulus $d = R - 1$ is a Hajós order (equivalently $(R - 1)/2$ is a prime power or has at most three prime factors), which is dense enough that assuming Conjecture 5.10 *only for Hajós orders* — where the group theory is tame — yields Theorem 5.12 with the same exponent up to $o(1)$.

5.4 A conditional program past the ceiling

Theorem 1.11 closes the fixed-list axis and Theorem 5.5 pushes the growing-list axis to its unconditional ceiling $\exp(-c(\log \log x)^2)$ (Remark 5.7); no list-based bound can reach emptiness (Remark 5.3). We record here, in compressed form, the one program we know that opens past the ceiling: a *p-adapted* family of residuals, selected through quadratic reciprocity, along which the all-residue failure mode is impossible by construction. Its measurements are archived (Appendix A); the fuller development, including the per-rung census data, is in the accompanying repository.

Lemma 5.14 (composite reciprocity). *Let $R \equiv 3 \pmod{4}$ be any admissible residual and q an odd prime with $q \mid (p + R)/4$ and $\gcd(q, R) = 1$. Then the Jacobi symbol satisfies $\left(\frac{q}{R}\right) = \left(\frac{p}{q}\right)$.*

Proof. $q \mid p + R$ gives $R \equiv -p \pmod{q}$; Jacobi reciprocity with $R \equiv 3 \pmod{4}$ gives $\left(\frac{q}{R}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{R}{q}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{-1}{q}\right) \left(\frac{p}{q}\right) = \left(\frac{p}{q}\right)$ — the proof of Theorem 1.9 verbatim, which never used primality of R . $\square$

Lemma 5.15 (Jacobi necessity). *For every admissible R : if no prime factor of $a = (p + R)/4$ has Jacobi symbol $(\frac{\cdot}{R}) = -1$, then (1.2) fails at R .*

Proof. All factors of a Jacobi residues gives $(\frac{a}{R}) = +1$; by consistency (1.3), $(\frac{p}{R}) = (\frac{4a}{R}) = +1$ as well, so every prime factor of $m = pa$ has symbol $+1$, hence every divisor $k \mid m^2$ has $(\frac{k}{R}) = +1$, while the target has $(\frac{-m}{R}) = (\frac{-1}{R}) = -1$. $\square$

Call R *pure* (at hard primes) if the converse also holds: a Jacobi non-residue factor present implies success. Theorems 1.2, 1.4 and 1.6 say exactly that $R = 3, 7, 15$ are pure; the census below shows they are the *only* pure residuals up to 107, with the conditional failure rate at the others ranging from 1.3% ($R = 23$) to 73% ($R = 67$).

Now run Lemma 5.14 backwards. For a hard prime p let $q = q(p)$ be its least Legendre non-residue (hard primes are residues at 3, 5, 7 by Corollary 2.3, so $q \geq 11$; Burgess bounds $q \ll p^{1/(4\sqrt{e})+\varepsilon}$ unconditionally [2] and GRH gives $q \ll (\log p)^2$ [1]; the observed maximum below 10^{11} is 83), and let the *selected residual* be $R_0 = (-p) \bmod 4q$, so that $R_0 \equiv 3 \pmod{4}$ automatically and $q \mid (p + R_0)/4$. By Lemma 5.14, $(\frac{q}{R_0}) = (\frac{p}{q}) = -1$: the all-residue failure mode of Lemma 5.15 is impossible at R_0 *by construction*. The q -ladder is the progression $R_j = R_0 + 4qj$, along which the same holds at every rung.

Computed over every one of the 1,587,581 hard primes below 10^9 , with mask-free window samples of 10,000 primes per half-decade beyond, and the complete 10^{11} deep tail (Appendix A): the selected residual R_0 alone certifies 94.78% of hard primes below 10^9 and 95.0–95.8% in every sampled window up to 10^{11} , with no drift; the ladder capped at $R_j \leq 400$ certifies 100% below 10^9 (two primes requiring the ladder of the second-least non-residue) and 100% in every window. On the adversarial 10^{11} deep tail ($R_{\min} \geq 43$, the 0.016% hardest), the first rung succeeds for 75.1% and the capped ladder for 99.99%, with 100% at $R \leq 435$. Among all 82,808 first-rung failures below 10^9 , 97.5% are *true budget* failures — the subgroup generated by the factor classes contains the target and only the bounded exponents miss it — 29% miss exactly one unit class, and none has more than 5 distinct non-identity factor classes, against the $(R - 3)/2$ that Theorem 1.11 would allow.

Theorem 5.16 (ladder chain). *Fix a prime $q \geq 11$ and $J \geq 1$. The number of hard primes $p \leq x$ with $(\frac{p}{q}) = -1$ failing every rung R_j , $0 \leq j < J$, of their q -ladder is $O_{q,J}(x/(\log x)^{1+\kappa_J})$, where $\kappa_J = \sum \delta_{R_j}$ over the rungs with at most two prime factors, with $\delta_R = \frac{1}{2}$ for R prime or composite ≤ 107 and $\delta_R = \frac{1}{2} - 1/\varphi(R)$ for larger composites.*

Proof. The shifted forms $(p + R_j)/4 = n + (R_0 - 3)/4 + qj$ are distinct integer shifts of $n = (p + 3)/4$, and Theorem 1.11 — part (ii) at prime rungs, part (iii) at composite rungs with $\omega \leq 2$ — makes failure at each usable rung a sifting condition of dimension $\geq \delta_{R_j}$, exactly as in the proof of Theorem 1.12, which applies verbatim to this p -adapted list. $\square$

Theorem 5.17 (first rung). *Fix a prime $q \geq 11$. Among hard primes $p \leq x$ with least non-residue $q(p) = q$, the proportion failing the selected residual R_0 is $\ll_q (\log x)^{-\eta_q} \rightarrow 0$, where $\eta_q = \min 1/\varphi(R_0) > 1/(4q)$, the minimum over the finitely many selected residuals $R_0 < 4q$ that arise for q .*

Proof. The condition $q(p) = q$ is a union of residue classes modulo $M = 840 \prod_{11 \leq \ell \leq q} \ell$; fix one, which fixes R_0 . By Proposition 2.2, any prime factor u of $(p + R_0)/4$ in the universal class $-4^{-1} \bmod R_0$ certifies success via $k = up^2$, so a failing p has $(p + R_0)/4$ free of an entire congruence class of primes: a sifting condition of dimension $1 + 1/\varphi(R_0)$ (with primality),

and the Fundamental Lemma bounds the failing count by $\ll x/(\log x)^{1+1/\varphi(R_0)}$ against class size $\asymp x/\log x$ (Siegel–Walfisz). Summing over the classes, the proportion is governed by the worst class, giving η_q ; and $\varphi(R_0) < R_0 < 4q$ gives $\eta_q > 1/(4q)$. The bound is uniform for $q \ll \log \log x$. $\square$

Write $E_j(x)$ for the set of hard primes in a fixed class that fail every rung through j , and $\tilde{E}_j(x)$ for the larger set whose shifted values $(p + R_i)/4$, $i \leq j$, all avoid prime factors in their universal classes. The same Fundamental Lemma argument bounds $\#\tilde{E}_J \ll_{q,J} x/(\log x)^{1+\kappa}$ with $\kappa = \sum_{j \leq J} 1/\varphi(R_j)$, and for κ below an absolute constant a beta-sieve in dimension $\kappa < 1$ at level $x^{1/2-\varepsilon}$ (Bombieri–Vinogradov) gives the matching lower bound: the *avoidance* ladder provably decays by $(\log x)^{-1/\varphi(R_{j+1})}$ at every rung. What is not provable by present methods is that the true failure sets E_j track their avoidance sets:

Hypothesis 5.18 (P; proportionality). *For each fixed q , class, and J there are $c_P > 0$ and x_0 , with c_P bounded below uniformly in the rung $j \leq J$ (uniformity in q is part of the hypothesis), such that $\#E_j(x) \geq c_P \#\tilde{E}_j(x)$ for all $j \leq J$ and $x \geq x_0$.*

Measured: $c_P \approx 0.097$ – 0.12 at rung 0 and 0.13 – 0.16 at rung 1, flat across two decades to 10^{11} (Appendix A). Since $E_{j+1} \subseteq \tilde{E}_{j+1}$, Hypothesis 5.18 converts the avoidance decay into decay of the true failure counts at every *fixed* rung, though not with uniformity in j and q .

The one place where the true failure count *is* bounded below unconditionally is at the start of the chain, and the reason is that the character-failure condition is a *norm-form* condition in disguise. For $r \equiv 3 \pmod{4}$, quadratic reciprocity gives $(\frac{-r}{\ell}) = (\frac{\ell}{r})$ for odd primes $\ell \nmid r$, so

$$\begin{aligned} \text{every prime factor } \ell \text{ of } a \text{ has } \left(\frac{\ell}{r}\right) = +1 &\iff -r \text{ is a square mod } 4a \\ &\iff a \text{ is primitively represented by} \\ &\quad \text{a binary quadratic form of discriminant } -r, \end{aligned} \tag{5.2}$$

the last step by the classical correspondence between representations and square roots of the discriminant (Cox [16, Lem. 2.5]); for r composite, $(\frac{\cdot}{r})$ is the Jacobi symbol throughout. When r is prime the discriminant $-r$ has a single genus and the whole form class group participates; at $r = 15$ there are two genera and two classes, and the counting below is applied to each form separately and summed, which is all that is needed. The relevant counting theorem is Iwaniec’s half-dimensional sieve [18] in its *prime-indexed* form: the number of primes $p \leq N$ with $p = B\varphi(\xi, \eta) + A$ for a primitive form φ is $\asymp N/(\log N)^{3/2}$ [19]. This is not obstructed by parity — at dimension $\kappa = \frac{1}{2}$ the lower-bound sieve function satisfies $\beta(\frac{1}{2}) = 1$.

Theorem 5.19 (the first link is sharp). *For $R \in \{3, 7, 15\}$,*

$$\#\{\text{hard primes } p \leq x \text{ failing } R\} \asymp \frac{x}{(\log x)^{3/2}},$$

and the same holds with $\ll, \gg$ separately at $R = 11$. In particular the chain of Theorem 1.12 is sharp at its first link, and no exponent better than $3/2$ is available from any single residual: by Theorem 1.11 every failure-sufficient family has sieve dimension $\geq \frac{1}{2}$, so $3/2$ is the ceiling of the method, not an artifact of the argument.

Proof sketch. At these residuals the exact criteria (Theorems 1.2, 1.4, 1.6) make failure equivalent to an all-residue condition on the factors of $a = (p + R)/4$, which by (5.2) is primitive

representability by a form of discriminant $-r$ ($r = 3, 7, 15$ respectively). The upper bound is the Fundamental Lemma as in Theorem 1.12; the lower bound is [19] applied to each class of forms, with the congruence conditions (H1)–(H3) absorbed into A and B . At $R = 11$ the case-(b) budget failures of Theorem 1.5 are not of this shape, which is why only the one-sided bounds are asserted. $\square$

A lower bound of the same order survives the passage to p -adapted residuals, but only on a sparse family of classes, and the reason is worth recording because it is the obstruction in miniature. Fix a class $c \bmod M$ (M as in the proof of Theorem 5.17) and suppose R_0 has a prime factor $r_1 \equiv 3 \pmod{4}$ with $\left(\frac{\ell}{r_1}\right) = +1$ for *every* prime $\ell \leq q$ that the class forces to divide a . Then every p in the class whose shifted value $(p + R_0)/4$ has all prime factors $\left(\frac{\cdot}{r_1}\right) = +1$ fails the first rung, by Proposition 2.1 at r_1 , and (5.2) with [19] counts those primes: $\gg x/(\log x)^{3/2}$ on any such class. The upgrade from representability by *some* form of the discriminant to *primitive* representability is not automatic; it is the subject of Fuchs–Hsu–Rickards–Schindler–Stange [17], and the cases used here fall under their Theorem 1.1 when $r_1 \equiv 3 \pmod{8}$ with a odd, or $r_1 \equiv 7 \pmod{8}$ with a even.

Remark 5.20 (how sparse). The forcing condition is restrictive in a way that is easy to miss: the class fixes a modulo every prime $\leq q$, so a single forced factor ℓ with $\left(\frac{\ell}{r_1}\right) = -1$ empties the family. For instance $q = 23$, $R_0 = 35$, $r_1 = 7$ passes the $\ell = q$ test, but $3 \mid a$ for every hard prime (as $p \equiv 1 \pmod{3}$ and $35 \equiv 2 \pmod{3}$) while $\left(\frac{3}{7}\right) = -1$: the set is empty. Among the 39,391 hard primes below 2×10^7 , the classes passing the $\ell = q$ test alone account for 6.6% of first rungs (2,593), those with a nonempty family for 0.023% (9, in four distinct classes (q, R_0, r_1)), and 93.4% admit no candidate r_1 whatever. One structural reason dominates: whenever R_0 is prime — 65.2% of first rungs in this range — its only prime factor is $r_1 = R_0$, and Theorem 1.9 forces $\left(\frac{q}{R_0}\right) = \left(\frac{p}{q}\right) = -1$, so no admissible r_1 exists by construction. All figures are reproduced by `erdos_straus.burgess_scan` `admissibility` and archived. Consequently Hypothesis 5.18 is bracketed two-sidedly only on that sparse family at rung 0, and one-sidedly elsewhere; Theorem 5.19 is the unconditional lower-bound statement of real reach.

Hypothesis 5.21 (B; per-rung decay). *Fix a rung-count function $J(x) \rightarrow \infty$. There are $\delta > 0$ and x_0 such that for $x \geq x_0$, all $q \leq (\log x)^2$ and all $0 \leq j < J(x)$: $\#E_{j+1}(x) \leq (1 - \delta) \#E_j(x) + O(\sqrt{x})$.*

The $j = 0$ case is Theorem 5.17; every fixed rung follows from Hypothesis 5.18 (non-uniformly in j and q); the measured survival factors $1 - \delta$ are 0.05–0.14 per rung — 86–95% of the surviving failures resolve at each rung — with no drift in p .

Theorem 5.22 (conditional almost-all bound). *Assume Hypothesis 5.21. Then all but $O\left(x e^{-\delta J(x)} + x 2^{-\pi((\log x)^2)}\right)$ hard primes $p \leq x$ satisfy $R_{\min}(p) \leq 4q(p)(J(x) + 1) \leq (\log x)^{2+o(1)}J(x)$.*

Proof. Hard primes with $q(p) > (\log x)^2$ are residues at every prime in $(7, (\log x)^2]$, a half-class condition per modulus, and the large sieve bounds their count by the second term. For the rest, iterate Hypothesis 5.21 $J(x)$ times; a survivor is the only way to lack a certificate on the ladder. $\square$

Remark 5.23 (scope). With $J(x) = (\log \log x)^2$ the exceptional set would be $\ll x \exp(-\delta(\log \log x)^2)$ — but that calibration is delivered *unconditionally*, and with the stronger bound $R_{\min} \leq$

$\varepsilon \log \log x$, by Theorem 5.5. So Hypothesis 5.21’s real content begins past the ceiling of Remark 5.7: with $J(x) = (\log x)^\theta$, $\theta > 2/3$, Theorem 5.22 would pass Vaughan’s bound. It is not unconditional, it is not a density record, and it is not a path to the full conjecture by itself: emptiness is beyond upper-bound methods (Remark 5.3), and excluding the individual conspiratorial prime — a super-record dodging $J(p)$ rungs, of which the record prime of Section 4 is an empirical foreshadowing — is a simultaneous-multiplicative-structure statement of prime-tuple difficulty. Two computations pin the walls down. (i) Budget failure is not a character event: at $R = 11$ the configurations with class set $\{2, 7\}$ and multiplicities $(1, 1)$ versus $(4, 2)$ have the same class set and the same product mod 11 — hence identical values under every Dirichlet character mod 11 — yet the first fails the target and the second attains it, and 97.5% of measured first-rung failures are of this budget type. (ii) Weighted detection is equivalent to Hypothesis 5.18: by Proposition 2.2 the certificate detector is already maximally relaxed, so almost-prime weakenings change nothing, and the first moment of any Bonferroni/Selberg weighting over E_j is itself a failure count.

6 Open problems

1. **Finite covering.** Prove that some finite admissible set S covers all hard primes (Theorem 1.14 gives density $(\log x)^{-A}$ for every A ; emptiness is beyond the sieve).
2. **Beyond fixed lists.** The support bound is proved unconditionally for every residual, prime and composite (Theorem 1.11), Theorem 5.5 carries list growth to length $\varepsilon \log \log x$, and the branch classification that governs everything further is now formulated in Section 5.3: the kernel case is proved (Theorem 5.9), the count form refuted (Theorem 5.8), and what remains open is exactly Conjecture 5.10 — the aperiodic window-container statement, whose difficulty is now localized to the control of non-critical deletion steps (the punctured-chain mechanism is empty by Proposition 5.13); its resolution — for Hajós orders alone — gives Theorem 5.12. The fixed-list format itself cannot reach emptiness (Remark 5.3). The proof of Theorem 1.11 shows the extremal failing supports are *subgroup-trapped* (the stabilizer H is nontrivial and S concentrates on H up to at most $d/|H| - 2$ outliers); exploiting this structure *uniformly in R* — fewer, p -independent branches — is now a well-posed additive-combinatorics question, and the structurally right target remains a slowly growing bound on $R_{\min}(p)$.
3. **The constant 5.17.** The joint $\{3, 7\}$ -exceptional density is measured as $\approx 5.17/\log x$, stable to 1% across three decades. Conjecturally 5.17 is a Selberg–Delange-type constant: the count should satisfy $E(x) \sim Cx/(\log x)^2$ with $C = c_0 \prod_q \delta_q$, the local factors δ_q recording, for each prime q , the probability that q divides neither form in a forbidden class. Proving the asymptotic (not just the upper bound) requires a two-form Selberg–Delange or shifted-convolution analysis — open, and a natural test problem for such methods (see also Remark 1.15 for the density context).
4. **Structural reformulations and complexity.** Theorem 1.9 invites a graph formulation: on the bipartite graph joining primes q to residuals R by $q \mid (p+R)/4$, the edge characters $(\frac{q}{R})$ collapse to vertex signs $(\frac{p}{q})$, and joint character-failure is an all-plus vertex assignment — a percolation-style object whose structure could organize the growth law. Separately, the enumeration complexity is unresolved: the DP state count grows empirically by $\approx 2.8 \times$

per successive prime residual (271 at $R = 19$ to 1.95×10^7 at $R = 107$); is the equivalence-state complexity genuinely exponential in R ? (For the support-bound property itself the question is settled: Theorem 1.11 is a constant-size certificate. The open version concerns the full exact criteria.)

5. **The record and the covering hypothesis.** Theorem 1.9 explains the record prime quantitatively (see Corollary 2.3) and yields a conditional resolution of the covering question: under Dickson’s prime-tuple conjecture, for every B there are infinitely many hard primes failing every residual $\leq B$ — fix the forced smooth parts of the shifted values (character-neutral by Corollary 2.3(i)), make the cofactors prime in prescribed residue classes, and apply Proposition 2.1. So no fixed finite list suffices, conditionally, and the record grows without bound (at the pure-character level, like $\log x / \log \log x$). Unconditional versions appear out of reach: already “infinitely many hard p fail $R = 3$ ” contains a two-form prime-tuple statement. What remains open is an unconditional upper bound on $R_{\min}(p)$ — and this is not an incremental item: by Proposition 1.1, $R_{\min}(p) \leq 2p$ is *equivalent* to the conjecture at p , so an unconditional bound $R_{\min}(p) \leq f(p)$ for any f whatsoever is the full problem.

A Verification details

This appendix documents every computer-assisted claim of the paper (their proof status is stated in Remark 1.8): the finite enumerations that are *proof components* (Appendix A.1), the large-scale certificate computations behind Section 4 and how each dataset is verified (Appendix A.2), and the commands that reproduce all of it (Appendix A.3). All code is in the accompanying repository; every check is deterministic and uses exact integer arithmetic only.

A.1 Proof components

Each component below is an exhaustive enumeration of a finite space whose sufficiency is established in the text (Theorem 1.3 for the exact criteria; for Lemma 1.10, monotonicity of the reachability of the target in support and multiplicities, checked at minimal exponents). Each run completes in seconds on commodity hardware, except the largest support-bound check ($R = 107$: 331 s); the entry points named are in the module `erdos_straus.theory`.

Trusted computing base. The certificate data does not rest on the solver: every certificate claim is re-checked by evaluating the integer identity $4abc = p(bc + ca + ab)$ directly, so the trusted base for the data is a bignum multiply-and-compare (the host language’s built-in integers), not the search pipeline. The finite case analyses trust more — the enumeration code itself, roughly sixty lines per check. Four mitigations apply. First, the enumerated spaces are defined by statements proved in the text (Theorem 1.3, Lemma 5.1), and the code mirrors those definitions one-to-one. Second, independent implementations agree where they overlap: the subset dynamic program and the direct binomial enumeration verify identical claims at $R = 19, 23$, and the exact criteria are cross-checked against ground-truth divisor search on up to 1,587,581 primes with zero discrepancies. Third, every check is deterministic, uses exact integer arithmetic only (no floating point, randomness, or external libraries), and runs from the repository in seconds (the largest, the $R = 107$ dynamic program, in 331 s) — an independent re-implementation is an afternoon’s work, and we would welcome one.

Fourth, much of the development is formalized in the accompanying repository (`lean/`) in Lean 4 [3] with its mathematical library `mathlib` [21]. For readers unfamiliar with proof assistants: Lean is a language in which theorem statements and proofs are written formally and every proof is checked by a small trusted *kernel*, so a theorem accepted by the kernel is correct relative to its formal statement, independently of who — or what — wrote the proof. Two Lean-specific mechanisms matter below: `native_decide` discharges a decidable proposition by running compiled code, which extends the trusted base from the kernel to Lean’s compiler and evaluator; and `#print axioms` reports, for any theorem, exactly which axioms (including any such evaluator appeal) its proof ultimately uses. The formalization comprises:

- *The elementary layer*, as ordinary kernel-checked proofs: certificate soundness and integrality, both directions of Theorem 1.2, both families of Proposition 1.13, Theorem 1.9, and the ladder’s elementary layer — Lemmas 5.14 and 5.15 in full, and the selected-residual identity of Section 5.4 (the formalization shows the coprimality hypothesis of Lemma 5.14 is not needed: for $q \mid R$ both symbols vanish).
- *The enumeration layer*: the reachability model underlying Theorem 1.3 is defined in the multiplicative structure of $\mathbb{Z}/R$ directly, and both halves of Lemma 5.1 are proved as theorems, with no computational content to trust. The finite checks behind Theorem 1.4 (1,536 configurations, checked as the 384 that remain after the neutral class collapses), Theorem 1.5 (497,664 capped configurations, including the exact two-case failure classification), and Lemma 1.10 at $R = 19, 23$ (437,580 resp. 7,759,752 checks), and Theorem 1.6 (349,920 configurations, via a product-index bridge for the non-cyclic unit group) are discharged by `native_decide`.
- *The coordinate bridges*: the larger enumerations are *evaluated* in discrete-log coordinates for speed, but the translation to the multiplicative model is itself a proved theorem, so the coordinates add nothing to the trusted base; the $R = 11, 19, 23$ statements hold in the multiplicative model, each reporting exactly its enumeration’s single evaluator axiom.
- *The reduction of Theorem 1.3*: membership of the target class in the reachability model is proved equivalent to the existence of a divisor $k \mid m^2$ in that class (via the fundamental theorem of arithmetic), and produces the explicit positive certificate (k, b, c) solving (1.1) — closing the loop between the abstract model the enumerations check and the integer certificates the conjecture asks for.
- *Two capstones*: a composed corollary (if $(p + 19)/4$ has nine prime factors in pairwise-distinct unit classes other than 1 modulo 19, an explicit solution of (1.1) at p exists, machine-checked end to end), and Lemma 1.10 at $R = 31$ by *certified dynamic programming* — the $\binom{29}{15} = 77,558,760$ supports are never enumerated; a machine-checked induction shows the dynamic program covers every support, and the compiled evaluator runs the dynamic program itself together with the final check that every state realized by ≥ 15 classes hits all 30 targets (3,001 states in the run) — the program’s *correctness* is proved; its evaluation is what `native_decide` trusts.

Every main theorem is audited with `#print axioms`, and helper lemmas are covered transitively through them. The elementary layer is now formalized outright, including Proposition 1.1 in both directions (a solution is a certificate at its own residual, and conversely — so the equivalence organizing this paper is machine-checked) and Propositions 2.1 and 2.2. Theorem 1.6’s finite

check is formalized by a product-index mask bridge over the non-cyclic $(\mathbb{Z}/15)^* \cong C_2 \times C_4$, and Theorem 1.11 fully symbolically: Kneser’s addition theorem vendored from its existing sorry-free formalization, an iterated many-summand form proved on top, and the support bound derived for an *arbitrary* finite abelian group with an explicit involution count — part (iii)’s general form, from which parts (i) and the odd-order variant follow as instances. So too the branch-classification layer: Theorem 5.9’s single kernel step in all three parts, with the quotient *constructed* rather than assumed and the container size bound $\leq d/2 - 1$ derived from the support bounds downstairs; Proposition 5.13 in full; and Theorem 5.8 end to end — the complete-partition machinery, the maximality of the constructed supports, their injectivity, and an explicit $2^{c\sqrt{d}}$ lower bound. (Not formalized: the transfer of Theorem 1.11(ii) to the multiplicative divisor-class model, the identity $t = 2^{\omega(R)} - 1$ for the involution count of $(\mathbb{Z}/R)^*$ — so the general bound is quoted with that t supplied by hand — the *iterated* kernel reduction, and the analytic arguments of Sections 5.1 and 5.4, which rest on the hand proofs alone.) For these enumerations the sixty-line trust base above is thus replaced by the Lean toolchain evaluating a formally stated proposition — and the Python and Lean implementations, written independently against the same statements, agree. (The development notes and Lean sources use working letter names for several results; the repository’s documentation includes the dictionary.)

In addition (as sanity checks, not proof components), the criteria of Theorems 1.2, 1.4, 1.5 and 1.6 were compared against ground-truth divisor search on the complete solvability data below 10^9 : agreement in 1,587,581/1,587,581 cases for $R = 3$, 41,421/41,421 sampled failure cases for $R = 7$ (all character failures, as Theorem 1.4 requires), 158,759/158,759 sampled cases for $R = 15$, and 79,380 sampled cases for $R = 11$, with zero discrepancies; the general composite- R engine (`solvable_exact_general`) likewise agrees on 7,938 sampled primes for every composite admissible residual, and the strong form of Theorem 1.11(ii) is confirmed numerically at every residual checked (`verify_support_bound_strong`). One check named in this paragraph’s vicinity is *not* a sanity check: `kneser_support_general` is a proof component of the $R \leq 107$ clause of Theorem 1.11(iii), and is listed as such in Table 5. The identity of Theorem 1.9 and its forced-prime corollary were likewise checked on 135,588 resp. 61,438 instances from the same data, with zero violations (Corollary 2.3).

A.2 Certificate datasets

The datasets underlying the census, and the verification applied to each, are tabulated in Section 4.1 (Table 2); all are archived in the repository under `data/`, with sha256 manifests.

A.3 Reproduction

With the repository installed (`pip install -e .`), the commands of Table 6 regenerate and re-verify everything above. Generation of the full 10^{10} map takes about 11 minutes on four cores; every other command runs in seconds to minutes.

The ladder census of Section 5.4 is reproduced by `python -m erdos_straus.burgess_scan` (the 10^9 population) and by the scaled drivers in the same module for the window samples and the deep tail; the archives are the five `burgess_*.json` files under `data/analysis/`.

The full test suite (`python -m pytest`; 56 tests) exercises the proof-component checks, the exact criteria against ground truth, and sampled dataset validation on every run. Table 7 summarizes the computational environment.

Data and code availability

All code, all datasets of Table 2, and the Lean development are openly available in the repository <https://github.com/jeffmhopkins/Erd-s-Straus-attack>.

result	proof inputs	proved in	machine-checked?
Prop. 1.1	elementary	§1.1	Lean, both directions
Props. 2.1, 2.2	elementary (characters; explicit k)	§2	Lean
Thm. 1.2 ($R=3$)	elementary	§3	Lean, both directions
Thm. 1.3	elementary	§3	model + reduction in Lean
Thm. 1.4 ($R=7$)	Prop. 2.1 + Thm. 1.3 + enumeration (1,536)	§3	Lean (<code>native_decide</code>)
Thm. 1.5 ($R=11$)	Thm. 1.3 + enumeration (497,664)	§3	Lean (<code>native_decide</code>)
Thm. 1.6 ($R=15$)	Thm. 1.3, composite form + enumeration (349,920)	§3	Lean (<code>native_decide</code>)
Thm. 1.9	quadratic reciprocity	§2	Lean
Lem. 5.10	Lem. 5.1 + enumerations / DP	§5.1	Lean at $R = 19, 23, 31$
Thm. 1.11	Kneser’s theorem + Lem. 5.1; (iii) upgrade at $R \leq 107$: enumeration	§5.1	Lean, parts (i) and (iii); (iii) upgrade: Python
Thm. 1.12	Thms. 1.2, 1.4–1.6, 1.11 + Fundamental Lemma	§5.1	no
Prop. 1.13	elementary + Thm. 1.12	§5.1	families in Lean
Thm. 1.14	Prop. 2.2 + Fundamental Lemma	§5.1	no
Thm. 5.5	Thm. 1.11(iii), general tier + Brun’s sieve, uniform constants	§5.1	no
Thm. 5.8	explicit construction (complete partitions)	§5.1	Lean
Thm. 5.9	Kneser + Lem. 5.1 + quotient induction	§5.1	Lean, single step
Thm. 5.11	Thm. 5.9 + Conjecture 5.10	§5.1	no (conditional)
Prop. 5.13	elementary (isolation of punctures)	§5.1	Lean
Thm. 5.12	Thm. 5.11 + the audit of Thm. 5.5	§5.1	no (conditional)
Lems. 5.14–5.15	Jacobi reciprocity, elementary	§5.4	Lean
Thms. 5.16, 5.17	Thm. 1.11 / Prop. 2.2 + sieve	§5.4	no
Thm. 5.19	norm-form bridge (5.2) + [19], [17] + Fundamental Lemma	§5.4	no
Thm. 5.22	Hypothesis 5.21 (measured) + large sieve	§5.4	no (conditional)

Table 1: Audit map: what each main result depends on, where it is proved, and its formal-verification status. “—” means the result is elementary enough that no machine check was warranted; “no” means the proof is by hand only. Theorem 5.19 is established in outline (proof sketch).

dataset	entries	verification performed
explicit triples, $p < 1.2 \times 10^8$	213,131	every triple checked exactly
minimal- R map, $p < 10^9$	1,587,581	<i>exhaustive</i> : every entry reconstructed and checked exactly
minimal- R map, $p < 10^{10}$	14,215,707	143,000 systematically sampled entries reconstructed and checked; additionally <i>all</i> 842 tail entries ($R \geq 51$) checked with full minimality (no smaller admissible R succeeds)
solvability masks, $p < 10^9$	$27 \times 1,587,581$	per-prime solvability of every admissible $R \leq 107$ by full divisor-class computation; consistency with the minimal- R map
R -sequence dataset, $p < 10^{11}$	128,671,219	257,343 systematically sampled entries reconstructed and checked; <i>all</i> 20,151 tail entries ($R \geq 43$) checked with full minimality; prime array pinned by sha256; pipeline validated byte-for-byte against the exhaustively verified 10^9 map
R -sequence dataset, $p < 10^{12}$	1,175,215,396	2,350,431 systematically sampled entries reconstructed and checked; <i>all</i> 118,210 tail entries ($R \geq 43$) checked with full minimality, including the three attaining the record $R = 111$; zero unsolved; both arrays pinned by sha256; single 10.5 h run on 32 cores, regenerated by <code>scripts/run_1e12.sh</code>

Table 2: Certificate datasets and the verification applied to each.

R	count	R	count	R	count
3	657,645,152	43	43,558	83	78
7	286,441,544	47	50,190	87	40
11	167,419,973	51	8,756	91	5
15	33,222,834	55	7,441	95	12
19	19,307,426	59	4,863	99	4
23	7,996,742	63	1,439	103	5
27	1,376,271	67	488	107	4
31	1,318,258	71	970	111	3
35	208,713	75	146		
39	160,273	79	208		

Table 3: Distribution of the minimal residual $R_{\min}$ over the 1,175,215,396 hard primes below 10^{12} (plotted in Figure 1). The value $R = 111$ is new at this scale.

bin ($p \sim$)	10^6	10^7	6×10^7	2.5×10^8	10^9
density $\times \log x$	5.14	5.19	5.18	5.17	5.16

Table 4: The proportion of hard primes near x that fail both $R = 3$ and $R = 7$, multiplied by $\log x$, in five bins spanning three decades of the 10^9 data (x is the bin's representative value). Constancy of the product is the decay rate predicted by the base case of Theorem 1.12.

component	configuration space	count	entry point
Thm. 1.4 ($R = 7$)	class multisets in $(\mathbb{Z}/7)^*$, multiplicities ≤ 3 , with the forced even factor, $p \bmod 7 \in \{1, 2, 4\}$, and the consistency relation	1,536	<code>verify_R7_finite</code>
Thm. 1.5 ($R = 11$)	dynamic program over $(\mathbb{Z}/11)^*$ with the forced factor 3; collapses to 25 equivalence states (16 succeed, 6 fail by character, 3 by budget)	25	<code>finite_criterion_dp</code>
Thm. 1.6 ($R = 15$)	class multisets in $(\mathbb{Z}/15)^*$, multiplicities ≤ 5 , with the forced even factor, $p \bmod 15 \in \{1, 4\}$, and the consistency relation (349,380 succeed, 540 fail all-residue, zero budget failures)	349,920	<code>verify_R15_finite</code>
Lem. 1.10 ($R = 19$)	all supports of size 9 over the nonzero logs $\times$ all classes of p , at minimal multiplicities	437,580	<code>verify_support_bound</code>
Lem. 1.10 ($R = 23$)	all supports of size 11, likewise	7,759,752	<code>verify_support_bound</code>
Lem. 1.10 ($R = 19$ –107, all twelve primes)	subset dynamic program over reachability-mask states; state counts 271 ($R=19$), 3,001 ($R=31$), 2,414,145 ($R=83$), 19,530,971 ($R=107$); zero violations throughout	$\approx 19.5\text{M}$	<code>verify_support_bound_dp</code>
Thm. 1.11(iii), $R \leq 107$ clause	maximal support of a proper product-set in $(\mathbb{Z}/R)^*$ equals $\varphi(R)/2 - 1$ for every composite admissible $R \leq 107$ (e.g. $R = 95$: max 35 over 583,455 realizable product-sets)	12 moduli	<code>kneser_support_general</code>

Table 5: Machine-verified proof components: enumerated configuration spaces and the entry points in `erdos_straus.theory`.

claim	command
Thm. 1.4 finite check	<code>python -c "from erdos_straus.theory \ import verify_R7_finite; \ print(verify_R7_finite())"</code>
Thm. 1.5 enumeration	<code>finite_criterion_dp(11, forced_logs=[_dlog_table(11)[1][3]]) from erdos_straus.theory (as in the test suite, tests/test_solver.py)</code>
Lem. 1.10 at $R = 19, 23$	<code>verify_support_bound(19), (23)</code>
Thm. 1.6 finite check	<code>python -c "from erdos_straus.theory \ import verify_R15_finite; \ print(verify_R15_finite())"</code>
Thm. 1.11 numerical confirmation	<code>verify_support_bound_strong(31), kneser_support_general(15)</code>
generate the 10^{10} map	<code>python -m erdos_straus.bulk_generate \ -max 10000000000 -segmented -store rmap</code>
verify a dataset exhaustively	<code>python -m erdos_straus.verify data/<file></code>
distribution / covering / taxonomy	<code>python -m erdos_straus.analyze dist \ -rmap data/hard_primes_1e9_minimalR.json.gz (also cover, tail); python -m erdos_straus.theory</code>
Lean formalization (elementary + enumeration layers)	<code>cd lean/ErdosStraus lake exe cache get && lake build</code>

Table 6: Reproduction commands. Shell commands are given verbatim, with `\` marking an explicit line continuation within a single command; the finite-check rows are Python calls from the module `erdos_straus.theory`.

implementation	Python 3, exact integer arithmetic only; no external libraries in any proof-component check (sympy appears only in the reference solver)
factorization	trial division by sieved primes up to $\sqrt{n}$ (numpy-sieved prime tables for bulk generation)
wall-clock times	10^{12} generation: 10.5 h on 32 cores (AMD Ryzen AI Max+ 395), of which 53 min sieving; 10^{11} : 2.85 h on 4 cores; 10^{10} : 656 s; 10^9 : 49 s; largest proof-component check ($R = 107$ dynamic program): 331 s
verification	every certificate re-checked via $4abc = p(bc+ca+ab)$; exhaustive below 10^9 ; 2,350,431 samples plus all 118,210 tail entries at 10^{12} (Table 2)
integrity	datasets archived in the repository; prime arrays and R -sequences pinned by SHA-256 in the corresponding meta.json (at 10^{12} , prefixes ae1e4e76... , 7736d10b...); the 10^{12} value array, 304 MB, is regenerated deterministically by scripts/run_1e12.sh and checked against its stored digest
formal layer	Lean 4.32.2 + mathlib (pinned toolchain); six native_decide checks; Theorem 1.11(i) fully symbolic (vendored Kneser); every main theorem audited by #print axioms

Table 7: Computational environment and integrity summary.

References

- [1] N. C. Ankeny, *The least quadratic non residue*, Ann. of Math. (2) **55** (1952), 65–72.
- [2] D. A. Burgess, *The distribution of quadratic residues and non-residues*, Mathematika **4** (1957), 106–112.
- [3] L. de Moura and S. Ullrich, *The Lean 4 theorem prover and programming language*, in: Automated Deduction – CADE 28, Lecture Notes in Computer Science, vol. 12699, Springer, Cham, 2021, pp. 625–635.
- [4] C. Elsholtz and T. Tao, *Counting the number of solutions to the Erdős–Straus equation on unit fractions*, J. Aust. Math. Soc. **94** (2013), 50–105.
- [5] P. Erdős, *Az $1/x_1 + 1/x_2 + \cdots + 1/x_n = a/b$ egyenlet egész számú megoldásairól* [On the integer solutions of the equation $1/x_1 + 1/x_2 + \cdots + 1/x_n = a/b$] (Hungarian), Mat. Lapok **1** (1950), 192–210.
- [6] J. Friedlander and H. Iwaniec, *Opera de Cribro*, American Mathematical Society Colloquium Publications, vol. 57, American Mathematical Society, Providence, RI, 2010.
- [7] W. Gao, Y. O. Hamidoune, A. Lladó and O. Serra, *Covering a finite abelian group by subset sums*, Combinatorica **23** (2003), 599–611.
- [8] H. Halberstam and H.-E. Richert, *Sieve Methods*, Academic Press, London, 1974.
- [9] K. Bradford, *A note on the Erdős–Straus conjecture*, Integers **21** (2021), #A24; preprint, arXiv:1906.00561 [math.NT], 2019.
- [10] M. Bello-Hernández, M. Benito and E. Fernández, *A divisor parametrization for the Erdős–Straus conjecture*, preprint, arXiv:2606.10922 [math.NT], 2026.
- [11] G. E. Andrews, G. Beck and B. Hopkins, *Complete partitions*, Fibonacci Quart. **36** (1998), 354–360.
- [12] D. J. Gryniewicz, *Structural Additive Theory*, Developments in Mathematics, vol. 30, Springer, Cham, 2013.
- [13] G. T. Diderrich and H. B. Mann, *Combinatorial problems in finite abelian groups*, in: J. N. Srivastava et al. (eds.), A Survey of Combinatorial Theory, North-Holland, Amsterdam, 1973, pp. 95–100.
- [14] M. Freeze, W. Gao and A. Geroldinger, *The critical number of finite abelian groups*, J. Number Theory **129** (2009), 2766–2777.
- [15] M. A. Lopez, *A complete congruence system for the Erdős–Straus conjecture*, preprint, arXiv:2404.01508 [math.NT], 2024.
- [16] D. A. Cox, *Primes of the Form $x^2 + ny^2$: Fermat, Class Field Theory, and Complex Multiplication*, 2nd ed., Wiley, Hoboken, NJ, 2013.
- [17] E. Fuchs, K. Hsu, J. Rickards, D. Schindler and K. Stange, *Primitive representations by binary quadratic forms*, preprint, arXiv:2504.20289 [math.NT], 2025.

- [18] H. Iwaniec, *The half dimensional sieve*, Acta Arith. **29** (1976), 69–95.
- [19] H. Iwaniec, *Primes of the type $\varphi(x, y) + A$ where φ is a quadratic form*, Acta Arith. **21** (1972), 203–234.
- [20] M. Kneser, *Abschätzung der asymptotischen Dichte von Summenmengen*, Math. Z. **58** (1953), 459–484.
- [21] The mathlib Community, *The Lean mathematical library*, in: Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2020), ACM, New York, 2020, pp. 367–381.
- [22] S. Mihnea and D. C. Bogdan, *Further verification and empirical evidence for the Erdős–Straus conjecture*, preprint, arXiv:2509.00128 [math.NT], 2025.
- [23] L. J. Mordell, *Diophantine Equations*, Academic Press, London, 1969.
- [24] S. E. Salez, *The Erdős–Straus conjecture: new modular equations and checking up to $N = 10^{17}$* , preprint, arXiv:1406.6307 [math.NT], 2014.
- [25] J. W. Sander, *On $4/n = 1/x + 1/y + 1/z$ and Rosser’s sieve*, Acta Arith. **59** (1991), 183–204.
- [26] A. Swett, *The Erdős–Straus conjecture*, online verification record to 10^{14} , 1999; formerly at <http://math.uindy.edu/swett/esc.htm> (archived snapshots available via the Internet Archive; accessed August 2026).
- [27] D. G. Terzi, *On a conjecture by Erdős–Straus*, Nordisk Tidskr. Informationsbehandling (BIT) **11** (1971), 212–216.
- [28] T. Tao and V. H. Vu, *Additive Combinatorics*, Cambridge Studies in Advanced Mathematics, vol. 105, Cambridge University Press, Cambridge, 2006.
- [29] R. C. Vaughan, *On a problem of Erdős, Straus and Schinzel*, Mathematika **17** (1970), 193–198.
- [30] W. A. Webb, *On $4/n = 1/x + 1/y + 1/z$* , Proc. Amer. Math. Soc. **25** (1970), 578–584.
- [31] K. Yamamoto, *On the Diophantine equation $4/n = 1/x + 1/y + 1/z$* , Mem. Fac. Sci. Kyushu Univ. Ser. A **19** (1965), 37–47.