

Residual certificates for the Erdős–Straus conjecture: exact solvability criteria, a sieve bound, and computations to 10^{11}

J. M. Hopkins*

August 2026

Abstract

The Erdős–Straus conjecture asserts that $4/n = 1/a + 1/b + 1/c$ is solvable in positive integers for every $n \geq 2$; the open cases concentrate on primes in six residue classes mod 840 (*hard primes*). We develop the *residual method*: a solution with first denominator $a = (p + R)/4$ exists if and only if some divisor $k \mid m^2$, $m = pa$, satisfies $k \equiv -m \pmod{R}$. Our central results make this criterion exact and finite: solvability at any fixed prime residual R is a computable function of the factor classes of $(p + R)/4$ and of $p \bmod R$ (a finite-state meta-theorem), yielding closed-form criteria at $R = 3, 7, 11$ and a reciprocity identity $\left(\frac{q}{R}\right) = \left(\frac{p}{q}\right)$ that explains how failures correlate across residuals. Consequences include an unconditional chain of sieve bounds — exceptional sets $O(x/(\log x)^\kappa)$ with κ up to $19/2$, each cut out by explicit conditions with constructive certificates. These are weaker in density than Vaughan’s classical $x \exp(-c(\log x)^{2/3})$, but mechanistic where mean-value methods are not: every bound produces certificates and exact local criteria, and the elementary and finite-enumeration layers of the proofs are machine-checked in Lean 4/mathlib. Computationally we exhibit certificates for all 128,671,219 hard primes below 10^{11} , verified exhaustively to 10^9 and by systematic sampling plus full tail verification beyond; the record minimal residual is 107, first attained below 10^7 and never exceeded, and the calibrated model correctly predicted the population of the last empty band of minimal residuals between 10^{10} and 10^{11} .

1 Introduction

Erdős and Straus conjectured (see [2]) that for every integer $n \geq 2$ the equation

$$\frac{4}{n} = \frac{1}{a} + \frac{1}{b} + \frac{1}{c} \tag{1.1}$$

has a solution in positive integers a, b, c . By multiplicativity it suffices to treat prime n , and classical polynomial identities settle every prime outside the six residue classes

$$\mathcal{H} \bmod 840 = \{1, 121, 169, 289, 361, 529\},$$

the squares of units modulo 840 (Mordell [6, Ch. 30]). We call primes in these classes *hard primes*. Direct verification has confirmed the conjecture far beyond any range considered here:

*jeff.m.hopkins@gmail.com. Computations and computer-assisted case analyses were carried out with the code repository accompanying this paper (<https://github.com/jeffmhopkins/Erd-s-Straus-attack>).

to 10^{14} by Swett [8], to 10^{17} by Salez [7], and to 10^{18} by Mihnea and Bogdan [5]; the present paper is logically independent of those bounds. The structure of the solution count was analyzed by Elsholtz and Tao [1]. Verification, however, carries no structural information; the aim here is a quantitative structural reduction. For calibration, Vaughan’s classical mean-value bound [9] gives an exceptional set $\ll x \exp(-c(\log x)^{2/3})$ for (1.1) — asymptotically stronger than any fixed power of $\log x$; the value of the present approach is not raw density but its explicit, certificate-generating, machine-verified structure, together with the exact criteria.

New contributions

- The *residual formulation* of the Erdős–Straus equation: solvability with first denominator $(p + R)/4$ is a divisor-class condition, complete over all solutions (Section 1.1).
- *Exact* solvability criteria at $R = 3, 7, 11$ (Theorems 1.1, 1.3, 1.4), and a meta-theorem (Theorem 1.2) showing every fixed prime residual admits a finite computable exact criterion.
- A reciprocity principle (Theorem 1.6) collapsing the per-residual conditions into one global statement about p — the organizing principle behind correlated failure and the record prime.
- A constructive sieve chain built on these criteria (Lemma 1.7, Theorem 1.8, Proposition 1.9): explicit, certificate-producing exceptional-set bounds, with the machine-verified support-bound lemma as the per-residual input; and an unconditional density reduction (Theorem 1.10): for every A , a finite residual set whose joint failure set is $O_A(x/(\log x)^{1+A})$.
- Computed and verified residual certificates for all 128,671,219 hard primes below 10^{11} , with a calibrated model whose deep-tail predictions were subsequently confirmed (Section 5).
- A machine-checked (Lean 4/mathlib) verification of the elementary layer — certificate soundness and integrality, Theorem 1.1 in both directions, both families of Proposition 1.9, and Theorem 1.6 — and of the finite-enumeration layer: the divisor-class reachability model with its monotonicity reductions (Lemma 4.1), the full finite verifications of Theorems 1.3 and 1.4, the support-bound checks of Lemma 1.7 at $R = 19$ and 23 (and at $R = 31$ by certified dynamic programming), the proved discrete-log bridges carrying the $R = 11, 19, 23$ enumerations into the multiplicative model, and the model-faithfulness reduction of Theorem 1.2 (reachability in the model is equivalent to the existence of a divisor certificate (1.2), and produces one explicitly) — in the accompanying repository.

1.1 The residual method

Fix a hard prime p . All six classes satisfy

$$(H1) \ p \equiv 1 \pmod{8}, \quad (H2) \ p \equiv 1 \pmod{3}, \quad (H3) \ p \bmod 7 \in \{1, 2, 4\}.$$

For an *admissible residual* $R \equiv 3 \pmod{4}$ put

$$a = \frac{p + R}{4} \in \mathbb{Z}, \quad m = pa.$$

Substituting a into (1.1) and clearing denominators shows that a solution with first denominator a exists if and only if

$$\exists k \mid m^2 \quad \text{with} \quad k \equiv -m \pmod{R}, \quad (1.2)$$

in which case $b = (k + m)/R$ and $c = (m^2/k + m)/R$ are positive integers solving (1.1). We write $R_{\min}(p)$ for the least admissible R for which (1.2) holds. The formulation is *complete*: any solution with denominators $a \leq b \leq c$ has $p/4 < a \leq 3p/4$, so $R = 4a - p \in (0, 2p]$ with $R \equiv 3 \pmod{4}$ automatic — every solution of (1.1) at p is a residual certificate, and $R_{\min}(p) \leq 2p$ is equivalent to the conjecture at p . Two elementary congruences drive everything below. From $4a = p + R$,

$$a \equiv 4^{-1}p, \quad m \equiv 4^{-1}p^2 \pmod{R}, \quad (1.3)$$

and since $\gcd(p, R) = 1$ (as $p > R$ is prime), also $\gcd(a, R) = \gcd(m, R) = 1$: the setting is always a unit computation modulo R . By (1.3), m is a *square* modulo every prime $r \mid R$; hence for every prime $r \mid R$ with $r \equiv 3 \pmod{4}$ (at least one exists, since $R \equiv 3 \pmod{4}$),

$$\left(\frac{-m}{r}\right) = \left(\frac{-1}{r}\right) = -1 : \quad (1.4)$$

the target class $-m$ is a quadratic non-residue mod r .

A worked example. Take the hard prime $p = 1009 \pmod{840}$ and the smallest admissible residual $R = 3$. Then $a = (1009 + 3)/4 = 253 = 11 \cdot 23$ and $m = pa = 255,277$. Since $11 \equiv 2 \pmod{3}$, Theorem 1.1 below promises a certificate: take $k = 11$, a divisor of m^2 with $k \equiv -m \equiv 2 \pmod{3}$. Then $b = (k + m)/3 = 85,096$ and $c = (m^2/k + m)/3 = 1,974,822,872$, and indeed

$$\frac{4}{1009} = \frac{1}{253} + \frac{1}{85,096} + \frac{1}{1,974,822,872}.$$

Had every prime factor of a been $\equiv 1 \pmod{3}$, every divisor of m^2 would be $\equiv 1 \pmod{3}$ and the required class 2 would be unreachable — the prototype of all the obstructions studied below.

The shape of what follows. Everything in this paper grows out of the picture just seen, read at three levels. *Locally* (one residual): whether the target class $-m$ is attainable from the factor classes of a is a question about a finite multiplicative state — the multiset of classes with bounded exponent budgets — so each fixed R admits an exact, machine-enumerable solvability criterion (Theorems 1.1–1.4 below). *Globally* (all residuals at once): the conditions entering these criteria are quadratic characters, and reciprocity folds the condition at every residual into signs $\left(\frac{p}{q}\right)$ indexed by the primes q dividing the shifted values $(p + R)/4$ — failures at different residuals correlate because shifted values share prime factors (Theorem 1.6). *Quantitatively*: each criterion says that failure forces $(p + R)/4$ to avoid prime factors in many residue classes, a sifting condition; stacking residuals multiplies the savings (Theorem 1.8), though it provably cannot reach emptiness (Remark 4.2) — the honest endpoint of the method, discussed in Section 6.

1.2 Main results

Theorem 1.1 ($R = 3$; exact). *Let $p \equiv 1 \pmod{4}$ with $3 \nmid p$, and $a = (p + 3)/4$. Then (1.2) holds for $R = 3$ if and only if a has a prime factor $q \equiv 2 \pmod{3}$. When it does, $k = q$ is a certificate.*

Modulo 3, character equals class and the criterion is elementary. For every other fixed residual, the reduction to a finite check is itself a theorem, and the next two results are its first showcase instances.

Theorem 1.2 (meta-theorem). *For every fixed prime $R \equiv 3 \pmod{4}$, the validity of (1.2) at a hard prime p depends only on (i) the multiset of classes mod R of the prime factors of $a = (p + R)/4$, with multiplicities capped at $\lceil (R - 1)/2 \rceil$, and (ii) $p \pmod{R}$. Consequently an exact solvability criterion exists for every R and is computable by finite enumeration.*

Theorem 1.3 ($R = 7$, hard primes; exact). *Let p be a hard prime and $a = (p + 7)/4$. Then (1.2) holds for $R = 7$ if and only if a has a prime factor that is a quadratic non-residue modulo 7.*

Theorem 1.4 ($R = 11$, hard primes; exact). *Let p be a hard prime and $a = (p + 11)/4$; note $3 \mid a$, $2 \nmid a$, $7 \nmid a$ automatically. Then (1.2) fails for $R = 11$ if and only if either*

- (a) *every prime factor of a is a quadratic residue mod 11; or*
- (b) *$v_3(a) = 1$, every other prime factor of a is $\equiv 1 \pmod{11}$ except a non-residue part — its factors in the non-residue classes mod 11 — of exactly one of the shapes*

$$\{q \equiv 2\}, \quad \{q \equiv 6\}, \quad \{q \equiv 2\} \cup \{q' \equiv 6\} \pmod{11},$$

each with multiplicity one, paired respectively with $p \equiv 2, 6, 1 \pmod{11}$.

Remark 1.5 (computer-assisted components). The exhaustive verifications behind Theorems 1.3, 1.4 and Lemma 1.7 below are *parts of the respective proofs*, on the same footing as a hand case analysis. Each is a deterministic enumeration of a finite configuration space whose finiteness and sufficiency are proved in advance (Theorem 1.2; for Lemma 1.7, monotonicity of target-reachability in support and multiplicities, checked at minimal exponents). The checks are short (seconds each, up to 331 s for the largest, $R = 107$), involve only exact integer arithmetic, and the code is available in the accompanying repository; Table 3 records the exact configuration counts and entry points. These enumerations are additionally formalized in Lean 4: the $R = 7$ verification of Theorem 1.3, the $R = 11$ criterion of Theorem 1.4, and the $R = 19, 23$ cases of Lemma 1.7 are discharged by Lean’s compiled evaluator (and the $R = 31$ case by a certified dynamic program), with the monotonicity reductions (Lemma 4.1) *and* the discrete-log coordinate translations proved as theorems, so that the four enumerations at $R = 7, 11, 19, 23$ hold as machine-checked statements directly in the multiplicative model (Appendix A).

The obstruction framework of Section 2 yields a reciprocity identity that converts the per-residual conditions into a single global statement about p — a structural interpretation of joint failure across residuals:

Theorem 1.6 (reciprocity structure). *Let p be a hard prime, $R \equiv 3 \pmod{4}$ prime, and q an odd prime factor of $a = (p + R)/4$ (necessarily $q \neq R$). Then*

$$\left(\frac{q}{R}\right) = \left(\frac{p}{q}\right).$$

The sieve consequences below need one input per residual: a statement that failure at R forces the shifted form $(p + R)/4$ to avoid prime factors in a positive-density set of residue classes. For $R = 3, 7, 11$ the exact criteria provide this; for larger R it is the following machine-verified lemma (proved in Section 4).

Lemma 1.7 (support bound; verified for every prime $R \equiv 3 \pmod{4}$, $19 \leq R \leq 107$). *Let R be one of the twelve primes 19, 23, 31, 43, 47, 59, 67, 71, 79, 83, 103, 107, and call the set of unit classes mod R occupied by the prime factors of $(p+R)/4$ its support. Every configuration for which residual R fails at a hard prime has support of size at most $(R-1)/2$, at most $(R-3)/2$ of whose classes differ from the identity class (which never affects solvability). Consequently failure at R forbids at least half the classes, through finitely many maximal-support branches.*

Theorem 1.8 (the chain). *The number of hard primes $p \leq x$ failing every residual in $\{3, 7\}$, in $\{3, 7, 11\}$, in $\{3, 7, 11, 19\}$, and in $\{3, 7, 11, 19, 23\}$ is, respectively,*

$$O\left(\frac{x}{(\log x)^2}\right), \quad O\left(\frac{x}{(\log x)^{5/2}}\right), \quad O\left(\frac{x}{(\log x)^3}\right), \quad O\left(\frac{x}{(\log x)^{7/2}}\right).$$

More generally, let $\mathcal{P} = \{3, 7, 11, 19, 23, 31, 43, 47, 59, 67, 71, 79, 83, 103, 107\}$ be the set of all fifteen primes $\equiv 3 \pmod{4}$ up to 107. The number of hard primes failing every residual in $\mathcal{P}$ is

$$O\left(\frac{x}{(\log x)^{17/2}}\right).$$

In particular all but a proportion $O((\log x)^{-2})$ of hard primes are solved within $\{3, 7, 11, 19\}$, and all but $O((\log x)^{-15/2})$ within $\mathcal{P}$ — one power of $\log x$ less than the counts above, since there are $\asymp x/\log x$ hard primes up to x .

Proposition 1.9 (aggregate identity families). *Let $p \equiv 1 \pmod{4}$ be prime and $R \equiv 3 \pmod{4}$. If $R \mid p+1$, then $k = ap^2$ is a certificate for residual R ; if $R \mid p+4$, then $k = a^2p$ is. Hence any prime factor $\equiv 3 \pmod{4}$ of $p+1$ or of $p+4$ yields an explicit solution, and these are the only a -independent divisor shapes $k = a^j p^i$ with this property. Adding the two corresponding linear forms to Theorem 1.8: the number of hard primes $p \leq x$ failing every residual in $\mathcal{P}$ and both families is $O(x/(\log x)^{19/2})$.*

Theorem 1.10 (density reduction). *For every $A > 0$ there is a finite admissible set $S(A)$ such that the number of hard primes $p \leq x$ for which every residual in $S(A)$ fails is $O_A(x/(\log x)^{1+A})$.*

Remark 1.11 (relation to Vaughan’s bound). Theorem 1.10 and the chain of Theorem 1.8 are *weaker in density* than Vaughan’s bound $\ll x \exp(-c(\log x)^{2/3})$ [9], which already exceeds every fixed power of $\log x$. Their interest is orthogonal: the exceptional sets here are cut out by finitely many *explicit* multiplicative conditions on shifted forms, each failure of which is constructively certified by an exhibited divisor k , and the underlying lemmas are exact or machine-verified. Vaughan’s method counts solutions on average and exhibits neither certificates nor criteria.

Organization. Proofs are arranged by logical dependency. Section 2 develops the obstruction framework (Propositions 2.1, 2.2) and proves Theorem 1.6, whose Corollary 2.3 depends on the Type-I/II dichotomy introduced there. Section 3 proves Theorem 1.1, then Theorem 1.2, then Theorems 1.3 and 1.4, which rest on it. Section 4 proves Theorem 1.8 — base case first, then, after the monotonicity reduction (Lemma 4.1) and Lemma 1.7, the higher exponents — followed by Proposition 1.9 and Theorem 1.10.

Computations (Section 5) provide verified certificates for all hard primes below 10^{11} and calibrate the base-case constant and decay exponents of the chain.

2 The obstruction framework

Proposition 2.1 (character obstruction). *Let $r \mid R$ be prime with $r \equiv 3 \pmod{4}$. If every prime factor q of $m = pa$ satisfies $\left(\frac{q}{r}\right) = +1$, then (1.2) fails for R .*

Proof. Every divisor k of m^2 is a product of prime factors of m , so $\left(\frac{k}{r}\right) = +1$. A certificate requires $k \equiv -m \pmod{r}$, whence $\left(\frac{k}{r}\right) = \left(\frac{-m}{r}\right) = -1$ by (1.4) — impossible. $\square$

We call failures produced by Proposition 2.1 *Type I*; failures where m possesses a non-residue prime factor modulo every relevant r , yet no divisor of m^2 attains the exact class $-m \pmod{R}$, are *Type II*.

Proposition 2.2 (guaranteed-success classes). *Let $t \equiv -4^{-1}p^2 \pmod{R}$. If a has a prime factor q with $q \equiv tp^{-i} \pmod{R}$ for some $i \in \{0, 1, 2\}$, then $k = qp^i$ is a certificate. In particular, since $tp^{-2} \equiv -4^{-1} \pmod{R}$, any prime factor of a in the fixed class $-4^{-1} \pmod{R}$ certifies success regardless of p .*

Proof. $k = qp^i$ divides m^2 (as $q \mid a$ and $i \leq 2$) and lies in class $t \equiv -m$. Integrality of b, c is automatic from $k \equiv -m$ and $m^2/k \equiv m^2(-m)^{-1} \equiv -m \pmod{R}$. $\square$

The contrapositive of Proposition 2.2 is the sieve input for Theorem 1.10: failure at R forces $(p + R)/4$ to avoid prime factors in up to three explicit classes modulo R .

The next identity converts the per-residual character conditions into a single global statement about p , and is the organizing principle behind joint failure across residuals.

Proof of Theorem 1.6. From $q \mid p + R$ we get $R \equiv -p \pmod{q}$. By quadratic reciprocity with $R \equiv 3 \pmod{4}$, $\left(\frac{q}{R}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{R}{q}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{-1}{q}\right) \left(\frac{p}{q}\right) = \left(\frac{p}{q}\right)$. $\square$

Corollary 2.3 (forced primes and correlated failure). *(i) The hard classes are the squares mod 840, so p is automatically a quadratic residue mod 3, 5, 7, and $p \equiv 1 \pmod{8}$ handles $q = 2$ (which divides a only for $R \equiv 7 \pmod{8}$, where $\left(\frac{2}{R}\right) = +1$): the small primes forced into the shifted values by the Mordell structure never break a Type-I failure. (ii) p fails Type-I at every prime residual $R \leq B$ simultaneously if and only if $\left(\frac{p}{R}\right) = +1$ for each such R and $\left(\frac{p}{q}\right) = +1$ for every unforced odd prime q (one not forced by the congruences of (H1)–(H3)) dividing any of the shifted values $(p + R)/4$, $R \leq B$ — one condition per distinct prime, shared across every residual it divides.*

Proof. Immediate from Theorem 1.6 together with (H1)–(H3) and the quadratic characters of the forced small primes. $\square$

Part (i) is a conceptual reading of why the hard classes are exactly the square classes. By part (ii), failures at distinct residuals correlate through shared prime factors: shifted values with few distinct prime factors leave few independent sign conditions. Section 5 quantifies this for the record prime. Both parts are verified at scale (135,588 instances of Theorem 1.6, 61,438 forced-prime instances, zero violations of either; Appendix A).

3 Exact criteria

3.1 Proof of Theorem 1.1

Proof. Modulo 3 the unit group has order 2, so quadratic character equals class. Since $p^2 \equiv 1 \pmod{3}$, (1.3) gives $m \equiv 1$ and the target is $-m \equiv 2 \pmod{3}$; also $3 \nmid a$. *Necessity:* if every prime factor of a is $\equiv 1 \pmod{3}$, then $a \equiv 1 \pmod{3}$, hence $p = 4a - 3 \equiv 1 \pmod{3}$ as well, so every prime factor of $m = pa$ is $\equiv 1 \pmod{3}$ and every divisor of m^2 is $\equiv 1 \pmod{3}$: class 2 is never attained. *Sufficiency:* $k = q \equiv 2 \pmod{3}$ divides m^2 , and the recovery formulae give integers as in Proposition 2.2. $\square$

3.2 Proof of Theorem 1.2

Proof. A certificate is a divisor $k \mid m^2$ in class $-m \pmod{R}$. The set of classes of divisors of m^2 is determined by the classes of the prime factors of $m = pa$ together with their multiplicities, each factor being usable with exponent at most twice its multiplicity; exponents are only meaningful modulo $\text{ord}_R(\text{class}) \leq R - 1$, so multiplicities beyond $\lceil (R - 1)/2 \rceil$ are equivalent to that cap (Remark 3.1 illustrates why). The target class and the consistency relation $\prod(\text{classes}) \equiv a \equiv 4^{-1}p \pmod{R}$ are functions of $p \pmod{R}$ by (1.3). Hence solvability is a function of the capped class multiset and $p \pmod{R}$, a finite amount of data. $\square$

Remark 3.1 (why the cap works, and the size of the space). The one step above that is not pure bookkeeping is the multiplicity cap, and a small example shows both why it is harmless and what it buys. Take $R = 7$ and a factor of a in a class x of order 6 (say $x = 3$): a factor of multiplicity μ may enter a divisor of m^2 with any exponent $0, 1, \dots, 2\mu$, and once $2\mu \geq 6$ these powers already sweep out the full cyclic group $\langle x \rangle$ — raising μ past $\lceil (R - 1)/2 \rceil = 3$ enlarges nothing that a divisor subproduct can reach. For classes of smaller order the cap is generous rather than tight (a class of order 2 stabilizes at $\mu = 1$), which is fine: the enumeration ranges over a finite over-approximation. What finiteness buys is concrete: at $R = 7$ the capped space has 4^6 multisets before the structural constraints (i)–(iii) in the proof of Theorem 1.3 below and 1,536 consistent configurations after them; at $R = 11$ it compresses to 25 equivalence states. The enumeration is decidable, not merely heuristic, precisely because the cap and the consistency relation are proved before any code runs.

3.3 Proof of Theorem 1.3

Proof. $(\Rightarrow)$ is Proposition 2.1 with $r = R = 7$. $(\Leftarrow)$ reduces to a finite verification via Theorem 1.2. Write $(\mathbb{Z}/7)^* \cong \mathbb{Z}/6$ in discrete-log coordinates. Three structural constraints hold at hard primes:

- (i) $p \pmod{7} \in \{1, 2, 4\}$, by (H3);
- (ii) $2 \mid a$: by (H1), $p + 7 \equiv 0 \pmod{8}$, so $a = (p + 7)/4$ is even — the class of 2 (a residue mod 7) is always present;
- (iii) the factor classes multiply to $a \equiv 4^{-1}p \pmod{7}$.

Enumerating all class multisets with multiplicities capped at 3 subject to (i)–(iii) yields 1,536 consistent configurations; for each, one checks whether the target class is attained by a bounded-exponent subproduct. The verification (code: `theory.verify_R7_finite`) confirms: the target is attainable precisely when some factor class is a non-residue. $\square$

Remark 3.2. The hard-class hypotheses are essential: for $p = 701 \equiv 1 \pmod{4}$ (not hard), $a = 177 = 3 \cdot 59$ has the non-residue factor 3, yet $R = 7$ fails. Constraint (ii) — the forced even factor — is exactly what rescues every would-be Type-II configuration at hard primes. Empirically (Section 5), among 41,421 sampled failures of $R = 7$ at hard primes, *all* are Type I.

3.4 Proof of Theorem 1.4

Proof. By Theorem 1.2 with $R = 11$, $d = 10$: multiplicities cap at 5, and at hard primes $3 \mid a$ (from (H2), $11 \equiv 2 \pmod{3}$), while $2 \nmid a$ (from (H1), $p + 11 \equiv 4 \pmod{8}$) and $7 \nmid a$ (from (H3), $-11 \equiv 3 \pmod{7} \notin \{1, 2, 4\}$). A dynamic-programming enumeration over the finite configuration space (code: `theory.finite_criterion_dp`) compresses it to 25 equivalence states: 16 succeed, 6 fail with all factor classes quadratic residues (case (a)), and exactly 3 fail with a non-residue present (case (b)); the three failing states decode to the three listed patterns. Two structural facts emerge from the enumeration. First, no state fails by a *proper-subgroup* obstruction: the only subgroup of $\mathbb{Z}/10$ containing an odd element properly is $\{0, 5\}$, and the consistency relation forces p 's class to extend it to all of $\mathbb{Z}/10$. Second, the case-(b) failures are pure *budget* failures — the exponent bounds, not the character obstruction, block the target: any of $v_3(a) \geq 2$, a factor in class $\{4, 5, 9\}$, or a second non-residue factor restores solvability. The criterion agrees with ground-truth divisor search for 158,759 of 158,759 hard primes tested below 10^9 . $\square$

4 The sieve bound

We now prove the chain of Theorem 1.8, beginning with the base case, and then Lemma 1.7, the aggregate identities, and the density reduction.

Proof of Theorem 1.8, base case $\{3, 7\}$. Step 1 (exact criteria). Let p be a hard prime and set $n = (p + 3)/4$, so that $(p + 7)/4 = n + 1$: *the two shifted forms are consecutive integers*. By Theorems 1.1 and 1.3, p is exceptional (both residuals fail) precisely when

$$n \text{ has no prime factor } \equiv 2 \pmod{3} \quad \text{and} \quad n + 1 \text{ has no prime factor } \equiv 3, 5, 6 \pmod{7},$$

with $p = 4n - 3$ prime.

Step 2 (congruence bookkeeping). Fix one of the six hard classes $h \pmod{840}$; then $n \equiv (h+3)/4 \pmod{210}$ is fixed, and the conditions at the primes 2, 3, 5, 7 are determined by the class (in particular $3 \nmid n$ and $7 \nmid n + 1$ automatically). It suffices to bound each class and sum.

Step 3 (containment in a sifted set). Let $z = x^{1/10}$. After discarding the $O(z)$ primes $p \leq z$, every exceptional p gives $n \leq (x + 3)/4$, $n \equiv n_0 \pmod{210}$, such that for every prime q with $7 < q \leq z$:

- (a) $q \equiv 2 \pmod{3} \implies n \not\equiv 0 \pmod{q}$;
- (b) $q \equiv 3, 5, 6 \pmod{7} \implies n \not\equiv -1 \pmod{q}$;
- (c) $q \nmid 4n - 3$ (since $4n - 3 = p$ is prime, $p > z$).

Conditions (a), (b) only use “no forbidden factor up to z ”, which is implied by the full criteria — an upper bound needs no more.

Step 4 (sieve of dimension 2). This is a standard upper-bound sieve for the three linear forms $n, n+1, 4n-3$, with sifting function

$$\omega(q) = \mathbf{1}_{q \equiv 2(3)} + \mathbf{1}_{q \text{ NQR}(7)} + 1 \quad (q > 7),$$

whose forbidden classes $0, -1, 3 \cdot 4^{-1} \pmod q$ are pairwise distinct for $q \nmid 21$. By Dirichlet, ω has average $\kappa = \frac{1}{2} + \frac{1}{2} + 1 = 2$ in the Halberstam–Richert sense $\sum_{q \leq w} \omega(q) \log q/q = \kappa \log w + O(1)$. The Fundamental Lemma of sieve theory (see [4, Thm. 2.5] and [3, Lem. 6.3]) with $z = x^{1/10}$ gives

$$\#E(x) \ll x \prod_{7 < q \leq z} \left(1 - \frac{\omega(q)}{q}\right) \asymp x (\log z)^{-1} (\log z)^{-1/2} (\log z)^{-1/2} \asymp \frac{x}{(\log x)^2},$$

by Mertens’ theorem and its arithmetic-progression form. Summing over the six classes preserves the bound. The higher exponents of the chain are deduced after Lemma 1.7 below. $\square$

The verification behind Lemma 1.7 reduces to finitely many checks through the following elementary but load-bearing observation, which we isolate for scrutiny.

Lemma 4.1 (monotonicity reduction). *For a class multiset C (classes v with multiplicities μ_v) let $D(C) = \{\sum_v e_v v : 0 \leq e_v \leq 2\mu_v\}$ denote the set of attainable divisor logs (before the p -shifts). If $\mu_v \leq \mu'_v$ for all v , then $D(C) \subseteq D(C')$; likewise enlarging the support enlarges D . Consequently a support S carries some failing configuration (any multiplicities ≥ 1 on S) if and only if the all-multiplicity-one configuration on S fails.*

Proof. Any admissible exponent vector for C is admissible for C' , giving the inclusions. If a configuration on S with multiplicities $\mu_v \geq 1$ fails, the target lies outside $D(C)$, hence outside the subset $D(C_1)$ attained at multiplicities one — so C_1 fails too; the converse inclusion is trivial. The same argument applies verbatim after adjoining the p -shifts. $\square$

Proof of Lemma 1.7. By Lemma 4.1 it suffices to check that for every support of size $(R-1)/2$ over the nonzero discrete logs and every class of $p \pmod R$, the target is reachable already at minimal multiplicities (exponent budget 2 per class, p -budget 2); success there kills every configuration on every superset support at any multiplicities. The check deliberately ignores the consistency relation $\prod(\text{classes}) \equiv 4^{-1}p$: it therefore examines a *superset* of the realizable configurations, and can only overstate the failing supports — the lemma is conservative. For $R = 19, 23$ a direct enumeration of all $\binom{R-2}{(R-1)/2}$ supports suffices (437,580 resp. 7,759,752 pairs; code: `theory.verify_support_bound`). For the larger residuals the binomial count explodes, and the check is instead performed by a subset dynamic program (`theory.verify_support_bound_dp`): fold all supports into a table mapping each reachability mask to the maximal support size achieving it — valid because, by the same monotonicity, failure at any multiplicities implies failure at multiplicity one. The realized masks are heavily structured, so the state count stays tractable (3,001 states at $R = 31$ against 77.5 million supports; 19,530,971 states at $R = 107$, 331 seconds). All twelve residuals verify with zero violations; exact state counts are tabulated in Table 3. The bound is tight: the all-residue (Type I) configurations occupy exactly $(R-3)/2$ nonzero classes. $\square$

Proof of Theorem 1.8, higher exponents. Write $n = (p+3)/4$; the relevant forms are $n, n+1, n+2 = (p+11)/4, n+4 = (p+19)/4, n+5 = (p+23)/4$, and $4n-3 = p$. Joint failure of

$\{3, 7, 11\}$ implies, by Theorems 1.1, 1.3, 1.4: n free of primes $\equiv 2 \pmod{3}$; $n+1$ free of non-residues mod 7; and, splitting by the two cases of Theorem 1.4: *branch (a)* — $n+2$ free of the five non-residue classes mod 11 (density $\frac{1}{2}$); *branch (b)* — $n+2$ free of the six classes $\{4, 5, 7, 8, 9, 10\}$ mod 11 (density $\frac{3}{5}$). Each branch is a four-form sieve problem of dimension $1 + \frac{1}{2} + \frac{1}{2} + \frac{1}{2} = \frac{5}{2}$, resp. $1 + \frac{1}{2} + \frac{1}{2} + \frac{3}{5} = \frac{13}{5}$. Since $\frac{13}{5} > \frac{5}{2}$, the Type-II branch contributes $O(x/(\log x)^{13/5}) = o(x/(\log x)^{5/2})$: the bound is the maximum of the branch exponents, attained by the character branch (a). The Fundamental Lemma applied to each branch gives $O(x/(\log x)^{5/2})$ in total.

For $\{3, 7, 11, 19\}$, Lemma 1.7 splits failure at 19 into finitely many maximal-support branches, each forbidding an explicit set of at least 9 of the 18 unit classes mod 19 on $n+4$ — dimension $\geq \frac{1}{2}$. Crossing with the branches at 11 leaves finitely many combined branches, each a five-form problem of dimension ≥ 3 (those involving the Type-II case at 11 have dimension $\geq 1 + \frac{1}{2} + \frac{1}{2} + \frac{3}{5} + \frac{1}{2} = \frac{31}{10} > 3$ and are again lower-order); the forbidden residues of the five forms are pairwise distinct for $q > 19$, smaller primes being absorbed into the congruence class of n . Summing the finitely many branches gives $O(x/(\log x)^3)$. The $\{3, 7, 11, 19, 23\}$ case adds the form $n+5$ with Lemma 1.7 at $R = 23$, one more factor of $(\log x)^{1/2}$. The full chain $\mathcal{P}$ is identical: the fifteen forms $n + (R-3)/4$, $R \in \mathcal{P}$, are distinct integer shifts of n (pairwise differences less than 27), each contributing dimension $\geq \frac{1}{2}$ through finitely many maximal-support branches, plus 1 for primality: total dimension $\geq 1 + \frac{15}{2} = \frac{17}{2}$, with forbidden residues pairwise distinct for $q > 107$ and smaller primes absorbed into the congruence class of n . $\square$

Proof of Proposition 1.9. $k = ap^2$ divides $m^2 = p^2a^2$ and $k \equiv -m$ if and only if $ap^2 \equiv -pa$, i.e. $p \equiv -1 \pmod{R}$; $k = a^2p$ divides m^2 and $a^2p \equiv -pa$ if and only if $a \equiv -1$, i.e. $p \equiv -4 \pmod{R}$ by (1.3). Integrality of b, c is automatic. Scanning all shapes $k = a^j p^i$ ($j, i \leq 2$) yields no further p -only condition. The failure of both families states that the two forms $(p+1)/2$ and $p+4$ (in n : $2n-1$ and $4n+1$) are free of primes $\equiv 3 \pmod{4}$ — two further sifting conditions of dimension $\frac{1}{2}$ each on forms distinct from all previous ones. Empirically the two families alone cover 74% of hard primes below 10^9 — and none of the four primes of Section 5 that admit a unique working residual ≤ 107 . $\square$

Remark 4.2. Each further prime residual $R \equiv 3 \pmod{4}$ whose support-bound lemma is verified adds $\frac{1}{2}$ to the exponent; the dynamic-program verification has now been carried through every prime residual up to 107. The chain is a concrete instantiation of Theorem 1.10 with explicit, verified constants. Two caveats frame its scope. First, the composite residuals $(15, 27, 35, \dots)$ are not covered by the present lemma; treating them requires the componentwise theory mod each prime-power factor. Second — and fundamentally — no fixed finite list can be pushed to *emptiness* by upper-bound sieving: the count bound $x/(\log x)^\kappa$ diverges for every fixed κ , and extrapolating the calibration of Section 5 to any fixed list yields a positive density constant, i.e. infinitely many expected exceptional primes with the record minimal residual growing slowly without bound. The correctly posed remaining problem is a slowly growing bound $R_{\min}(p) \ll f(p)$, which requires existence (lower-bound) technology beyond upper-bound sieves.

Proof of Theorem 1.10. Fix $A > 0$. Choose $B = B(A)$ so large that

$$\sum_{\substack{3 \leq R \leq B \\ R \equiv 3 \pmod{4}}} \frac{1}{\varphi(R)} \geq \sum_{\substack{3 \leq R \leq B \\ R \equiv 3 \pmod{4}}} \frac{1}{R} > A; \quad (4.1)$$

this is possible since the reciprocal sum over the progression $R \equiv 3 \pmod{4}$ diverges, and $\varphi(R) \leq R$. Let $S(A) = \{R \equiv 3 \pmod{4} : 3 \leq R \leq B\}$ and $M = 840 \cdot \prod_{R \in S(A)} R$.

Write $n = (p + 3)/4$ as before, so $(p + R)/4 = n + (R - 3)/4$ for each $R \in S(A)$: a family of $|S(A)|$ integer linear forms in n with pairwise distinct shifts, together with the form $4n - 3 = p$. Fix a residue class $n \equiv n_0 \pmod{M}$ compatible with p hard; there are finitely many, and it suffices to bound each. Within such a class, for every $R \in S(A)$ the quantities $t \equiv -4^{-1}p^2$ and $p^{-1} \pmod{R}$ are constant, so the class set $S_R(p) = \{t, tp^{-1}, tp^{-2}\} \pmod{R}$ of Proposition 2.2 is a fixed nonempty set of $t_R = t_R(n_0) \in \{1, 2, 3\}$ classes. By the contrapositive of Proposition 2.2, if every residual in $S(A)$ fails at p , then for each $R \in S(A)$ the form $n + (R - 3)/4$ has *no* prime factor in $S_R(p)$; and $4n - 3$, being prime and $> z$, has no prime factor $\leq z$.

Sift the class $\{n \leq (x + 3)/4 : n \equiv n_0 \pmod{M}\}$ by the primes $q \leq z = x^{1/10}$, $q \nmid M$. A prime q forbids: the class $n \equiv 3 \cdot 4^{-1} \pmod{q}$ (from the primality of $4n - 3$), and, for each $R \in S(A)$ with $q \pmod{R} \in S_R(p)$, the single class $n \equiv -\frac{R-3}{4} \pmod{q}$ (a prime factor q of the form $n + (R - 3)/4$ in a guaranteed-success class cannot occur). Thus

$$\omega(q) = 1 + \#\{R \in S(A) : q \pmod{R} \in S_R(p)\},$$

which is bounded by $|S(A)| + 1$; the forbidden residues of distinct forms are pairwise distinct for $q > B$, since the shifts $(R - 3)/4$ differ by less than B and $q \nmid M$. For each R , the density of primes q with $q \pmod{R} \in S_R(p)$ is $t_R/\varphi(R)$ by Dirichlet's theorem, so ω has average

$$\kappa = 1 + \sum_{R \in S(A)} \frac{t_R}{\varphi(R)} \geq 1 + \sum_{R \in S(A)} \frac{1}{\varphi(R)} > 1 + A,$$

in the Halberstam–Richert sense, by (4.1). The Fundamental Lemma [4, Thm. 2.5] in dimension κ then bounds the sifted count by $O_A\left(x \prod_{q \leq z} (1 - \omega(q)/q)\right) = O_A(x(\log x)^{-\kappa}) = O_A(x/(\log x)^{1+A})$, and summing over the finitely many classes $n_0 \pmod{M}$ preserves the bound. $\square$

Remark 4.3. Theorem 1.10 is the strongest statement the pure sieve can give: density $O((\log x)^{-A})$ for every A , but not emptiness. The conjecture for hard primes is equivalent to a *finite covering hypothesis*: some finite S admits, for every hard prime, a residual $R \in S$ with (1.2). The data to 10^{11} is covered by $S = \{3, 7, 11, \dots, 107\}$ — yet both the calibration (Remark 4.2) and the conditional construction of Section 6 predict that any fixed list eventually fails.

5 Computations to 10^{11}

All code and data are in the accompanying repository. Certificates are found by direct integer factorization of a (trial division; the factorization of $m^2 = p^2 a^2$ follows) and scanning divisors for the class condition (1.2); every certificate is re-verified in exact integer arithmetic, $4abc = p(bc + ac + ab)$.

5.1 Data

Every one of the 128,671,219 hard primes $p < 10^{11}$ possesses a residual certificate with $R \leq 107$. The minimal residuals distribute as follows (10^{11} data; Figure 1):

Salient features:

R	count	R	count	R	count
3	69,951,190	39	25,258	79	50
7	31,335,473	43	6,850	83	23
11	19,439,771	47	8,811	87	8
15	3,974,286	51	1,450	91	3
19	2,441,687	55	1,369	95	5
23	1,070,808	59	954	99	1
27	187,963	63	283	103	1
31	193,848	67	104	107	1
35	30,784	71	209		
		75	29		

Table 1: Distribution of the minimal residual $R_{\min}$ over the 128,671,219 hard primes below 10^{11} .

- $R_{\max} = 107$, attained by the single prime $8,803,369 < 10^7$, *unchanged from 5×10^7 to 10^{11}* . In the language of Corollary 2.3: its 27 shifted values contain only 43 distinct odd primes (typical: 45–52), and 81% of them satisfy $(\frac{p}{q}) = +1$ (typical: 44–60%) — a $\approx 4\sigma$ fluctuation in provably the right statistic.
- **A confirmed prediction.** Below 10^{10} no prime had minimal $R \in \{87, 91, 95, 99, 103\}$; the calibrated independence model attributed this gap to small-number statistics (expected counts ≤ 0.12 per value at 10^9) and predicted that deep-tail primes (failing every $R \leq 83$) land at $87, 91, 95, 99, 103, \geq 107$ with probabilities 0.53, 0.11, 0.22, 0.03, 0.05, 0.06 (the last being the total mass at or beyond 107). Between 10^{10} and 10^{11} the band populated with 8, 3, 5, 1, 1 primes respectively (all in $(1.3 \times 10^{10}, 10^{11})$), and none passed 103: the observed conditional distribution over the 19 deep-tail primes (8, 3, 5, 1, 1, 1 including the record) matches the prediction. The gap confirmed its own explanation by disappearing.
- Among all 1,587,581 hard primes below 10^9 , exactly *four* admit a unique working residual ≤ 107 (8,803,369 with 107; 142,361,209 with 59; 287,567,281 with 83; 794,037,841 with 63) — the entire obstruction to a shorter covering list. A greedy cover of the full solvability table uses 18 residuals; a packing argument gives a lower bound of 12.

5.2 Calibration of the base case of Theorem 1.8

The proportion of hard primes failing both $R = 3$ and $R = 7$, multiplied by $\log x$, is constant to within 2% across three decades of the 10^9 data:

bin ($p \sim$)	10^6	10^7	6×10^7	2.5×10^8	10^9
density $\times \log x$	5.14	5.19	5.18	5.17	5.16

Table 2: Joint $\{3, 7\}$ -exceptional density times $\log x$ across dyadic bins of the 10^9 data.

That is, the joint exceptional density is $\approx 5.17/\log x$: the sieve exponent of the base case of Theorem 1.8 is exactly right, and 5.17 estimates the (conjectural) Selberg–Delange constant. The exceptional set is density zero but decays slowly: 26.5% of hard primes below 10^9 fail both




Figure 1: Distribution of the minimal residual $R_{\min}$ over the 128,671,219 hard primes below 10^{11} (logarithmic count scale). The head is dominated by $R = 3, 7, 11$; the tail thins geometrically; the record $R = 107$ is attained by a single prime below 10^7 .

residuals; all are settled by deeper residuals of the list. The fitted per-residual decay exponents likewise match the sieve prediction $\kappa = \frac{1}{2}$ where exact criteria exist (e.g. $\hat{\kappa} = 0.539$ for $R = 7$).

5.3 Failure taxonomy

Sampling the complete solvability table at 10^9 : failures at $R = 7$ are 100% Type I (as Theorem 1.3 requires); at $R = 11$, 89.7% Type I and 10.3% the budget pattern of Theorem 1.4(b); the Type-I share broadly decreases with $\varphi(R)$ (e.g. 51% at $R = 103$, though not monotonically), as hitting one specific class among $\varphi(R)$ grows harder while the character dichotomy stays binary. Taken together, the data of this section is the empirical face of Sections 2–4: every distributional feature — the head, the geometric tail, the once-empty band, the static record — now has a structural interpretation attached.

6 Open problems

1. **Finite covering.** Prove that some finite admissible set S covers all hard primes (Theorem 1.10 gives density $(\log x)^{-A}$ for every A ; emptiness is beyond the sieve).
2. **Beyond the prime chain.** The support-bound lemma is now verified for every prime residual up to 107 (exponent $\frac{17}{2}$); extending past 107 meets the growth of the DP state space ($\approx 2.8\times$ per residual), and the composite residuals require the componentwise theory. More importantly, the fixed-list format itself cannot reach emptiness (Remark 4.2): the structurally right target is a slowly growing bound on $R_{\min}(p)$.
3. **The constant 5.17.** The joint $\{3, 7\}$ -exceptional density is measured as $\approx 5.17/\log x$, stable to 2% across three decades. Conjecturally 5.17 is a Selberg–Delange-type constant: the count should satisfy $E(x) \sim Cx/(\log x)^2$ with $C = c_0 \prod_q \delta_q$, the local factors δ_q

recording, for each prime q , the probability that q divides neither form in a forbidden class. Proving the asymptotic (not just the upper bound) requires a two-form Selberg–Delange or shifted-convolution analysis — open, and a natural test problem for such methods (see also Remark 1.11 for the density context).

4. **Structural reformulations and complexity.** Theorem 1.6 invites a graph formulation: on the bipartite graph joining primes q to residuals R by $q \mid (p + R)/4$, the edge characters $(\frac{q}{R})$ collapse to vertex signs $(\frac{p}{q})$, and joint Type-I failure is an all-plus vertex assignment — a percolation-style object whose structure could organize the growth law. Separately, the meta-theorem’s enumeration complexity is unresolved: the DP state count grows empirically by $\approx 2.8\times$ per successive prime residual (271 at $R = 19$ to 1.95×10^7 at $R = 107$); is the equivalence-state complexity genuinely exponential in R , and is there a polynomial-size certificate for the support-bound property?
5. **The record and the covering hypothesis.** Theorem 1.6 explains the record prime quantitatively (see Corollary 2.3) and yields a conditional resolution of the covering question: under Dickson’s prime-tuple conjecture, for every B there are infinitely many hard primes failing every residual $\leq B$ — fix the forced smooth parts of the shifted values (character-neutral by Corollary 2.3(i)), make the cofactors prime in prescribed residue classes, and apply Proposition 2.1. So no fixed finite list suffices, conditionally, and the record grows without bound (at the pure-Type-I level, like $\log x / \log \log x$). Unconditional versions appear out of reach: already “infinitely many hard p fail $R = 3$ ” contains a two-form prime-tuple statement. What remains open is an unconditional upper bound on $R_{\min}(p)$ — and this is not an incremental item: by the completeness observation of Section 1.1, $R_{\min}(p) \leq 2p$ is *equivalent* to the conjecture at p , so an unconditional bound $R_{\min}(p) \leq f(p)$ for any f whatsoever is the full problem.

A Verification details

This appendix documents every computer-assisted claim of the paper (their proof status is stated in Remark 1.5): the finite enumerations that are *proof components* (Appendix A.1), the large-scale certificate computations behind Section 5 and how each dataset is verified (Appendix A.2), and the commands that reproduce all of it (Appendix A.3). All code is in the accompanying repository; every check is deterministic and uses exact integer arithmetic only.

A.1 Proof components

Each component below is an exhaustive enumeration of a finite space whose sufficiency is established in the text (Theorem 1.2 for the exact criteria; for Lemma 1.7, monotonicity of the reachability of the target in support and multiplicities, checked at minimal exponents). Each run completes in seconds on commodity hardware, except the largest support-bound check ($R = 107$: 331 s); the entry points named are in the module `erdos_straus.theory`.

Trusted computing base. The certificate data does not rest on the solver: every certificate claim is re-checked by evaluating the integer identity $4abc = p(bc + ca + ab)$ directly, so the trusted base for the data is a bignum multiply-and-compare (the host language’s built-in integers), not the search pipeline. The finite case analyses trust more — the enumeration code itself, roughly

sixty lines per check. Four mitigations apply. First, the enumerated spaces are defined by statements proved in the text (Theorem 1.2, Lemma 4.1), and the code mirrors those definitions one-to-one. Second, independent implementations agree where they overlap: the subset dynamic program and the direct binomial enumeration verify identical claims at $R = 19, 23$, and the exact criteria are cross-checked against ground-truth divisor search on up to 1,587,581 primes with zero discrepancies. Third, every check is deterministic, uses exact integer arithmetic only (no floating point, randomness, or external libraries), and completes in seconds from the repository — an independent re-implementation is an afternoon’s work, and we would welcome one.

Fourth, the elementary layer and the enumeration layer are formalized in Lean 4/mathlib in the accompanying repository (`lean/`). The elementary layer covers certificate soundness and integrality (`certificate_sound`, `certificate_integrality`), both directions of Theorem 1.1 (`theoremA_necessity`, `theoremA_sufficiency`), both families of Proposition 1.9 (`family_p_plus_one`, `family_p_plus_four`, as integer identities), and Theorem 1.6 (`reciprocity_structure`). For the enumeration layer, the reachability model underlying Theorem 1.2 is defined in the multiplicative structure of $\mathbb{Z}/R$ directly, both halves of Lemma 4.1 are proved as theorems (`reach_mono`, `reach_sublist`) with no computational content to trust, and the full enumerations behind Theorem 1.3 (1,536 configurations, checked as the 384 that remain after the neutral class collapses), Theorem 1.4 (497,664 capped configurations, including the exact two-case failure classification), and the $R = 19, 23$ cases of Lemma 1.7 (437,580 resp. 7,759,752 checks) are discharged by Lean’s compiled evaluator (`native_decide`, which trusts the Lean compiler in addition to the kernel). The three larger enumerations are *evaluated* in the discrete-log coordinates in which the proofs and `verify_support_bound` work, but the coordinates add nothing to the trusted base: the translation is itself a proved theorem (the mask semantics — rotation is multiplication by g^s , the mask fold is the divisor-class reachability set), and the formal development derives the $R = 11, 19, 23$ statements back in the multiplicative model, where each reports exactly its enumeration’s single evaluator axiom. The reduction of Theorem 1.2 is also formalized (`reach_certificate`): membership of the target class in the reachability model is equivalent to the existence of a divisor $k \mid m^2$ in that class — via the fundamental theorem of arithmetic — and yields the explicit positive certificate (k, b, c) with the identity of (1.1); this closes the loop between the abstract model that the enumerations check and the integer certificates the conjecture asks for, with no computational content to trust.

Two further layers are formalized on top. First, the composed corollary: if $(p+19)/4$ has nine prime factors in pairwise-distinct unit classes other than 1 modulo 19, an explicit solution of (1.1) at p exists — Lemma 1.7 applied end to end, machine-checked through the enumeration, both bridges, and the budget-consolidation lemmas. Second, the $R = 31$ case of Lemma 1.7 is verified by *certified dynamic programming*: the $\binom{29}{15} = 77,558,760$ supports are never enumerated — the formal development implements the subset dynamic program of `verify_support_bound_dp` in Lean, a machine-checked soundness induction shows it covers every support, and the single step delegated to the evaluator is the final check that every state realized by ≥ 15 classes hits all 30 targets (3,001 states in the run). Every main theorem is audited with `#print axioms`, and helper lemmas are covered transitively through them. For these enumerations the sixty-line trust base above is thus replaced by the Lean toolchain evaluating a formally stated proposition — and the Python and Lean implementations, written independently against the same statements, agree.

In the development notes and Lean sources accompanying this paper, several results carry working letter names; the dictionary is: Theorem A = Theorem 1.1, A' = Theorem 1.3, A'' =

Theorem 1.4, Theorem J = Theorem 1.6, Theorem I = Proposition 1.9, Lemma S = Lemma 1.7, and the monotonicity reduction = Lemma 4.1.

component	configuration space	count	entry point
Thm. 1.3 ($R = 7$)	class multisets in $(\mathbb{Z}/7)^*$, multiplicities ≤ 3 , with the forced even factor, $p \bmod 7 \in \{1, 2, 4\}$, and the consistency relation	1,536	<code>verify_R7_finite</code>
Thm. 1.4 ($R = 11$)	dynamic program over $(\mathbb{Z}/11)^*$ with the forced factor 3; collapses to 25 equivalence states (16 succeed, 6 fail by character, 3 by budget)	25	<code>finite_criterion_dp</code>
Lem. 1.7 ($R = 19$)	all supports of size 9 over the nonzero logs $\times$ all classes of p , at minimal multiplicities	437,580	<code>verify_support_bound</code>
Lem. 1.7 ($R = 23$)	all supports of size 11, likewise	7,759,752	<code>verify_support_bound</code>
Lem. 1.7 ($R = 19$ –107, all twelve primes)	subset dynamic program over reachability-mask states; state counts 271 ($R=19$), 3,001 ($R=31$), 2,414,145 ($R=83$), 19,530,971 ($R=107$); zero violations throughout	$\leq 19.5\text{M}$	<code>verify_support_bound_dp</code>

Table 3: Machine-verified proof components: enumerated configuration spaces and the entry points in `erdos_straus.theory`.

In addition (as sanity checks, not proof components), the criteria of Theorems 1.1, 1.3 and 1.4 were compared against ground-truth divisor search on the complete solvability data below 10^9 : agreement in 1,587,581/1,587,581 cases for $R = 3$, 41,421/41,421 sampled failure cases for $R = 7$ (all Type-I, as Theorem 1.3 requires), and 158,759/158,759 sampled cases for $R = 11$, with zero discrepancies. The identity of Theorem 1.6 and its forced-prime corollary were likewise checked on 135,588 resp. 61,438 instances from the same data, with zero violations (Corollary 2.3).

A.2 Certificate datasets and their verification

The computational claims of Section 5 rest on the following datasets, all archived in the repository under `data/`. A certificate (R, a, b, c) for p is *accepted* only if $4abc = p(bc + ac + ab)$ holds in exact integer arithmetic and $R = 4a - p$; for the compact minimal- R maps, the verifier *reconstructs* the triple from (p, R) by the divisor search of (1.2) and then applies the same exact check, so a map entry is verified as a genuine certificate, not taken on faith.

The masks dataset of Table 4 is the basis for the covering-set statements (the 18-residual cover, the lower bound 12, the four critical primes), the failure taxonomy, and the calibration of Section 5; the minimal- R maps are the basis for Table 1, the record, and the gap analysis.

dataset	entries	verification performed
explicit triples, $p < 1.2 \times 10^8$	213,131	every triple checked exactly
minimal- R map, $p < 10^9$	1,587,581	<i>exhaustive</i> : every entry reconstructed and checked exactly
minimal- R map, $p < 10^{10}$	14,215,707	143,000 systematically sampled entries reconstructed and checked; additionally <i>all</i> 842 tail entries ($R \geq 51$) checked with full minimality (no smaller admissible R succeeds)
solvability masks, $p < 10^9$	$27 \times 1,587,581$	per-prime solvability of every admissible $R \leq 107$ by full divisor-class computation; consistency with the minimal- R map
R -sequence dataset, $p < 10^{11}$	128,671,219	257,343 systematically sampled entries reconstructed and checked; <i>all</i> 20,151 tail entries ($R \geq 43$) checked with full minimality; prime array pinned by sha256; pipeline validated byte-for-byte against the exhaustively verified 10^9 map

Table 4: Certificate datasets and the verification applied to each.

A.3 Reproduction

With the repository installed (`pip install -e .`), the following commands regenerate and re-verify everything above. Generation of the full 10^{10} map takes about 11 minutes on four cores; every other command runs in seconds to minutes.

The full test suite (`python -m pytest`; 46 tests) exercises the proof-component checks, the exact criteria against ground truth, and sampled dataset validation on every run.

References

- [1] C. Elsholtz and T. Tao, *Counting the number of solutions to the Erdős–Straus equation on unit fractions*, J. Aust. Math. Soc. **94** (2013), 50–105.
- [2] P. Erdős, *Az $1/x_1 + 1/x_2 + \dots + 1/x_n = a/b$ egyenlet egész számú megoldásairól* [On the integer solutions of the equation $1/x_1 + 1/x_2 + \dots + 1/x_n = a/b$] (Hungarian), Mat. Lapok **1** (1950), 192–210.
- [3] J. Friedlander and H. Iwaniec, *Opera de Cribro*, American Mathematical Society Colloquium Publications, vol. 57, American Mathematical Society, Providence, RI, 2010.
- [4] H. Halberstam and H.-E. Richert, *Sieve Methods*, Academic Press, London, 1974.
- [5] S. Mihnea and D. C. Bogdan, *Further verification and empirical evidence for the Erdős–Straus conjecture*, preprint, arXiv:2509.00128 [math.NT], 2025.
- [6] L. J. Mordell, *Diophantine Equations*, Academic Press, London, 1969.

claim	command
Thm. 1.3 finite check	<code>python -c 'from erdos_straus.theory \</code> <code>import verify_R7_finite; \</code> <code>print(verify_R7_finite())'</code>
Thm. 1.4 enumeration	<code>finite_criterion_dp(11, ...)</code> as in the test suite (tests/test_solver.py)
Lem. 1.7 at $R = 19, 23$	<code>verify_support_bound(19), (23)</code>
generate the 10^{10} map	<code>python -m erdos_straus.bulk_generate \</code> <code>-max 10000000000 -segmented -store rmap</code>
verify a dataset exhaustively	<code>python -m erdos_straus.verify data/<file></code>
distribution / covering / taxonomy	<code>python -m erdos_straus.analyze dist (also</code> <code>cover, tail); python -m erdos_straus.theory</code>
Lean formalization (elementary + enumeration layers)	<code>cd lean/ErdosStraus</code> <code>lake exe cache get && lake build</code>

Table 5: Reproduction commands (one shell command per line; \ marks an explicit line continuation within a single command).

- [7] S. E. Salez, *The Erdős–Straus conjecture: new modular equations and checking up to $N = 10^{17}$* , preprint, arXiv:1406.6307 [math.NT], 2014.
- [8] A. Swett, *The Erdős–Straus conjecture*, online verification record to 10^{14} , 1999; formerly at <http://math.uindy.edu/swett/esc.htm> (archived snapshots available via the Internet Archive; accessed August 2026).
- [9] R. C. Vaughan, *On a problem of Erdős, Straus and Schinzel*, *Mathematika* **17** (1970), 193–198.