

Residual certificates for the Erdős–Straus conjecture: exact solvability criteria, almost-all certificate bounds, and computations to 10^{11}

J. M. Hopkins*

August 2026

Abstract

The Erdős–Straus conjecture asserts that $4/n = 1/a + 1/b + 1/c$ is solvable in positive integers for every $n \geq 2$; the open cases concentrate on primes in six residue classes mod 840 (*hard primes*). We develop the *residual method*: a solution with first denominator $a = (p + R)/4$ exists if and only if some divisor $k \mid m^2$, $m = pa$, satisfies $k \equiv -m \pmod{R}$, and every solution arises this way — so the conjecture at p is exactly a bound on the minimal working residual $R_{\min}(p)$. The paper’s central theorem is a *support theorem*: failure at R forbids at least half the unit classes mod R on the shifted form $(p + R)/4$, proved for every residual at once via Kneser’s addition theorem. Its uniform consequence bounds $R_{\min}$ itself: all but $O(x \exp(-c(\log \log x)^2))$ hard primes $p \leq x$ satisfy $R_{\min}(p) \leq \varepsilon \log \log x$, with explicit certificates. This is not a sharper exceptional-set estimate — Vaughan’s classical mean-value bound is denser, and remains the record on that axis — but a statement about a different quantity, one that mean-value methods do not reach: they count the primes that fail without exhibiting a denominator for those that succeed. Around these sit exact solvability criteria at $R = 3, 7, 11, 15$, a reciprocity identity $(\frac{q}{R}) = (\frac{R}{q})$ governing correlated failure, a branch-classification conjecture that would raise the almost-all bound to $R_{\min}(p) \leq (\log x)^{c_0}$, and a measured ladder program past that ceiling. Identifying the failure families as norm-form families makes the count at the chain’s first link sharp, $\asymp x/(\log x)^{3/2}$. The elementary layer and the core finite enumerations are machine-checked in Lean 4/mathlib, and we exhibit certificates for all 128,671,219 hard primes below 10^{11} ; the record minimal residual is 107.

2020 Mathematics Subject Classification. Primary 11D68; Secondary 11N35, 11Y50, 68V20.

Keywords. Erdős–Straus conjecture, Egyptian fractions, unit fractions, sieve methods, formal verification.

1 Introduction

Erdős and Straus conjectured (see [5]) that for every integer $n \geq 2$ the equation

$$\frac{4}{n} = \frac{1}{a} + \frac{1}{b} + \frac{1}{c} \tag{1.1}$$

*Independent researcher. jeff.m.hopkins@gmail.com. Computations and computer-assisted case analyses were carried out with the code repository accompanying this paper: <https://github.com/jeffmhopkins/Erd-s-Straus-attack>.

has a solution in positive integers a, b, c . By multiplicativity it suffices to treat prime n , and classical polynomial identities settle every prime outside the six residue classes

$$\mathcal{H} \bmod 840 = \{1, 121, 169, 289, 361, 529\},$$

the squares of units modulo 840 (Mordell [15, Ch. 30]; the quadratic-residue nature of the obstruction goes back to Yamamoto [22]). We call primes in these classes *hard primes*. Direct verification has confirmed the conjecture far beyond any range considered here: to 10^{14} by Swett [18], to 10^{17} by Salez [16], and to 10^{18} by Mihnea and Bogdan [14]; the present paper is logically independent of those bounds. The structure of the solution count was analyzed by Elsholtz and Tao [4]. Verification, however, carries no structural information; the aim here is a quantitative structural reduction. For calibration: the first density bounds on the exceptional set are due to Webb [21], sieve bounds were advanced by Sander [17], and Vaughan’s classical mean-value bound [20] gives an exceptional set $\ll x \exp(-c(\log x)^{2/3})$ for (1.1) — asymptotically stronger than any fixed power of $\log x$, and still the record on that axis.

We are not competing on it. The distinction is worth stating plainly, because it determines what the results below are for. A mean-value bound counts the integers that *fail*; it exhibits no denominator for the ones that succeed, and says nothing about how large a first denominator a given prime requires. By Proposition 1.1 the conjecture at p is precisely the assertion that the least admissible residual $R_{\min}(p)$ is finite — indeed $\leq 2p$ — so $R_{\min}$, not the exceptional set, is the quantity the problem is about. The theorems here bound $R_{\min}$: for every hard prime through 10^{11} by explicit certificate, and for almost every hard prime by Theorem 4.5, at a residual of size $O(\log \log x)$. Density statements follow as corollaries, and are weaker than Vaughan’s; the bounds on $R_{\min}$, the exact local criteria that produce them, and the certificates they exhibit have no counterpart in the mean-value literature. Vaughan’s theorem is the benchmark against which we calibrate the sieve, not the target we are trying to overtake.

New contributions

Main theorems.

- The *residual formulation* of the Erdős–Straus equation: solvability with first denominator $(p + R)/4$ is a divisor-class condition, complete over all solutions (Proposition 1.1).
- *Exact* solvability criteria at $R = 3, 7, 11$ (Theorems 1.2, 1.4, 1.5) and at the first composite residual $R = 15$ (Theorem 1.6: a Jacobi-character dichotomy with no budget cases), and a meta-theorem (Theorem 1.3) showing every fixed residual — prime or composite — admits a finite computable exact criterion.
- A reciprocity principle (Theorem 1.8) collapsing the per-residual conditions into one global statement about p — which we read as the organizing principle behind correlated failure and the record prime.
- The per-residual sieve input proved unconditionally for *every* residual at once (Theorem 1.10): via Kneser’s addition theorem, any failing configuration occupies at most $(R - 3)/2$ nonzero factor classes, with composite residuals covered through the general abelian form; the per-residual machine verification to $R = 107$ (Lemma 1.9) serves as its independent confirmation.

- A constructive sieve chain built on these criteria (Theorem 1.11, Proposition 1.12): explicit, certificate-producing exceptional-set bounds of exponent $1 + |\mathcal{P}|/2$ for every finite admissible list $\mathcal{P}$ of primes and of composites up to 107, with marginally smaller exponents for larger composites with $\omega(R) \leq 2$ (29/2 on the 27 residuals up to 107; 31/2 with the two identity families of Proposition 1.12 adjoined); an unconditional density reduction (Theorem 1.13): for every A , a finite residual set whose joint failure set is $O_A(x/(\log x)^{1+A})$; and the *uniform chain* (Theorem 4.5), the strongest unconditional statement here: all but $O(xe^{-c(\log \log x)^2})$ hard primes $p \leq x$ have $R_{\min}(p) \leq \varepsilon \log \log x$ — together with the branch-classification program behind its ceiling (Section 4.2): the kernel case proved (Theorem 4.8), the count form refuted (Theorem 4.7), and a heavily verified container conjecture (Conjecture 4.9) whose resolution raises the bound to $R_{\min} \leq (\log x)^{c_0}$ (Theorem 4.11).

Further contributions.

- Computed and verified residual certificates for all 128,671,219 hard primes below 10^{11} , with a calibrated model whose deep-tail predictions were subsequently confirmed (Section 6).
- The *Burgess–reciprocity ladder* (Section 5): a p -adapted residual family on which the all-residue failure mode is impossible by construction; the ladder chain and its first rung proved (Theorems 5.3, 5.4), the avoidance ladder two-sided (Theorem 5.5), and an almost-all bound past every fixed power of $\log x$ conditional on one measured hypothesis (Theorem 5.11).
- The identification of Type-I failure families as *norm-form* families (5.1), which makes Iwaniec’s prime-indexed half-dimensional sieve [11] the natural counting tool and yields matching upper and lower bounds at the first link of the chain (Theorem 5.7: the failure count at $R = 3, 7, 15$ is $\asymp x/(\log x)^{3/2}$, and $3/2$ is the ceiling of the method); along the ladder the same mechanism applies on a sparse family of classes (Theorem 5.8 and Remark 5.9).
- A machine-checked (Lean 4 [3] with mathlib [13]) formalization, in the accompanying repository, of the elementary layer of the proofs, of the finite enumerations behind Theorems 1.4 and 1.5 and Lemma 1.9 at $R = 19, 23, 31$, and of the reduction of Theorem 1.3 itself. The complete inventory, and what each layer contributes to the trusted base, is in Appendix A.

Conceptually, the paper has three layers: (1) an exact reformulation of solvability — every solution of (1.1) at a hard prime is a residual certificate (Proposition 1.1); (2) finite, machine-checkable inputs for individual residuals — the exact criteria (Theorems 1.2–1.6), the reciprocity identity (Theorem 1.8), and the support bounds (Lemma 1.9, Theorem 1.10); and (3) sieve consequences obtained by combining those inputs (Theorems 1.11–1.13), capped by the p -adapted ladder program of Section 5. The computations of Section 6 calibrate layer (3) and test its predictions; the formal development machine-checks the core of layers (1)–(2). Within this structure one result carries the technical weight, and is worth naming at the outset: Theorem 1.10, the support theorem — a single per-residual input, proved for all residuals at once by Kneser’s addition theorem, with no computation and no restriction on R . Everything downstream is what it buys: Theorem 4.5’s bound on $R_{\min}$ for almost every hard prime, the chain of exceptional-set bounds, and the conditional program of Section 5. The exact criteria supply its small-residual instances, and the computations calibrate what it predicts.

1.1 The residual method

Fix a hard prime p . All six classes satisfy

$$(H1) \ p \equiv 1 \pmod{8}, \quad (H2) \ p \equiv 1 \pmod{3}, \quad (H3) \ p \bmod 7 \in \{1, 2, 4\}.$$

For an *admissible residual* $R \equiv 3 \pmod{4}$ put

$$a = \frac{p+R}{4} \in \mathbb{Z}, \quad m = pa.$$

Substituting a into (1.1) and clearing denominators shows that a solution with first denominator a exists if and only if

$$\exists k \mid m^2 \quad \text{with} \quad k \equiv -m \pmod{R}, \quad (1.2)$$

in which case $b = (k+m)/R$ and $c = (m^2/k+m)/R$ are positive integers solving (1.1). We write $R_{\min}(p)$ for the least admissible R for which (1.2) holds. The formulation loses no solutions:

Proposition 1.1 (completeness). *Every solution of (1.1) at a hard prime p , ordered $a \leq b \leq c$, has $p/4 < a \leq 3p/4$; hence $R = 4a - p$ is an admissible residual in $(0, 2p]$ and the solution is a residual certificate at R . Consequently $R_{\min}(p) \leq 2p$ if and only if the conjecture holds at p .*

Proof. With $a \leq b \leq c$, the equation gives $1/a < 4/p \leq 3/a$, so $p/4 < a \leq 3p/4$ and $R = 4a - p \in (0, 2p]$; $R \equiv 3 \pmod{4}$ is automatic from $p \equiv 1 \pmod{4}$. By the displayed equivalence, the solution yields a divisor k as in (1.2), so $R_{\min}(p) \leq R \leq 2p$. Conversely (1.2) at any admissible R produces a solution. $\square$

A word on provenance. The divisor-congruence identity underlying (1.2) is classical: criteria of this shape are implicit in Mordell's treatment [15, Ch. 30] and in Yamamoto's analysis of the quadratic obstruction [22]. What we claim as new is not the identity but what is built on it: the completeness statement of Proposition 1.1 (every solution is a residual certificate, so the conjecture *is* a bound on $R_{\min}$), the residual-indexed program — exactness at fixed R , reciprocity across residuals, the support theorem uniform in R — and the sieve consequences those inputs support.

Two elementary congruences drive everything below. From $4a = p + R$,

$$a \equiv 4^{-1}p, \quad m \equiv 4^{-1}p^2 \pmod{R}, \quad (1.3)$$

and since $\gcd(p, R) = 1$ (as p is prime and $R \neq p$, $R \equiv 3 \not\equiv p \pmod{4}$), also $\gcd(a, R) = \gcd(m, R) = 1$: the setting is always a unit computation modulo R . By (1.3), m is a *square* modulo every prime $r \mid R$; hence for every prime $r \mid R$ with $r \equiv 3 \pmod{4}$ (at least one exists, since $R \equiv 3 \pmod{4}$),

$$\left(\frac{-m}{r}\right) = \left(\frac{-1}{r}\right) = -1 \quad (1.4)$$

— the target class $-m$ is a quadratic non-residue mod r .

A worked example. Take the hard prime $p = 1009 \pmod{840}$ and the smallest admissible residual $R = 3$. Then $a = (1009 + 3)/4 = 253 = 11 \cdot 23$ and $m = pa = 255,277$. Since $11 \equiv 2 \pmod{3}$, Theorem 1.2 below promises a certificate: take $k = 11$, a divisor of m^2 with $k \equiv -m \equiv 2 \pmod{3}$. Then $b = (k + m)/3 = 85,096$ and $c = (m^2/k + m)/3 = 1,974,822,872$, and indeed

$$\frac{4}{1009} = \frac{1}{253} + \frac{1}{85,096} + \frac{1}{1,974,822,872}.$$

Had every prime factor of a been $\equiv 1 \pmod{3}$, every divisor of m^2 would be $\equiv 1 \pmod{3}$ and the required class 2 would be unreachable — the prototype of all the obstructions studied below.

The shape of what follows. Everything in this paper grows out of the picture just seen, read at three levels. *Locally* (one residual): whether the target class $-m$ is attainable from the factor classes of a is a question about a finite multiplicative state — the multiset of classes with bounded exponent budgets — so each fixed R admits an exact, machine-enumerable solvability criterion (Theorems 1.2–1.5 below). *Globally* (all residuals at once): the conditions entering these criteria are quadratic characters, and reciprocity folds the condition at every residual into signs $(\frac{p}{q})$ indexed by the primes q dividing the shifted values $(p+R)/4$ — failures at different residuals correlate because shifted values share prime factors (Theorem 1.8). *Quantitatively*: each criterion says that failure forces $(p+R)/4$ to avoid prime factors in many residue classes, a sifting condition; stacking residuals multiplies the savings (Theorem 1.11), though it provably cannot reach emptiness (Remark 4.3) — the honest endpoint of the method, discussed in Section 7.

1.2 Main results

Theorem 1.2 ($R = 3$; exact). *Let $p \equiv 1 \pmod{4}$ be prime with $3 \nmid p$, and $a = (p+3)/4$. Then (1.2) holds for $R = 3$ if and only if a has a prime factor $q \equiv 2 \pmod{3}$. When it does, $k = q$ is a certificate.*

Modulo 3, character equals class and the criterion is elementary. For every other fixed residual, the reduction to a finite check is itself a theorem, and the next two results are its first showcase instances.

Theorem 1.3 (meta-theorem). *For every fixed prime $R \equiv 3 \pmod{4}$, the validity of (1.2) at a hard prime p depends only on (i) the multiset of classes mod R of the prime factors of $a = (p+R)/4$, with multiplicities capped at $\lceil (R-1)/2 \rceil$, and (ii) $p \pmod{R}$. Consequently an exact solvability criterion exists for every R and is computable by finite enumeration.*

Theorem 1.4 ($R = 7$, hard primes; exact). *Let p be a hard prime and $a = (p+7)/4$. Then (1.2) holds for $R = 7$ if and only if a has a prime factor that is a quadratic non-residue modulo 7.*

Theorem 1.5 ($R = 11$, hard primes; exact). *Let p be a hard prime and $a = (p+11)/4$; note $3 \mid a$, $2 \nmid a$, $7 \nmid a$ automatically. Then (1.2) fails for $R = 11$ if and only if either*

- (a) *every prime factor of a is a quadratic residue mod 11; or*
- (b) *$v_3(a) = 1$, every other prime factor of a is $\equiv 1 \pmod{11}$ except a non-residue part — its factors in the non-residue classes mod 11 — of exactly one of the shapes*

$$\{q \equiv 2\}, \quad \{q \equiv 6\}, \quad \{q \equiv 2\} \cup \{q' \equiv 6\} \pmod{11},$$

each with multiplicity one, paired respectively with $p \equiv 2, 6, 1 \pmod{11}$.

For composite R the certificate condition couples the prime-power components of R through the Chinese Remainder Theorem, and the divisor-class computation lives in the (generally non-cyclic) unit group $(\mathbb{Z}/R)^*$. Theorem 1.3 holds verbatim with $(\mathbb{Z}/R)^*$ in place of the cyclic group (multiplicities capped once exponent ranges cover each class's cyclic span; see Remark 3.1). The first composite case turns out to be *simpler* than $R = 11$: a pure character dichotomy, with the quadratic character modulo a prime replaced by the Jacobi character modulo 15, and no budget cases at all.

Theorem 1.6 ($R = 15$, hard primes; exact). *Let p be a hard prime and $a = (p + 15)/4$; note $p \bmod 15 \in \{1, 4\}$ and $2 \mid a$ automatically. Then (1.2) holds for $R = 15$ if and only if a has a prime factor q with $(\frac{q}{15}) = -1$, i.e. $q \equiv 7, 11, 13, 14 \pmod{15}$.*

Remark 1.7 (computer-assisted components). The exhaustive verifications behind Theorems 1.4, 1.5, 1.6 and Lemma 1.9 below are *parts of the respective proofs*, on the same footing as a hand case analysis. Each is a deterministic enumeration of a finite configuration space whose finiteness and sufficiency are proved in advance (Theorem 1.3; for Lemma 1.9, monotonicity of target-reachability in support and multiplicities, checked at minimal exponents). The checks are short (seconds each, up to 331 s for the largest, $R = 107$), involve only exact integer arithmetic, and the code is available in the accompanying repository; Table 5 records the exact configuration counts and entry points. The enumerations at $R = 7, 11$ and the $R = 19, 23, 31$ cases of Lemma 1.9 are additionally formalized in Lean 4, with the reductions they rely on proved as theorems rather than trusted (the $R = 15$ enumeration and the remaining cases of Lemma 1.9 run in the Python layer only); Appendix A gives the inventory and the exact trust base.

The obstruction framework of Section 2 yields a reciprocity identity that converts the per-residual conditions into a single global statement about p — a structural interpretation of joint failure across residuals:

Theorem 1.8 (reciprocity structure). *Let p be a hard prime, $R \equiv 3 \pmod{4}$ prime, and q an odd prime factor of $a = (p + R)/4$ (necessarily $q \neq R$). Then*

$$\left(\frac{q}{R}\right) = \left(\frac{p}{q}\right).$$

The sieve consequences below need one input per residual: a statement that failure at R forces the shifted form $(p + R)/4$ to avoid prime factors in a positive-density set of residue classes. For $R = 3, 7, 11, 15$ the exact criteria provide this; for every other residual it is the support bound, which we state first in the per-residual machine-verified form in which it was originally established (Lemma 1.9) and then prove unconditionally for every residual at once (Theorem 1.10; both proved in Section 4).

Lemma 1.9 (support bound; verified for every prime $R \equiv 3 \pmod{4}$, $19 \leq R \leq 107$). *Let R be one of the twelve primes 19, 23, 31, 43, 47, 59, 67, 71, 79, 83, 103, 107, and call the set of unit classes mod R occupied by the prime factors of $(p + R)/4$ its support. Every configuration for which residual R fails at a hard prime has support of size at most $(R - 1)/2$, at most $(R - 3)/2$ of whose classes differ from the identity class (which never affects solvability). Consequently failure at R forbids at least half the classes, through finitely many maximal-support branches.*

Theorem 1.10 (unconditional support bound). (i) *Let $d \geq 4$ be even and $S \subseteq \mathbb{Z}/d \setminus \{0\}$. If the bounded subset-sum set*

$$M(S) = \left\{ \sum_{v \in S} e_v v : e_v \in \{0, 1, 2\} \right\}$$

is not all of $\mathbb{Z}/d$, then $|S| \leq d/2 - 1$, and the bound is attained ($S =$ the nonzero even classes). (ii) *Consequently, for every prime $R \equiv 3 \pmod{4}$ and every p : if the prime factors of $(p + R)/4$ occupy at least $(R - 1)/2$ nonzero unit classes mod R , then every unit class is a divisor class of m^2 and (1.2) holds; failure forces the conclusion of Lemma 1.9, with no upper limit on R .* (iii) *In a finite abelian group G of even order g with t involutions (multiplicative form, $M(S) = \prod_{v \in S} \{1, v, v^2\}$): if $M(S) \neq G$ then $|S| \leq \max(\lfloor (g + t - 2)/2 \rfloor, g/2 - 1)$. For $G =$*

$(\mathbb{Z}/R)^*$ with $\omega(R) \leq 2$ this is $\leq g/2$, so failure at any admissible composite R forbids at least $\varphi(R)/2 - 1$ unit classes on $(p+R)/4$ — sieve dimension $\geq \frac{1}{2} - 1/\varphi(R)$. For the admissible composites $R \leq 107$ an exhaustive enumeration (`kneser_support_general`, a proof component of this clause) sharpens the maximal failing support to $g/2 - 1$, so those residuals forbid at least $\varphi(R)/2$ classes and join the chain with full dimension $\geq \frac{1}{2}$ each. For arbitrary $\omega(R)$ the same bound reads $|S| \leq \max(\lfloor (g+t-2)/2 \rfloor, g/2 - 1)$ with $t = 2^{\omega(R)} - 1$, so failure forbids at least $(g-t)/2 - 1 \geq \varphi(R)/2 - 2^{\omega(R)-1} - 1$ unit classes — dimension $\geq \frac{1}{2} - (2^{\omega(R)-1} + 1)/\varphi(R)$, and since $2^{\omega(R)} = R^{o(1)}$ this is ≥ 0.44 for every admissible R beyond an absolute bound: every residual, with no restriction on ω , is a usable sieve input.

Theorem 1.11 (the chain). *The number of hard primes $p \leq x$ failing every residual in $\{3, 7\}$, in $\{3, 7, 11\}$, in $\{3, 7, 11, 15\}$, and in $\{3, 7, 11, 15, 19\}$ is, respectively,*

$$O\left(\frac{x}{(\log x)^2}\right), \quad O\left(\frac{x}{(\log x)^{5/2}}\right), \quad O\left(\frac{x}{(\log x)^3}\right), \quad O\left(\frac{x}{(\log x)^{7/2}}\right).$$

More generally, by Theorem 1.10, for every finite set $\mathcal{P}$ of admissible residuals $R \equiv 3 \pmod{4}$ — prime, or composite with $\omega(R) \leq 2$ — the number of hard primes failing every residual in $\mathcal{P}$ is $O_{\mathcal{P}}\left(x/(\log x)^{1+\kappa(\mathcal{P})}\right)$, where $\kappa(\mathcal{P}) = \sum_{R \in \mathcal{P}} \delta_R$ with $\delta_R = \frac{1}{2}$ for R prime or composite ≤ 107 and $\delta_R = \frac{1}{2} - 1/\varphi(R)$ for composite $R > 107$ (the two tiers of Theorem 1.10(iii)); in particular $\kappa(\mathcal{P}) = |\mathcal{P}|/2$ for every list drawn from the residuals up to 107. For the fifteen primes $\equiv 3 \pmod{4}$ up to 107 the exponent is $17/2$, and for the full 27-element admissible list $\{3, 7, 11, \dots, 107\}$ it is $29/2$. All but a proportion $O((\log x)^{-2})$ of hard primes are solved within $\{3, 7, 11, 15\}$ — one power of $\log x$ less than the counts above, since there are $\asymp x/\log x$ hard primes up to x .

Proposition 1.12 (aggregate identity families). *Let $p \equiv 1 \pmod{4}$ be prime and $R \equiv 3 \pmod{4}$. If $R \mid p+1$, then $k = ap^2$ is a certificate for residual R ; if $R \mid p+4$, then $k = a^2p$ is. Hence any prime factor $\equiv 3 \pmod{4}$ of $p+1$ or of $p+4$ yields an explicit solution, and no divisor shape $k = a^j p^i$ yields a further p -only condition. Adding the two corresponding linear forms to Theorem 1.11: the number of hard primes $p \leq x$ failing every prime residual up to 107 and both families is $O(x/(\log x)^{19/2})$; with the full 27-element admissible list the exponent is $31/2$.*

Theorem 1.13 (density reduction). *For every $A > 0$ there is a finite admissible set $S(A)$ such that the number of hard primes $p \leq x$ for which every residual in $S(A)$ fails is $O_A\left(x/(\log x)^{1+A}\right)$.*

Remark 1.14 (relation to Vaughan's bound). As exceptional-set estimates, Theorem 1.13, the chain of Theorem 1.11, and Theorem 4.5 are all weaker than Vaughan's bound $\ll x \exp(-c(\log x)^{2/3})$ [20], which already exceeds every fixed power of $\log x$. They are not offered as improvements on it. Two differences of kind: first, the exceptional sets here are cut out by finitely many *explicit* multiplicative conditions on shifted forms, each failure constructively certified by an exhibited divisor k , with the underlying lemmas exact or machine-verified. Second — and this is the substantive point — Theorem 4.5 is not primarily a density statement at all: it bounds $R_{\min}(p)$, the quantity Proposition 1.1 shows is the conjecture at p , by $\varepsilon \log \log x$ for almost every hard prime. A mean-value argument counts the primes that fail; it produces no admissible residual, and hence no bound of this shape, for the primes that succeed. On that axis the comparison does not arise: no bound on $R_{\min}$ was previously available by any method.

Organization. Proofs are arranged by logical dependency. Section 2 develops the obstruction framework (Propositions 2.1, 2.2) and proves Theorem 1.8, whose Corollary 2.3 depends on the Type-I/II dichotomy introduced there. Section 3 proves Theorem 1.2, then Theorem 1.3, then Theorems 1.4, 1.5 and 1.6, which rest on it. Section 4 proves Theorem 1.11 — base case first, then, after the monotonicity reduction (Lemma 4.1), Lemma 1.9 and its unconditional upgrade Theorem 1.10, the higher exponents — followed by Proposition 1.12, Theorem 1.13, and the uniform chain (Theorem 4.5). Section 5 develops the Burgess–reciprocity ladder on top of Theorem 1.10, with its own ladder census (Table 2).

Table 1 records, for each main result, exactly what its proof consumes and its formal-verification status, so that each layer can be audited independently.

Three printed results are retained as stepping stones although later results subsume them, and the reader may want the map stated once: Theorem 1.13 is superseded by Theorem 4.5 (which gives $(\log x)^{-A}$ for every A at once), but its fixed-list proof is simpler and is the template the uniform chain scales; Lemma 1.9 is subsumed by Theorem 1.10, and is retained as the machine-verified confirmation — it is the form carried by the Lean development; and Theorem 5.3 is dominated unconditionally by Theorem 4.5 (the fixed chain beats the p -adapted ladder unconditionally), its role being structural: it opens the conditional program of Section 5, the only known frame whose measured inputs would pass the ceiling of Remark 4.6.

Computations (Section 6) provide verified certificates for all hard primes below 10^{11} and calibrate the base-case constant and decay exponents of the chain.

2 The obstruction framework

Proposition 2.1 (character obstruction). *Let $r \mid R$ be prime with $r \equiv 3 \pmod{4}$. If every prime factor q of $m = pa$ satisfies $\left(\frac{q}{r}\right) = +1$, then (1.2) fails for R .*

Proof. Every divisor k of m^2 is a product of prime factors of m , so $\left(\frac{k}{r}\right) = +1$. A certificate requires $k \equiv -m \pmod{r}$, whence $\left(\frac{k}{r}\right) = \left(\frac{-m}{r}\right) = -1$ by (1.4) — impossible. $\square$

We call failures produced by Proposition 2.1 *Type I*; failures where m possesses a non-residue prime factor modulo every relevant r , yet no divisor of m^2 attains the exact class $-m \bmod R$, are *Type II*.

Proposition 2.2 (guaranteed-success classes). *Let $t \equiv -4^{-1}p^2 \pmod{R}$. If a has a prime factor q with $q \equiv tp^{-i} \pmod{R}$ for some $i \in \{0, 1, 2\}$, then $k = qp^i$ is a certificate. In particular, since $tp^{-2} \equiv -4^{-1} \pmod{R}$, any prime factor of a in the fixed class $-4^{-1} \bmod R$ certifies success regardless of p .*

Proof. $k = qp^i$ divides m^2 (as $q \mid a$ and $i \leq 2$) and lies in class $t \equiv -m$. Integrality of b, c is automatic from $k \equiv -m$ and $m^2/k \equiv m^2(-m)^{-1} \equiv -m \pmod{R}$. $\square$

The contrapositive of Proposition 2.2 is the sieve input for Theorem 1.13: failure at R forces $(p + R)/4$ to avoid prime factors in up to three explicit classes modulo R .

The next identity converts the per-residual character conditions into a single global statement about p , collapsing joint failure across residuals into one set of quadratic-character conditions.

Proof of Theorem 1.8. From $q \mid p + R$ we get $R \equiv -p \pmod{q}$. By quadratic reciprocity with $R \equiv 3 \pmod{4}$, $\left(\frac{q}{R}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{R}{q}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{-1}{q}\right) \left(\frac{p}{q}\right) = \left(\frac{p}{q}\right)$. $\square$

Corollary 2.3 (forced primes and correlated failure). *(i) The hard classes are the squares mod 840, so p is automatically a quadratic residue mod 3, 5, 7, and $p \equiv 1 \pmod{8}$ handles $q = 2$ (which divides a only for $R \equiv 7 \pmod{8}$, where $(\frac{2}{R}) = +1$): no prime forced into a shifted value $(p+R)/4$ by the congruences (H1)–(H3) carries a negative character. (ii) p fails Type-I at every prime residual $R \leq B$ simultaneously if and only if $(\frac{p}{R}) = +1$ for each such R and $(\frac{p}{q}) = +1$ for every unforced odd prime q (one not forced by the congruences of (H1)–(H3)) dividing any of the shifted values $(p+R)/4$, $R \leq B$ — one condition per distinct prime, shared across every residual it divides.*

Proof. Immediate from Theorem 1.8 together with (H1)–(H3) and the quadratic characters of the forced small primes. $\square$

In words, part (i) says that the small primes forced into the shifted values by the Mordell structure never break a Type-I failure — a conceptual reading of why the hard classes are exactly the square classes. By part (ii), failures at distinct residuals correlate through shared prime factors: shifted values with few distinct prime factors leave few independent sign conditions. Section 6 quantifies this for the record prime. Both parts are verified at scale (135,588 instances of Theorem 1.8, 61,438 forced-prime instances, zero violations of either; Appendix A).

3 Exact criteria

The individual criteria proved in this section are instances; the conceptual point is the framework behind them — every residual, prime or composite, admits a finite computable exact criterion (Theorem 1.3), and every criterion doubles as a sieve input. The four residuals worked out here are the ones the computations show carry 95% of the hard primes.

3.1 Proof of Theorem 1.2

Proof. Modulo 3 the unit group has order 2, so quadratic character equals class. Since $p^2 \equiv 1 \pmod{3}$, (1.3) gives $m \equiv 1$ and the target is $-m \equiv 2 \pmod{3}$; also $3 \nmid a$. *Necessity:* if every prime factor of a is $\equiv 1 \pmod{3}$, then $a \equiv 1 \pmod{3}$, hence $p = 4a - 3 \equiv 1 \pmod{3}$ as well, so every prime factor of $m = pa$ is $\equiv 1 \pmod{3}$ and every divisor of m^2 is $\equiv 1 \pmod{3}$: class 2 is never attained. *Sufficiency:* $k = q \equiv 2 \pmod{3}$ divides m^2 , and the recovery formulae give integers as in Proposition 2.2. $\square$

3.2 Proof of Theorem 1.3

Proof. A certificate is a divisor $k \mid m^2$ in class $-m \pmod{R}$. The set of classes of divisors of m^2 is determined by the classes of the prime factors of $m = pa$ together with their multiplicities, each factor being usable with exponent at most twice its multiplicity; exponents are only meaningful modulo $\text{ord}_R(\text{class}) \leq R - 1$, so multiplicities beyond $\lceil (R - 1)/2 \rceil$ are equivalent to that cap (Remark 3.1 illustrates why). The target class and the consistency relation $\prod(\text{classes}) \equiv a \equiv 4^{-1}p \pmod{R}$ are functions of $p \pmod{R}$ by (1.3). Hence solvability is a function of the capped class multiset and $p \pmod{R}$, a finite amount of data. $\square$

Remark 3.1 (why the cap works, and the size of the space). The one step above that is not pure bookkeeping is the multiplicity cap, and a small example shows both why it is harmless and what it buys. Take $R = 7$ and a factor of a in a class x of order 6 (say $x = 3$): a factor of

multiplicity μ may enter a divisor of m^2 with any exponent $0, 1, \dots, 2\mu$, and once $2\mu \geq 6$ these powers already sweep out the full cyclic group $\langle x \rangle$ — raising μ past $\lceil (R-1)/2 \rceil = 3$ enlarges nothing that a divisor subproduct can reach. For classes of smaller order the cap is generous rather than tight (a class of order 2 stabilizes at $\mu = 1$), which is fine: the enumeration ranges over a finite over-approximation. What finiteness buys is concrete: at $R = 7$ the capped space has 4^6 multisets before the structural constraints (i)–(iii) in the proof of Theorem 1.4 below and 1,536 consistent configurations after them; at $R = 11$ it compresses to 25 equivalence states. The enumeration is decidable, not merely heuristic, precisely because the cap and the consistency relation are proved before any code runs.

3.3 Proof of Theorem 1.4

Proof. $(\Rightarrow)$ is Proposition 2.1 with $r = R = 7$. $(\Leftarrow)$ reduces to a finite verification via Theorem 1.3. Write $(\mathbb{Z}/7)^* \cong \mathbb{Z}/6$ in discrete-log coordinates. Three structural constraints hold at hard primes:

- (i) $p \bmod 7 \in \{1, 2, 4\}$, by (H3);
- (ii) $2 \mid a$: by (H1), $p+7 \equiv 0 \pmod{8}$, so $a = (p+7)/4$ is even — the class of 2 (a residue mod 7) is always present;
- (iii) the factor classes multiply to $a \equiv 4^{-1}p \pmod{7}$.

Enumerating all class multisets with multiplicities capped at 3 subject to (i)–(iii) yields 1,536 consistent configurations; for each, one checks whether the target class is attained by a bounded-exponent subproduct. The verification (code: `theory.verify_R7_finite`) confirms: the target is attainable precisely when some factor class is a non-residue. $\square$

Remark 3.2. The hard-class hypotheses are essential: for $p = 701 \equiv 1 \pmod{4}$ (not hard), $a = 177 = 3 \cdot 59$ has the non-residue factor 3, yet $R = 7$ fails. Constraint (ii) — the forced even factor — is exactly what rescues every would-be Type-II configuration at hard primes. Empirically (Section 6), among 41,421 sampled failures of $R = 7$ at hard primes, *all* are Type I.

3.4 Proof of Theorem 1.5

Proof. By Theorem 1.3 with $R = 11$ (the unit group has order 10): multiplicities cap at 5, and at hard primes $3 \mid a$ (from (H2), $11 \equiv 2 \pmod{3}$), while $2 \nmid a$ (from (H1), $p+11 \equiv 4 \pmod{8}$) and $7 \nmid a$ (from (H3), $-11 \bmod 7 = 3 \notin \{1, 2, 4\}$). A dynamic-programming enumeration over the finite configuration space (code: `theory.finite_criterion_dp`) compresses it to 25 equivalence states: 16 succeed, 6 fail with all factor classes quadratic residues (case (a)), and exactly 3 fail with a non-residue present (case (b)); the three failing states decode to the three listed patterns. Two structural facts emerge from the enumeration. First, no state fails by a *proper-subgroup* obstruction: the only subgroup of $\mathbb{Z}/10$ containing an odd element properly is $\{0, 5\}$, and the consistency relation forces p 's class to extend it to all of $\mathbb{Z}/10$. Second, the case-(b) failures are pure *budget* failures — the exponent bounds, not the character obstruction, block the target: any of $v_3(a) \geq 2$, a factor in one of the classes 4, 5, 9, or a second non-residue factor restores solvability. $\square$

3.5 Proof of Theorem 1.6

Proof. ($\Rightarrow$, the Jacobi obstruction.) By (1.3), $m \equiv 4^{-1}p^2 = (2^{-1}p)^2 \pmod{15}$ is the square of a unit, so the Jacobi symbol satisfies $\left(\frac{m}{15}\right) = +1$; and $\left(\frac{-1}{15}\right) = \left(\frac{-1}{3}\right)\left(\frac{-1}{5}\right) = -1$. Hence $\left(\frac{-m}{15}\right) = -1$. The classes with $\left(\frac{q}{15}\right) = +1$ are exactly $\{1, 2, 4, 8\} = \langle 2 \rangle$, a subgroup of $(\mathbb{Z}/15)^*$ containing p 's class ($p \pmod{15} \in \{1, 4\}$, since $p \equiv 1 \pmod{3}$ by (H2) and $p \equiv \pm 1 \pmod{5}$ for all six hard classes). If every prime factor of a lies in $\langle 2 \rangle$, then every divisor of m^2 does, and the target $-m \equiv 11 \pmod{15}$ (an equality valid for both p -classes) is unreachable. This composite-modulus form of Proposition 2.1 is strictly stronger than the proposition applied to the single prime $r = 3 \mid 15$.

($\Leftarrow$.) By the composite form of Theorem 1.3 over $(\mathbb{Z}/15)^* \cong \mathbb{Z}/2 \times \mathbb{Z}/4$ (Carmichael exponent $\lambda(15) = 4$, so reachability saturates at multiplicity 2; multiplicities 0..5 capture every pair of saturation state and class product). The structural constraints at hard primes are: $p \pmod{15} \in \{1, 4\}$; $2 \mid a$ (from (H1), $8 \mid p + 15$), so the class of 2 always appears; and consistency $\prod(\text{classes}) \equiv a \equiv 4p \pmod{15}$. Enumerating the full space (code: `theory.verify_R15_finite`) yields 349,920 consistent configurations: 349,380 succeed and 540 fail — and *every* failing configuration has all factor classes in $\langle 2 \rangle$. In contrast to $R = 11$ there are no budget failures at all: the consistency relation forces $\prod \left(\frac{q}{15}\right)^{\mu_q} = \left(\frac{4p}{15}\right) = +1$, so non-residue factors occur with even total multiplicity, and the enumeration confirms that any two non-residue units (or one squared), together with the forced powers of 2, always reach the target. $\square$

Remark 3.3. The general composite- R engine (`theory.solvable_exact_general`), which evaluates the criterion of the composite meta-theorem directly in $(\mathbb{Z}/R)^*$, agrees with ground-truth divisor search on *every* composite admissible residual $R \in \{15, 27, 35, \dots, 99\}$ (7,938 sampled hard primes per residual below 10^9 , zero disagreements).

4 The sieve bound

We now prove the chain of Theorem 1.11, beginning with the base case, and then Lemma 1.9, its unconditional upgrade Theorem 1.10, the aggregate identities, and the density reduction.

Proof of Theorem 1.11, base case $\{3, 7\}$. Step 1 (exact criteria). Let p be a hard prime and set $n = (p + 3)/4$, so that $(p + 7)/4 = n + 1$: *the two shifted forms are consecutive integers*. By Theorems 1.2 and 1.4, p is exceptional (both residuals fail) precisely when

$$n \text{ has no prime factor } \equiv 2 \pmod{3} \quad \text{and} \quad n + 1 \text{ has no prime factor } \equiv 3, 5, 6 \pmod{7},$$

with $p = 4n - 3$ prime.

Step 2 (congruence bookkeeping). Fix one of the six hard classes $h \pmod{840}$; then $n \equiv (h + 3)/4 \pmod{210}$ is fixed, and the conditions at the primes 2, 3, 5, 7 are determined by the class (in particular $3 \nmid n$ and $7 \nmid n + 1$ automatically). It suffices to bound each class and sum.

Step 3 (containment in a sifted set). Let $z = x^{1/10}$. After discarding the $O(z)$ primes $p \leq z$, every exceptional p gives $n \leq (x + 3)/4$, $n \equiv n_0 \pmod{210}$, such that for every prime q with $7 < q \leq z$:

- (a) $q \equiv 2 \pmod{3} \implies n \not\equiv 0 \pmod{q}$;
- (b) $q \equiv 3, 5, 6 \pmod{7} \implies n \not\equiv -1 \pmod{q}$;

(c) $q \nmid 4n - 3$ (since $4n - 3 = p$ is prime, $p > z$).

Conditions (a), (b) only use “no forbidden factor up to z ”, which is implied by the full criteria — an upper bound needs no more.

Step 4 (sieve of dimension 2). This is a standard upper-bound sieve for the three linear forms $n, n + 1, 4n - 3$, with sifting function

$$\omega(q) = \mathbf{1}_{q \equiv 2(3)} + \mathbf{1}_{q \text{ NQR}(7)} + 1 \quad (q > 7),$$

whose forbidden classes $0, -1, 3 \cdot 4^{-1} \bmod q$ are pairwise distinct for $q \nmid 21$. By Dirichlet, ω has average $\kappa = \frac{1}{2} + \frac{1}{2} + 1 = 2$ in the Halberstam–Richert sense $\sum_{q \leq w} \omega(q) \log q / q = \kappa \log w + O(1)$. The Fundamental Lemma of sieve theory (see [7, Thm. 2.5] and [6, Ch. 6]) with $z = x^{1/10}$ bounds the exceptional set $E(x)$ of hard primes $p \leq x$ failing both residuals:

$$\#E(x) \ll x \prod_{7 < q \leq z} \left(1 - \frac{\omega(q)}{q}\right) \asymp x (\log z)^{-1} (\log z)^{-1/2} (\log z)^{-1/2} \asymp \frac{x}{(\log x)^2},$$

by Mertens’ theorem and its arithmetic-progression form. Summing over the six classes preserves the bound. The higher exponents of the chain are deduced after Lemma 1.9 below. $\square$

The verification behind Lemma 1.9 reduces to finitely many checks through the following elementary but load-bearing observation, which we isolate for scrutiny.

Lemma 4.1 (monotonicity reduction). *For a class multiset C (classes v with multiplicities μ_v) let $D(C) = \{\sum_v e_v v : 0 \leq e_v \leq 2\mu_v\}$ denote the set of attainable divisor logs (before the p -shifts). If $\mu_v \leq \mu'_v$ for all v , then $D(C) \subseteq D(C')$, where C' carries multiplicities μ'_v on the same classes; likewise enlarging the support enlarges D . Consequently a support S carries some failing configuration (any multiplicities ≥ 1 on S) if and only if the all-multiplicity-one configuration on S fails.*

Proof. Any admissible exponent vector for C is admissible for C' , giving the inclusions. If a configuration on S with multiplicities $\mu_v \geq 1$ fails, the target lies outside $D(C)$, hence outside the subset $D(C_1)$ attained at multiplicities one — so C_1 fails too; the converse inclusion is trivial. The same argument applies verbatim after adjoining the p -shifts. $\square$

Proof of Lemma 1.9. By Lemma 4.1 it suffices to check that for every support of size $(R - 1)/2$ over the nonzero discrete logs and every class of $p \bmod R$, the target is reachable already at minimal multiplicities (exponent budget 2 per class, p -budget 2); success there kills every configuration on every superset support at any multiplicities. The check deliberately ignores the consistency relation $\prod(\text{classes}) \equiv 4^{-1}p$: it therefore examines a *superset* of the realizable configurations, and can only overstate the failing supports — the lemma is conservative. For $R = 19, 23$ a direct enumeration of all $(R - 1) \binom{R-2}{(R-1)/2}$ (support, class of p) pairs suffices (437,580 resp. 7,759,752 pairs; code: `theory.verify_support_bound`). For the larger residuals the binomial count explodes, and the check is instead performed by a subset dynamic program (`theory.verify_support_bound_dp`): fold all supports into a table mapping each reachability mask to the maximal support size achieving it — valid because, by the same monotonicity, failure at any multiplicities implies failure at multiplicity one. The realized masks are heavily structured, so the state count stays tractable (3,001 states at $R = 31$ against 77.6 million supports; 19,530,971 states at $R = 107$, 331 seconds). All twelve residuals verify with zero violations; exact state counts are tabulated in Table 5. The bound is tight: the all-residue (Type I) configurations occupy exactly $(R - 3)/2$ nonzero classes. $\square$

Proof of Theorem 1.10. (i) Write $M = A_1 + \cdots + A_k$ with $A_i = \{0, v_i, 2v_i\}$, $k = |S|$, and let $H = \text{Stab}(M) = \{h : M + h = M\}$, a subgroup of $\mathbb{Z}/d$. Kneser's addition theorem [12] (in the many-summand form obtained by iterating [19, Thm. 5.5]) gives

$$|M| \geq \sum_{i=1}^k |A_i + H| - (k-1)|H|.$$

Case $H = \{0\}$. Every $|A_i| = 3$, except that the (at most one) index i with $v_i = d/2$ has $|A_i| = 2$ (there $2v_i = 0$); so $|M| \geq (3k-1) - (k-1) = 2k$. Since $M \neq \mathbb{Z}/d$, $2k \leq d-1$, and d even forces $k \leq d/2 - 1$. *Case $|H| = h > 1$.* M is a union of H -cosets and $M \neq \mathbb{Z}/d$, so $|M| \leq d-h$. Split $k_0 = |S \cap H| \leq h-1$ (distinct nonzero elements of H) and $k_1 = k - k_0$. For $v \in H$, $|A_v + H| = h$; for $v \notin H$, $A_v + H \supseteq H \cup (v + H)$, so $|A_v + H| \geq 2h$. Kneser's bound gives $d-h \geq h(k_0 + 2k_1 - k + 1) = h(k_1 + 1)$, so $k_1 \leq d/h - 2$ and $k \leq (h-1) + (d/h - 2) = h + d/h - 3$. For every divisor h with $2 \leq h \leq d/2$, the quadratic $h^2 - (\frac{d}{2} + 2)h + d \leq 0$ (its roots are 2 and $d/2$), i.e. $h + d/h - 3 \leq d/2 - 1$. Tightness: for S the nonzero even classes, $M(S)$ is the even subgroup and $|S| = d/2 - 1$.

(ii) In discrete-log coordinates on $(\mathbb{Z}/R)^* \cong \mathbb{Z}/d$, $d = R-1$, the divisor classes of m^2 contain the set $M(S)$ over the nonzero support S of $(p+R)/4$ at exponent budget 2 per class (Lemma 4.1: multiplicities ≥ 1 only enlarge budgets). If $|S| \geq d/2$, part (i) forces $M(S) = \mathbb{Z}/d$: every class — in particular the target $-m$ — is a divisor class of m^2 , without invoking the p -powers, and (1.2) holds. Contrapositively, failure at R forces $|S| \leq d/2 - 1 = (R-3)/2$, which is the conclusion of Lemma 1.9 for arbitrary R .

(iii) The argument in (i) never used cyclicity. In a finite abelian G of even order g : in the H -trivial case each involution $v \in S$ has $|\{1, v, v^2\}| = 2$, so $|M| \geq (3k-t) - (k-1) = 2k - t + 1$ and $M \neq G$ gives $k \leq \lfloor (g+t-2)/2 \rfloor$; the $|H| = h > 1$ case is verbatim, giving $k \leq h + g/h - 3 \leq g/2 - 1$ over proper subgroups. For $G = (\mathbb{Z}/R)^*$, $t = 2^{\omega(R)} - 1$; with $\omega(R) \leq 2$, $t \leq 3$ and the bound is $\leq g/2$, so failure forbids at least $(g-1) - g/2 = \varphi(R)/2 - 1$ non-identity unit classes — dimension $\geq \frac{1}{2} - 1/\varphi(R)$. The general- t clause is the same computation without the specialization: failure forbids at least $(g-1) - (g+t-2)/2 = (g-t)/2$ non-identity classes (minus one when the max is attained by $g/2 - 1$), and $2^{\omega(R)} \leq R^{O(1/\log \log R)} = R^{o(1)}$ gives the stated 0.44 threshold beyond an absolute bound. The upgrade to $\varphi(R)/2$ forbidden classes for $R \leq 107$ is the enumeration cited in the statement: the maximum support of a proper product-set is exactly $g/2 - 1$ for every composite admissible $R \leq 107$ (`theory.kneser_support_general`; e.g. $R = 95$: $g = 72$, maximum 35 over 583,455 realizable product-sets). That enumeration is a proof component of the $R \leq 107$ clause, on the same footing as the checks of Remark 1.7; the $\frac{1}{2} - 1/\varphi(R)$ tier needs Kneser alone. $\square$

Remark 4.2. Theorem 1.10(ii) is strictly stronger than what the sieve needs: at support $\geq (R-1)/2$ every class is reachable, for every target and every p . The strong form is also confirmed numerically (`theory.verify_support_bound_strong`): the maximum support of a non-full reachability mask equals $(R-3)/2$ exactly at every prime residual checked. Lemma 1.9 and its dynamic-program verifications stand as independent machine confirmations of special cases (and remain the form carried by the Lean development).

Proof of Theorem 1.11, higher exponents. Write $n = (p+3)/4$; the relevant forms are n , $n+1$, $n+2 = (p+11)/4$, $n+3 = (p+15)/4$, $n+4 = (p+19)/4$, and $4n-3 = p$. Joint failure of $\{3, 7, 11\}$ implies, by Theorems 1.2, 1.4, 1.5: n free of primes $\equiv 2 \pmod{3}$; $n+1$ free of non-residues mod

7; and, splitting by the two cases of Theorem 1.5: *branch (a)* — $n + 2$ free of the five non-residue classes mod 11 (density $\frac{1}{2}$); *branch (b)* — $n + 2$ free of the six classes $\{4, 5, 7, 8, 9, 10\}$ mod 11 (density $\frac{3}{5}$). Each branch is a four-form sieve problem of dimension $1 + \frac{1}{2} + \frac{1}{2} + \frac{1}{2} = \frac{5}{2}$, resp. $1 + \frac{1}{2} + \frac{1}{2} + \frac{3}{5} = \frac{13}{5}$. Since $\frac{13}{5} > \frac{5}{2}$, the Type-II branch contributes $O(x/(\log x)^{13/5}) = o(x/(\log x)^{5/2})$: the bound is the maximum of the branch exponents, attained by the character branch (a). The Fundamental Lemma applied to each branch gives $O(x/(\log x)^{5/2})$ in total.

Adjoining $R = 15$: by Theorem 1.6, failure at 15 is the single condition “ $n + 3$ free of the four classes $\{7, 11, 13, 14\}$ mod 15” — one branch, density $\frac{4}{8} = \frac{1}{2}$ over the $\varphi(15) = 8$ unit classes — raising each combined branch to dimension ≥ 3 : joint failure of $\{3, 7, 11, 15\}$ is $O(x/(\log x)^3)$.

For $\{3, 7, 11, 15, 19\}$, Theorem 1.10 (equivalently, at this residual, Lemma 1.9) splits failure at 19 into finitely many maximal-support branches, each forbidding an explicit set of at least 9 of the 18 unit classes mod 19 on $n + 4$ — dimension $\geq \frac{1}{2}$. Crossing with the branches at 11 leaves finitely many combined branches, each a six-form problem of dimension $\geq \frac{7}{2}$ (those involving the Type-II case at 11 have dimension $\geq \frac{31}{10} + \frac{1}{2} > \frac{7}{2}$ and are again lower-order); the forbidden residues of the six forms are pairwise distinct for $q > 19$, smaller primes being absorbed into the congruence class of n . Summing the finitely many branches gives $O(x/(\log x)^{7/2})$.

The general case is identical: for a finite admissible list $\mathcal{P}$ (prime, or composite with $\omega(R) \leq 2$), the forms $n + (R - 3)/4$, $R \in \mathcal{P}$, are distinct integer shifts of n ; Theorems 1.2, 1.4, 1.5, 1.6 handle $R = 3, 7, 11, 15$ and Theorem 1.10 handles every other residual (part (ii) for primes, part (iii) for composites), each contributing dimension δ_R through finitely many maximal-support branches — $\delta_R = \frac{1}{2}$ for primes and for composites ≤ 107 , $\delta_R = \frac{1}{2} - 1/\varphi(R)$ for larger composites (the two tiers of part (iii)) — plus 1 for primality: total dimension $\geq 1 + \kappa(\mathcal{P})$, with forbidden residues pairwise distinct for $q > \max \mathcal{P}$ and smaller primes absorbed into the congruence class of n . For the fifteen primes up to 107 this gives exponent $\frac{17}{2}$; for the full 27-element admissible list, $\frac{29}{2}$. $\square$

Proof of Proposition 1.12. $k = ap^2$ divides $m^2 = p^2a^2$ and $k \equiv -m$ if and only if $ap^2 \equiv -pa$, i.e. $p \equiv -1 \pmod{R}$; $k = a^2p$ divides m^2 and $a^2p \equiv -pa$ if and only if $a \equiv -1$, i.e. $p \equiv -4 \pmod{R}$ by (1.3). Integrality of b, c is automatic. Scanning all shapes $k = a^j p^i$ ($j, i \leq 2$) yields no further p -only condition. The failure of both families states that the two forms $(p + 1)/2$ and $p + 4$ (in n : $2n - 1$ and $4n + 1$) are free of primes $\equiv 3 \pmod{4}$ — two further sifting conditions of dimension $\frac{1}{2}$ each on forms distinct from all previous ones. $\square$

Empirically the two families alone cover 74% of hard primes below 10^9 — and none of the four primes of Section 6 that admit a unique working residual ≤ 107 .

Remark 4.3. By Theorem 1.10, every residual — prime or composite — adds $\frac{1}{2}$ to the exponent, so the chain’s exponent is bounded only by the number of residuals used; the constants $O_{\mathcal{P}}$, however, grow with $|\mathcal{P}|$ (branch counts and the remainder term of the Fundamental Lemma). The chain is a concrete instantiation of Theorem 1.13 with explicit constants. One caveat frames its scope, and it is fundamental: no fixed finite list can be pushed to *emptiness* by upper-bound sieving: the count bound $x/(\log x)^\kappa$ diverges for every fixed κ , and extrapolating the calibration of Section 6 to any fixed list yields a positive density constant, i.e. infinitely many expected exceptional primes with the record minimal residual growing slowly without bound. The correctly posed remaining problem is a slowly growing bound $R_{\min}(p) \ll f(p)$, which requires existence (lower-bound) technology beyond upper-bound sieves.

Proof of Theorem 1.13. Fix $A > 0$. Choose $B = B(A)$ so large that

$$\sum_{\substack{3 \leq R \leq B \\ R \equiv 3 \pmod{4}}} \frac{1}{\varphi(R)} \geq \sum_{\substack{3 \leq R \leq B \\ R \equiv 3 \pmod{4}}} \frac{1}{R} > A; \quad (4.1)$$

this is possible since the reciprocal sum over the progression $R \equiv 3 \pmod{4}$ diverges, and $\varphi(R) \leq R$. Let $S(A) = \{R \equiv 3 \pmod{4} : 3 \leq R \leq B\}$ and $M = 840 \cdot \prod_{R \in S(A)} R$.

Write $n = (p + 3)/4$ as before, so $(p + R)/4 = n + (R - 3)/4$ for each $R \in S(A)$: a family of $|S(A)|$ integer linear forms in n with pairwise distinct shifts, together with the form $4n - 3 = p$. Set $z = x^{1/10}$ and discard the $O(z)$ primes $p \leq z$. Fix a residue class $n \equiv n_0 \pmod{M}$ compatible with p hard; there are finitely many, and it suffices to bound each. Within such a class, for every $R \in S(A)$ the quantities $t \equiv -4^{-1}p^2$ and $p^{-1} \pmod{R}$ are constant, so the class set $S_R(p) = \{t, tp^{-1}, tp^{-2}\} \pmod{R}$ of Proposition 2.2 is a fixed nonempty set of $t_R = t_R(n_0) \in \{1, 2, 3\}$ classes. By the contrapositive of Proposition 2.2, if every residual in $S(A)$ fails at p , then for each $R \in S(A)$ the form $n + (R - 3)/4$ has *no* prime factor in $S_R(p)$; and $4n - 3$, being prime and greater than z , has no prime factor $\leq z$.

Sift the class $\{n \leq (x + 3)/4 : n \equiv n_0 \pmod{M}\}$ by the primes $q \leq z$, $q \nmid M$. A prime q forbids: the class $n \equiv 3 \cdot 4^{-1} \pmod{q}$ (from the primality of $4n - 3$), and, for each $R \in S(A)$ with $q \pmod{R} \in S_R(p)$, the single class $n \equiv -\frac{R-3}{4} \pmod{q}$ (a prime factor q of the form $n + (R - 3)/4$ in a guaranteed-success class cannot occur). Thus

$$\omega(q) = 1 + \#\{R \in S(A) : q \pmod{R} \in S_R(p)\},$$

which is bounded by $|S(A)| + 1$; the forbidden residues of distinct forms are pairwise distinct for $q > B$, since the shifts $(R - 3)/4$ differ by less than B and $q \nmid M$. For each R , the density of primes q with $q \pmod{R} \in S_R(p)$ is $t_R/\varphi(R)$ by Dirichlet's theorem, so ω has average

$$\kappa = 1 + \sum_{R \in S(A)} \frac{t_R}{\varphi(R)} \geq 1 + \sum_{R \in S(A)} \frac{1}{\varphi(R)} > 1 + A,$$

in the Halberstam–Richert sense, by (4.1). The Fundamental Lemma [7, Thm. 2.5] in dimension κ then bounds the sifted count by $O_A\left(x \prod_{q \leq z} (1 - \omega(q)/q)\right) = O_A(x(\log x)^{-\kappa}) = O_A\left(x/(\log x)^{1+A}\right)$, and summing over the finitely many classes $n_0 \pmod{M}$ preserves the bound. $\square$

Remark 4.4. Theorem 1.13 is the strongest statement the pure sieve can give: density $O((\log x)^{-A})$ for every A , but not emptiness. The conjecture for hard primes would follow from a *finite covering hypothesis*: some finite S admits, for every hard prime, a residual $R \in S$ with (1.2) (the converse is not automatic — the conjecture bounds $R_{\min}(p)$ by $2p$, not uniformly). The data to 10^{11} is covered by $S = \{3, 7, 11, \dots, 107\}$ — yet both the calibration (Remark 4.3) and the conditional construction of Section 7 predict that any fixed list eventually fails.

4.1 The uniform chain

Theorem 1.13 fixes the list $S(A)$ once A is fixed; letting the list *grow with* x is purely a question of uniformity in the constants of Theorem 1.11 — the branch count per residual, the class

modulus, and the sieve constants in growing dimension. Auditing those constants shows the growth is affordable for lists of length a small multiple of $\log \log x$, which yields the strongest unconditional statement of this paper.

Theorem 4.5 (uniform chain). *There are absolute constants $\varepsilon, c > 0$ and an (ineffective) x_0 such that for all $x \geq x_0$,*

$$\#\{\text{hard primes } p \leq x : R_{\min}(p) > \varepsilon \log \log x\} \leq x \exp(-c(\log \log x)^2).$$

In particular, all but $O(xe^{-c(\log \log x)^2})$ hard primes $p \leq x$ possess an explicit, divisor-exhibited residual certificate at some admissible $R \leq \varepsilon \log \log x$ — unconditionally.

Proof. Write $L = \log \log x$, $B = \varepsilon L$, and let $\mathcal{R}$ be the admissible residuals $3 \leq R \leq B$ ($|\mathcal{R}| \sim B/4$). Discard the $O(\sqrt{x})$ primes $p \leq \sqrt{x}$. A hard prime with $R_{\min}(p) > B$ fails every $R \in \mathcal{R}$; we bound the count of such p by a single sieve application per *branch*, as follows.

Dimensions. Each $R \in \mathcal{R}$ contributes a sifting dimension δ_R : the exact criteria at $R = 3, 7, 11, 15$ and Theorem 1.10(ii) at primes give $\delta_R = \frac{1}{2}$; composites ≤ 107 give $\frac{1}{2}$ by the sharpened tier of Theorem 1.10(iii); and every remaining composite — with *no* restriction on $\omega(R)$ — gives $\delta_R \geq \frac{1}{2} - (2^{\omega(R)-1} + 1)/\varphi(R) \geq 0.44$ beyond an absolute bound R_1 , by the general- t clause of Theorem 1.10(iii). Hence $\kappa := 1 + \sum_{R \in \mathcal{R}} \delta_R \geq 1 + 0.44(|\mathcal{R}| - R_1) \geq 1 + B/10$ for B large.

Branches. For each R , failure confines the support of $(p+R)/4$ to one of at most $2^{\varphi(R)}$ failing supports S_R (crudely, all subsets); fixing $(S_R)_{R \in \mathcal{R}}$ and the class n_0 of $n = (p+3)/4$ modulo $\widetilde{M} = \text{lcm}(840, \prod_{\ell \leq B} \ell) = e^{O(B)}$ costs a factor $\prod_R 2^{\varphi(R)} \cdot e^{O(B)} = e^{O(B^2)}$ branches.

Sieve. Within a branch, sift the integers $n \leq (x+3)/4$, $n \equiv n_0 \pmod{\widetilde{M}}$, by the primes $\ell \in (B, z]$, $z = x^{1/s}$: the forbidden classes mod ℓ are $n \equiv 3 \cdot 4^{-1}$ (so that $p = 4n - 3$ keeps no factor $\leq z$) and, for each R with $\ell \bmod R$ outside S_R , $n \equiv -(R-3)/4$ — pairwise distinct for $\ell > B$, so $\omega(\ell) \leq |\mathcal{R}| + 1 < \ell/2$. The density condition holds with the dimension κ above and (Ω) -constant $O(B \log B)$, by Mertens' theorem in the progressions modulo $R \leq B$ (Siegel–Walfisz range, whence the ineffective x_0). Apply Brun's sieve of order $r \asymp \kappa$ ([7, Ch. 2]; for constants uniform in the dimension, [6, Thm. 6.9]) at $s = 10r$, so that $z^r = x^{1/10}$: since the sifted sequence consists of *integers* in a progression — no primality is detected anywhere — the remainder is $\sum_{d \leq z^r} (B+2)^{\omega(d)} \ll x^{1/10} e^{O(BL)} = x^{1/10+o(1)}$ for a suitable absolute choice of the implied constants, and the main term is

$$\ll \frac{x}{\widetilde{M}} \prod_{B < \ell \leq z} \left(1 - \frac{\omega(\ell)}{\ell}\right) \leq x \exp(-\kappa(L - O(\log B)) + O(B \log B)).$$

Accounting. Multiplying by the branch count,

$$\begin{aligned} \#\{p \leq x : R_{\min}(p) > B\} &\leq e^{O(B^2)} x \exp(-\kappa L + O(B \log B) + O(\kappa \log B)) \\ &\leq x \exp(-BL/10 + O(B^2)), \end{aligned}$$

and with $B = \varepsilon L$ the exponent is $-\varepsilon L^2/10 + O(\varepsilon^2 L^2) \leq -\varepsilon L^2/20$ for ε below an absolute threshold. $\square$

Remark 4.6 (the ceiling, and the one lever). Theorem 4.5 sits at the boundary set by its *branch factor*: $e^{\Theta(B^2)}$ from the crude all-subsets bound on failing supports, against the sieve gain $e^{-BL/10}$, forces $B \ll L$ and the ceiling $\exp(-cL^2)$. A closer audit (Section 4.2) shows the branch

factor is the *only* binding wall: the class modulus $\widetilde{M} = e^{O(B)}$ never multiplies the main term (the classes partition the integers), and the Mertens/Siegel–Walfisz inputs remain valid for B up to any fixed power of $\log x$ at the already-paid price of ineffectivity. Classifying the failing supports into polynomially many structured branches — the subject of Section 4.2 — would therefore raise Theorem 4.5 to an exceptional set $x \exp(-(\log x)^{c_0})$ with $R_{\min} \leq (\log x)^{c_0}$ (Theorem 4.11), with the family ceiling strictly below exponent $\frac{1}{2}$ (see the discussion there). Vaughan’s bound $x \exp(-c(\log x)^{2/3})$ [20] remains stronger as a pure density statement at every tier of this family; what Theorem 4.5 adds is structural: an explicit certificate at a small residual for almost every hard prime, a statement mean-value methods do not address. The calibrations $J = (\log x)^\theta$ of Theorem 5.11 remain out of reach of every unconditional variant considered here.

4.2 The branch classification

Fix a prime residual R , $d = R - 1$, and work in the discrete-log model of Theorem 1.10(i): a *support* $S \subseteq \mathbb{Z}/d \setminus \{0\}$ at exponent budget 2 per class has reach $M(S) = \sum_{v \in S} \{0, v, 2v\}$, and S *fails* for a target t if $t \notin M(S)$ (the p -budget is absorbed as one more summand). By Lemma 4.1 every failing configuration lives on a failing support, and every failing support extends to a *maximal* one (S failing, $S \cup \{w\}$ succeeding for all w). Two facts calibrate the problem. First, maximality forces no size concentration: maximal failing supports exist at every size scale from $\Theta(\log d)$ (complete-partition constructions realizing $M(S) = \mathbb{Z}/d \setminus \{t\}$, a *singleton miss*) up to the extremal $d/2 - 1$. Second, *enumerating* them is hopeless:

Theorem 4.7 (branch count blow-up). *There is $c > 0$ such that for every large even d , the number of maximal failing supports for the target $t = -1$ in $\mathbb{Z}/d$ is at least $2^{c\sqrt{d}}$: the supports $[1, r] \cup T' \cup \{X\}$, $T' \subseteq [r + 1, 2r]$, $|T'| = \lceil r/2 \rceil$, $r = \lceil \sqrt{d/4} \rceil$, with X completing the sum to $d/2 - 1$, are pairwise distinct complete partitions (for d large, $X > 2r$, so X collides with no other part and T' is recoverable from the support; the chain condition at the part X reduces to $d/2 - 2 \leq 3r(r + 1)$, which holds since $r^2 \geq d/4$), and each has reach exactly $\mathbb{Z}/d \setminus \{-1\}$, hence is maximal failing.*

So a useful classification must produce *containers* — polynomially many structured sets covering all failing supports — not an enumeration. The kernel (periodic) case is under complete control:

Theorem 4.8 (kernel branch; lossless reduction). *Let S fail for t and $H = \text{Stab}(M(S)) \neq \{0\}$, and write $\varphi_H: \mathbb{Z}/d \rightarrow \mathbb{Z}/\bar{d}$. Then $t \notin M(S) + H$, the projection $\bar{T} = \varphi_H(S \setminus H)$ is a failing support for $\bar{t} = \varphi_H(t) \neq 0$ at the same budget, and $S \subseteq (H \setminus \{0\}) \cup \varphi_H^{-1}(\bar{T})$; if S is maximal then moreover $H \setminus \{0\} \subseteq S$. Iterating: every failing support of $\mathbb{Z}/d$ is contained in $\psi^{-1}(\{0\} \cup T') \setminus \{0\}$ for a composed quotient $\psi: \mathbb{Z}/d \rightarrow \mathbb{Z}/d'$ and an aperiodic maximal failing support T' downstairs, and this container has at most $\lfloor d/2 \rfloor - 1$ elements. (The support bound extends to odd moduli, as needed for the quotients: for odd $\bar{d}$ no summand is degenerate, so in the aperiodic case Kneser’s bound gives $|T| \leq (\bar{d} - 3)/2$; for nontrivial stabilizer of order h the count is $\leq h + \bar{d}/h - 3 \leq \bar{d}/3 \leq (\bar{d} - 3)/2$ for $\bar{d} \geq 9$, smaller odd $\bar{d}$ having no proper nontrivial subgroups.)*

Proof. $M + H = M$ gives $H \subseteq M$ and $t \notin M + H$ (else $t \in M$), so $\bar{t} \neq 0$. The projection $\varphi_H(M(S)) = \sum_v \{0, \bar{v}, 2\bar{v}\}$ contains $M(\bar{T})$ computed downstairs at budget 2 (collisions only enlarge budgets; Lemma 4.1), and H -saturation of $M(S)$ turns $t \notin M(S)$ into $\bar{t} \notin M(\bar{T})$: $\bar{T}$ fails. The containment is definitional. For maximality: any $w \in H \setminus (S \cup \{0\})$ has $A_w \subseteq H$,

so $M(S \cup \{w\}) = M(S) + A_w = M(S)$ still fails — w must already lie in S . Iterate on the quotient (strong induction on d , extending $\bar{T}$ to a maximal support downstairs first); the size bookkeeping per step is $(h-1) + h|T'| \leq h-1 + h\lfloor(\bar{d}-2)/2\rfloor \leq d/2 - 1$, exact at even $\bar{d}$, with strict inequality at odd $\bar{d}$. $\square$

The open core is the aperiodic case. All the evidence points to one container shape — a coset-progression *window*: for $K \leq \mathbb{Z}/d$, a unit direction and a circular window W of quotient classes, set $C(K, W) = \varphi_K^{-1}(\{0\} \cup W) \setminus \{0\}$. The family $\{C(K, W) : |C(K, W)| \leq \lfloor d/2 \rfloor - 1\}$ has fewer than d^3 members (a container is determined by the kernel size $\kappa \mid d$, a unit direction, a window start, and a window length $\leq \bar{d}/2$: at most $\sum_{\kappa \mid d} \bar{d}^3/2 \leq \zeta(3)d^3/2$ choices).

Conjecture 4.9 (aperiodic branch classification). *Every maximal failing support S with $\text{Stab}(M(S)) = \{0\}$ for a target t in $\mathbb{Z}/d$ satisfies $S \subseteq C(K, W)$ for some window container with $|C(K, W)| \leq \lfloor d/2 \rfloor - 1$.*

The evidence: exhaustive verification for *all* moduli $d \leq 60$ and all targets (24.25 million orbit-weighted maximal supports; zero violations; the bound $\lfloor d/2 \rfloor - 1$ is attained, hence optimal); exhaustive verification at $d = 72$, the smallest *non-Hajós* cyclic order — the one place aperiodic factorization pathologies could first appear — over all 11 target orbits and 14.99 million maximal supports, with extensive stratified sampling (including supports seeded from explicit aperiodic de Bruijn factorizations of $\mathbb{Z}/72$) at every non-Hajós order up to 240; independent verification in the residual model at $R = 19, 23, 31, 43, 47, 59$ (495,782 maximal supports at $R = 59$) and at the composite moduli $R = 15, 35, 39$ (entry point `erdos_straus.branch_enum`; archived); a proved subfamily (the signed complete partitions, windows with room to spare); and a rank-2 witness ($S = \{11, 17, 19, 29\} \subseteq \mathbb{Z}/30$, a singleton miss contained in no short one-dimensional window but in a $K = \langle 10 \rangle$ window container of size 14) showing the subgroup component of the shape is necessary. Combining Theorem 4.8 with Conjecture 4.9 in the quotients (windows pull back to windows along composed quotients):

Theorem 4.10 (container form; conditional on Conjecture 4.9). *Every failing support for (d, t) is contained in one of fewer than d^3 explicit window containers of size $\leq \lfloor d/2 \rfloor - 1$; hence failure of residual R confines the factor classes of $(p+R)/4$ to one of $O(R^3)$ branches, each forbidding at least half the nonzero classes — sieve dimension $\frac{1}{2}$ per residual, with the branch count reduced from $2^{\varphi(R)}$ to $O(R^3)$.*

Theorem 4.11 (conditional uniform chain, polylog tier). *Assume Conjecture 4.9 (for all moduli). Then there is an absolute $c_0 > 0$ such that all but $O(x \exp(-(\log x)^{c_0}))$ hard primes $p \leq x$ satisfy $R_{\min}(p) \leq (\log x)^{c_0}$.*

Proof sketch. Rerun the proof of Theorem 4.5 with list length $B = (\log x)^\theta$. The branch factor over the list drops from $e^{\Theta(B^2)}$ to $\prod_{R \leq B} O(R^3) = e^{(3+o(1)) \cdot \frac{B}{4} \log B}$; the class conditioning never multiplies the main term (the classes partition the sifted integers, so the per-class main terms resum); and Siegel–Walfisz covers the Mertens-in-progressions inputs for moduli up to any fixed power of $\log x$. The per-residual balance is then a sieve gain $\delta(1-\theta)L$ against a branch-and-slop cost $(3+c_M)\theta L$, where $L = \log \log x$, $\delta = \frac{1}{2}$, and c_M collects the per-residual slop — the Mertens-in-progressions constants, the discarded initial segment $\ell \leq B$ of sieving primes, and the (Ω) -condition constant, together of order 1. This is admissible for every $\theta < \delta/(3+c_M+\delta)$ — with the constants proved here, any θ below roughly $1/9$. The family ceiling sits strictly below $\frac{1}{2}$: no container scheme of this shape has fewer than $\sim d$ members (distinct progression

directions already require distinct containers), so the count exponent is at least 1 and the balance caps θ at $\delta/(1 + c_M + \delta) < \frac{1}{3}$ — consistent with, and sharper than, the $\exp(-(\log x)^{1/2+o(1)})$ outer ceiling of the wider method family, itself short of Vaughan’s $2/3$. A quasi-polynomial classification ($d^{O(\log d)}$ containers) would still give exceptional sets $x \exp(-\exp(c\sqrt{\log \log x}))$, past every $\exp(-C(\log \log x)^A)$. $\square$

Conjecture 4.9 is thus one well-posed additive-combinatorics statement — a Kneser/Freiman-type container theorem for bounded-exponent subset sums — standing between Theorem 4.5 and its polylog tier. Where its difficulty lives can now be pinned exactly. Near-critical Kneser analysis is *vacuous* here (aperiodic failing supports with $|S| \geq 3$ are never Kneser-tight — the governing mechanism is exact covering by complete sequences, not small doubling); the singleton-miss supports satisfy the symmetry $2 \sum_{v \in S} v = 2t$, which pins their windows (and at *odd* moduli singleton misses do not exist at all: the miss set always has even size, so the floor family there is $|F| = 2$, auto-maximal and window-shaped). The one mechanism previously feared — chaining of *punctured* (cofinite) intermediate reaches, where direction information degenerates and non-Hajós factorization pathologies could hide — is provably empty:

Proposition 4.12 (β -vacuity). *Let S fail for t in $\mathbb{Z}/d$, let $v \in S$, and suppose the deletion step at v is Kneser-critical ($|M(S)| = |M(S \setminus \{v\})| + 2$) of punctured type: the two new elements are isolated punctures of the complement. Then $F - v = F$ for the miss set $F = \mathbb{Z}/d \setminus M(S)$, so $v \in \text{Stab}(M(S))$. In particular, for an aperiodic maximal failing support every critical deletion step is of window (α) type.*

Proof. Write $A = M(S \setminus \{v\})$ and $\mathbb{Z}/d \setminus A = F \sqcup \{x, y\}$ with x, y the punctures. For $f \in F$: since $f \notin M(S) = A \cup (A + v) \cup (A + 2v)$, both $f - v$ and $f - 2v$ avoid A , so $f - v \in F \cup \{x, y\}$. If $f - v = x$, then $x - v = f - 2v \notin A$, contradicting the isolation of x ($x - v \in A$; the $v = d/2$ two-cycle case is the same contradiction). Likewise for y . So $F - v \subseteq F$, hence $F - v = F$ by finiteness. $\square$

What remains open in Conjecture 4.9 is therefore only the control of *non-critical* deletion steps — slack-consuming extensions, a question about exact-covering efficiency that is indifferent to the Hajós/non-Hajós divide (consistent with the exhaustive verification at $d = 72$ above). And the payoff is robust to the conjecture’s scope: by a Chen-type P_3 count there are $\gg B/(\log B)^2$ admissible residuals $R \leq B$ whose modulus $d = R - 1$ is a Hajós order (equivalently $(R - 1)/2$ is a prime power or has at most three prime factors), which is dense enough that assuming Conjecture 4.9 *only for Hajós orders* — where the group theory is tame — yields Theorem 4.11 with the same exponent up to $o(1)$.

5 A conditional program: the Burgess–reciprocity ladder

Theorem 1.10 closed the fixed-list axis, and Theorem 4.5 pushed the growing-list axis to its unconditional ceiling $\exp(-c(\log \log x)^2)$ (Remark 4.6); no list-based bound can reach emptiness (Remark 4.3). This section develops the program that opens past the ceiling: a *p-adapted* family of residuals, chosen through quadratic reciprocity, along which the failure modes are localized, partially proved, and reduced to a single measured hypothesis. Entry points and archives for all measurements cited here are in Appendix A.

5.1 Necessity, purity, and the selected residual

Lemma 5.1 (composite reciprocity). *Let $R \equiv 3 \pmod{4}$ be any admissible residual and q an odd prime with $q \mid (p + R)/4$ and $\gcd(q, R) = 1$. Then the Jacobi symbol satisfies $\left(\frac{q}{R}\right) = \left(\frac{p}{q}\right)$.*

Proof. $q \mid p + R$ gives $R \equiv -p \pmod{q}$; Jacobi reciprocity with $R \equiv 3 \pmod{4}$ gives $\left(\frac{q}{R}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{R}{q}\right) = (-1)^{\frac{q-1}{2}} \left(\frac{-1}{q}\right) \left(\frac{p}{q}\right) = \left(\frac{p}{q}\right)$ — the proof of Theorem 1.8 verbatim, which never used primality of R . $\square$

Lemma 5.2 (Jacobi necessity). *For every admissible R : if no prime factor of $a = (p + R)/4$ has Jacobi symbol $\left(\frac{\cdot}{R}\right) = -1$, then (1.2) fails at R .*

Proof. All factors of a Jacobi residues gives $\left(\frac{a}{R}\right) = +1$; by consistency (1.3), $\left(\frac{p}{R}\right) = \left(\frac{4a}{R}\right) = +1$ as well, so every prime factor of $m = pa$ has symbol $+1$, hence every divisor $k \mid m^2$ has $\left(\frac{k}{R}\right) = +1$, while the target has $\left(\frac{-m}{R}\right) = \left(\frac{-1}{R}\right) = -1$. $\square$

Call R *pure* (at hard primes) if the converse also holds: a Jacobi non-residue factor present implies success. Theorems 1.2, 1.4 and 1.6 say exactly that $R = 3, 7, 15$ are pure; the census below shows they are the *only* pure residuals up to 107, with the conditional failure rate at the others ranging from 1.3% ($R = 23$) to 73% ($R = 67$).

Now run Lemma 5.1 backwards. For a hard prime p let $q = q(p)$ be its least Legendre non-residue (hard primes are residues at 3, 5, 7 by Corollary 2.3, so $q \geq 11$; Burgess bounds $q \ll p^{1/(4\sqrt{e})+\varepsilon}$ unconditionally [2] and GRH gives $q \ll (\log p)^2$ [1]; the observed maximum below 10^{11} is 83), and let the *selected residual* be $R_0 = (-p) \bmod 4q$, so that $R_0 \equiv 3 \pmod{4}$ automatically and $q \mid (p + R_0)/4$. By Lemma 5.1, $\left(\frac{q}{R_0}\right) = \left(\frac{p}{q}\right) = -1$: the all-residue failure mode of Lemma 5.2 is impossible at R_0 *by construction*. The *q-ladder* is the progression $R_j = R_0 + 4qj$, along which the same holds at every rung.

5.2 The census

Computed over every one of the 1,587,581 hard primes below 10^9 , with mask-free window samples of 10,000 primes per half-decade beyond, and the complete 10^{11} deep tail (Appendix A), the census is summarized in Table 2.

Failure anatomy (all 82,808 first-rung failures below 10^9): 97.5% are *true budget* failures — the subgroup generated by the factor classes contains the target and only the bounded exponents miss it — and 29% miss exactly one unit class, the target itself. No failure has more than 5 distinct non-identity factor classes (Theorem 1.10 would allow $(R - 3)/2$), with the conditional failure probability decaying by a factor ≈ 7 per additional class; and the Jacobi non-residue mass sits at the parity minimum forced by consistency in 97.9% of failures — the mechanism of Theorem 1.6’s proof, observed at every residual.

5.3 What is proved

Theorem 5.3 (ladder chain). *Fix a prime $q \geq 11$ and $J \geq 1$. The number of hard primes $p \leq x$ with $\left(\frac{p}{q}\right) = -1$ failing every rung R_j , $0 \leq j < J$, of their q -ladder is $O_{q,J}(x/(\log x)^{1+\kappa_J})$, where $\kappa_J = \sum \delta_{R_j}$ over the rungs with at most two prime factors, with $\delta_R = \frac{1}{2}$ for R prime or composite ≤ 107 and $\delta_R = \frac{1}{2} - 1/\varphi(R)$ for larger composites; in particular $\kappa_J \geq J'/2 - \sum_j 1/\varphi(R_j)$ with J' the number of such rungs.*

Proof. The shifted forms $(p + R_j)/4 = n + (R_0 - 3)/4 + qj$ are distinct integer shifts of $n = (p + 3)/4$, and Theorem 1.10 — part (ii) at prime rungs, part (iii) at composite rungs with $\omega \leq 2$ — makes failure at each usable rung a sifting condition of dimension $\geq \delta_{R_j}$ (the two tiers of part (iii) at composite rungs), exactly as in the proof of Theorem 1.11, which applies verbatim to this p -adapted list. $\square$

Theorem 5.4 (first rung). *Fix a prime $q \geq 11$. Among hard primes $p \leq x$ with least non-residue $q(p) = q$, the proportion failing the selected residual R_0 is $\ll_q (\log x)^{-\eta_q} \rightarrow 0$, where $\eta_q = \min 1/\varphi(R_0) > 1/(4q)$, the minimum over the finitely many selected residuals $R_0 < 4q$ that arise for q .*

Proof. The condition $q(p) = q$ is a union of residue classes modulo $M = 840 \prod_{11 \leq \ell \leq q} \ell$; fix one, which fixes R_0 . By Proposition 2.2, any prime factor u of $(p + R_0)/4$ in the universal class $-4^{-1} \bmod R_0$ certifies success via $k = up^2$, so a failing p has $(p + R_0)/4$ free of an entire congruence class of primes: a sifting condition of dimension $1 + 1/\varphi(R_0)$ (with primality), and the Fundamental Lemma bounds the failing count by $\ll x/(\log x)^{1+1/\varphi(R_0)}$ against class size $\asymp x/\log x$ (Siegel–Walfisz). Summing over the classes, the proportion is governed by the worst class, giving η_q ; and $\varphi(R_0) < R_0 < 4q$ gives $\eta_q > 1/(4q)$. The bound is uniform for $q \ll \log \log x$. $\square$

Theorem 5.5 (proxy ladder, two-sided). *Fix q , a class (fixing $R_0, R_1, \dots$), and J with $\kappa = \sum_{j \leq J} 1/\varphi(R_j)$ below an absolute constant. Let $\tilde{E}_J(x)$ be the set of hard primes in the class whose shifted values $(p + R_j)/4$, $0 \leq j \leq J$, all avoid prime factors in their universal classes (so $E_j \subseteq \tilde{E}_j$ for every $j \leq J$). Then*

$$x/(\log x)^{1+\kappa} \ll_{q,J} \#\tilde{E}_J(x) \ll_{q,J} x/(\log x)^{1+\kappa} :$$

the avoidance ladder decays by the factor $(\log x)^{-1/\varphi(R_{j+1})}$ at every rung.

Proof sketch. The upper bound is the Fundamental Lemma as in Theorem 5.4. The lower bound is a beta-sieve in dimension $\kappa < 1$ on the prime-supported sequence, with level $x^{1/2-\varepsilon}$ from Bombieri–Vinogradov; integers with a forbidden prime factor above the sifting range are removed by Brun–Titchmarsh upper bounds in the progressions they define, whose total is absorbable for κ small. $\square$

Hypothesis 5.6 (P; proportionality). *For each fixed q , class, and J there are $c_P > 0$ and x_0 , with c_P bounded below uniformly in the rung $j \leq J$ (uniformity in q is part of the hypothesis), such that $\#E_j(x) \geq c_P \#\tilde{E}_j(x)$ for all $j \leq J$ and $x \geq x_0$, where E_j is the set of true ladder failures through rung j : failures are a positive proportion of their necessary-condition set.*

Measured: $c_P \approx 0.097$ – 0.12 at rung 0 and 0.13 – 0.16 at rung 1, flat across two decades to 10^{11} (Appendix A). Under Hypothesis 5.6, since $E_{j+1} \subseteq \tilde{E}_{j+1}$,

$$\#E_{j+1} \leq \#\tilde{E}_{j+1} \ll (\log x)^{-1/\varphi(R_{j+1})} \#\tilde{E}_j \leq (\log x)^{-1/\varphi(R_{j+1})} c_P^{-1} \#E_j,$$

so the ladder failure count decays at every *fixed* rung, for $x \geq x_0(q, j, \delta)$ — the per-rung decay Hypothesis 5.10 below posits, obtained from Hypothesis 5.6 rung by rung, though not with the uniformity in j and q that Hypothesis 5.10 demands.

The key to counting true failures is that the Type-I failure condition is a *norm-form* condition in disguise. For $r \equiv 3 \pmod{4}$, quadratic reciprocity gives $(\frac{-r}{\ell}) = (\frac{\ell}{r})$ for odd primes $\ell \neq r$, so

$$\begin{aligned} \text{every prime factor of } a \text{ is a QR mod } r &\iff -r \text{ is a square mod } 4a \\ &\iff a \text{ is } \textit{primitively} \text{ represented by} \\ &\quad \text{a binary quadratic form of discriminant } -r, \end{aligned} \tag{5.1}$$

the last step by the classical correspondence between representations and square roots of the discriminant (Cox [8, Lem. 2.5]; the discriminant $-r$ is prime, hence has a single genus, so the whole class group participates). Failure families are thus families of shifted values represented by binary quadratic forms, and the relevant counting theorem is Iwaniec's half-dimensional sieve in its *prime-indexed* form: the number of primes $p \leq N$ with $p = B\varphi(\xi, \eta) + A$ for a primitive form φ is $\asymp N/(\log N)^{3/2}$ [11]. This is not obstructed by parity — at dimension $\kappa = \frac{1}{2}$ the lower-bound sieve function satisfies $\beta(\frac{1}{2}) = 1$ — and it is the tool the following statements use.

Theorem 5.7 (the first link is sharp). *For $R \in \{3, 7, 15\}$,*

$$\#\{\text{hard primes } p \leq x \text{ failing } R\} \asymp \frac{x}{(\log x)^{3/2}},$$

and the same holds with $\ll, \gg$ separately at $R = 11$. In particular the chain of Theorem 1.11 is sharp at its first link, and no exponent better than $3/2$ is available from any single residual: by Theorem 1.10 every failure-sufficient family has sieve dimension $\geq \frac{1}{2}$, so $3/2$ is the ceiling of the method, not an artifact of the argument.

Proof sketch. At these residuals the exact criteria (Theorems 1.2, 1.4, 1.6) make failure equivalent to an all-residue condition on the factors of $a = (p + R)/4$, which by (5.1) is primitive representability by a form of discriminant $-r$ ($r = 3, 7, 15$ respectively — at $R = 15$ the Jacobi condition of Theorem 1.6 plays the role of the Legendre condition). The upper bound is the Fundamental Lemma as in Theorem 1.11; the lower bound is [11] applied to each class of forms, with the congruence conditions (H1)–(H3) absorbed into A and B . At $R = 11$ the case-(b) budget failures of Theorem 1.5 are not of this shape, which is why only the one-sided bounds are asserted. $\square$

Along the ladder the same mechanism applies, but only on classes whose local conditions are consistent — a much more delicate requirement than it first appears.

Theorem 5.8 (ladder failure lower bound). *Call a class $c \pmod{M}$ (M as in the proof of Theorem 5.4) admissible if R_0 has a prime factor $r_1 \equiv 3 \pmod{4}$ with (A1) $(\frac{a}{r_1}) = +1$; (A2) $(\frac{\ell}{r_1}) = +1$ for every prime $\ell \leq q$ that the class forces to divide a ; (A3) either $r_1 \equiv 3 \pmod{8}$ with a odd, or $r_1 \equiv 7 \pmod{8}$ with a even. On every admissible class,*

$$\begin{aligned} \#E_0(x) &\geq \#\{p \leq x : p \equiv c \pmod{M}, \text{ every prime factor of } (p + R_0)/4 \text{ a QR mod } r_1\} \\ &\gg \frac{x}{(\log x)^{3/2}}. \end{aligned}$$

Proof sketch. Membership forces failure by Proposition 2.1 at r_1 : all factors of a are residues mod r_1 by hypothesis, and $(\frac{p}{r_1}) = +1$ follows for free from consistency $a \equiv 4^{-1}p$, so every divisor of m^2 is a residue while $(\frac{-m}{r_1}) = (\frac{-1}{r_1}) = -1$. For the count, apply (5.1) and [11] with $A = -R_0$,

$B = 4 \prod_{\ell|a, \ell \leq q} \ell$ and discriminant $-r_1$; condition (A3) is what makes the 2-adic constraint solvable, and (A2) is what makes the family nonempty. One point is not a routine citation: the sieve delivers representability by *some* form of the discriminant — non-split primes occurring to even multiplicity — whereas (5.1) needs *primitive* representability, no non-split prime at all. That upgrade is the subject of Fuchs–Hsu–Rickards–Schindler–Stange [9], where it is shown to be neither immediate nor universally true; the cases used here fall under their Theorem 1.1. $\square$

Remark 5.9 (how much of the ladder this reaches). Condition (A2) is restrictive in a way that is easy to miss: the class fixes a modulo every prime $\leq q$, so a single forced factor ℓ with $(\frac{\ell}{r_1}) = -1$ empties the family. For instance $q = 23$, $R_0 = 35$, $r_1 = 7$ satisfies (A1), but $3 \mid a$ for every hard prime (as $p \equiv 1 \pmod{3}$ and $35 \equiv 2 \pmod{3}$) while $(\frac{3}{7}) = -1$: the set is empty. Counting properly, among hard primes below 2×10^7 the classes passing (A1) account for 6.6% of first rungs, but those passing (A1)–(A3) with a nonempty family account for only 0.023%. Theorem 5.8 is therefore an existence statement about a sparse family of ladders, not a bound covering a positive proportion of them; and whenever R_0 is prime — which is the case for the large majority of hard primes — no r_1 of this kind exists at all. Consequently the bracket on Hypothesis 5.6 is two-sided only on that sparse family at rung 0, and one-sided (upper bound only) elsewhere. Theorem 5.7 is the unconditional lower-bound statement of real reach; what Theorem 5.8 adds is that the same mechanism survives the passage to p -adapted residuals where it survives at all.

5.4 The conditional theorem, and honest scope

Hypothesis 5.10 (B; per-rung decay). *Fix a rung-count function $J(x) \rightarrow \infty$ (as calibrated after Theorem 5.11). There are $\delta > 0$ and x_0 such that for $x \geq x_0$, all $q \leq (\log x)^2$ and all $0 \leq j < J(x)$: $\#E_{j+1}(x) \leq (1 - \delta) \#E_j(x) + O(\sqrt{x})$.*

The $j = 0$ case is Theorem 5.4; every fixed rung follows from Hypothesis 5.6 (non-uniformly in j and q); the measured survival factors $1 - \delta$ are 0.05–0.14 per rung — 86–95% of the surviving failures resolve at each rung — with no drift in p .

Theorem 5.11 (conditional almost-all bound). *Assume Hypothesis 5.10. Then all but*

$$O\left(x e^{-\delta J(x)} + x 2^{-\pi((\log x)^2)}\right)$$

hard primes $p \leq x$ satisfy $R_{\min}(p) \leq 4q(p)(J(x) + 1) \leq (\log x)^{2+o(1)}J(x)$.

Proof. Hard primes with $q(p) > (\log x)^2$ are residues at every prime in $(7, (\log x)^2]$, a half-class condition per modulus, and the large sieve bounds their count by the second term. For the rest, iterate Hypothesis 5.10 $J(x)$ times; a survivor is the only way to lack a certificate on the ladder. $\square$

With $J(x) = (\log \log x)^2$ — a mild extrapolation of the nine measured rungs — the exceptional set is $\ll x \exp(-\delta(\log \log x)^2)$: this calibration is now delivered *unconditionally*, and with the stronger bound $R_{\min} \leq \varepsilon \log \log x$, by Theorem 4.5, so Hypothesis 5.10’s real content begins past the ceiling of Remark 4.6: with $J(x) = (\log x)^\theta$, $\theta > 2/3$, Theorem 5.11 would pass Vaughan’s bound. The hypothesis-strength required scales with the ambition, and each increment is testable before it is assumed.

Remark 5.12 (scope). Theorem 5.11 is an almost-all statement under a measured hypothesis; three things it is not. It is not unconditional: Hypothesis 5.6 is bounded below unconditionally only on the sparse family of Remark 5.9, and one-sidedly elsewhere — in particular whenever R_0 is prime, which is the typical case. It is not a density record: Vaughan’s bound remains asymptotically stronger unconditionally. And it is not a path to the full conjecture by itself: emptiness is beyond upper-bound methods (Remark 4.3), and excluding the individual conspiratorial prime — a super-record dodging $J(p)$ rungs, of which the record prime of Section 6 is an empirical foreshadowing — is a simultaneous-multiplicative-structure statement of prime-tuple difficulty. What the ladder does is localize the conjecture for hard primes into one measured counting inequality, bracketed on both sides where present methods reach and above everywhere, with every sub-statement current tools can decide now decided.

Remark 5.13 (why the walls are real: two obstruction computations). Two computations pin the walls down. (i) *Budget failure is not a character event*. In the model at $R = 11$ (generator 2), the configurations with class set $\{2, 7\}$ and multiplicities $(1, 1)$ versus $(4, 2)$ have the *same* class set and the same product $a \equiv 3 \pmod{11}$ — hence identical values under every Dirichlet character mod 11 — yet the first fails the target 8 ($\equiv -m$ for the consistent $p \equiv 1 \pmod{11}$) and the second attains it. No character-based machinery (pretentious dichotomies included) can separate the two, and 97.5% of measured first-rung ladder failures are of this budget type. (ii) *Weighted detection is equivalent to Hypothesis 5.6*. The certificate detector is already maximally relaxed — by Proposition 2.2, any *divisor* of the shifted value in the universal class certifies, so almost-prime weakenings of the detected object change nothing — and the first moment of any Bonferroni/Selberg weighting over E_j is itself a failure count: carrying the scheme through is equivalent, up to constants, to Hypothesis 5.6 itself. The parity gap of Theorem 5.8 is therefore intrinsic to every relaxation of this type, not an artifact of the chosen detector.

6 Computations to 10^{11}

All code and data are in the accompanying repository. Certificates are found by direct integer factorization of a (trial division; the factorization of $m^2 = p^2 a^2$ follows) and scanning divisors for the class condition (1.2); every certificate is re-verified in exact integer arithmetic, $4abc = p(bc + ac + ab)$.

6.1 Data

Every one of the 128,671,219 hard primes $p < 10^{11}$ possesses a residual certificate with $R \leq 107$. The verification behind this claim is graded by dataset: exhaustive reconstruction and exact re-checking of every entry below 10^9 ; systematic samples (143,000 resp. 257,343 entries) at 10^{10} and 10^{11} , together with the *complete* tails (842 entries with $R \geq 51$ resp. 20,151 with $R \geq 43$) re-checked with full minimality — Table 6 gives the per-dataset detail. The minimal residuals distribute as in Table 3 and Figure 1.

Three features of the distribution deserve comment.

- $R_{\max} = 107$, attained by the single prime $8,803,369 < 10^7$ and never exceeded through 10^{11} . In the language of Corollary 2.3: its 27 shifted values contain only 43 distinct odd primes (typical: 45–52), and 81% of them satisfy $\left(\frac{p}{q}\right) = +1$ (typical: 44–60%) — a $\approx 4\sigma$ fluctuation ($z = 4.1$ against the binomial null of independent unit signs) in the statistic that Theorem 1.8 identifies as governing joint failure.




Figure 1: The data of Table 3 on a logarithmic count scale. The head is dominated by $R = 3, 7, 11$; the tail thins geometrically; the shaded band 87–103 was empty below 10^{10} and its late population matched the calibrated prediction (see the confirmed-prediction bullet below); the record $R = 107$ is attained by a single prime below 10^7 .

- **A confirmed prediction.** The *calibrated independence model* treats the per-residual failure events as independent across R , with per-residual failure rates fitted from the complete 10^9 solvability masks and a single deep-tail correlation factor calibrated on the $R > 83$ counts; it is a heuristic, not a theorem, and its role is to make falsifiable predictions. Below 10^{10} no prime had minimal $R \in \{87, 91, 95, 99, 103\}$; the model attributed this gap to small-number statistics (expected counts ≤ 0.12 per value at 10^9) and predicted that deep-tail primes (failing every $R \leq 83$) land at $87, 91, 95, 99, 103, \geq 107$ with probabilities $0.53, 0.11, 0.22, 0.03, 0.05, 0.06$ (the last being the total mass at or beyond 107). Between 10^{10} and 10^{11} the five values acquired 8, 3, 5, 1, 1 primes respectively (all in $(1.3 \times 10^{10}, 10^{11})$), and none passed 103: the observed conditional distribution over the 19 deep-tail primes (8, 3, 5, 1, 1, 1 including the record) matches the prediction.
- Among all 1,587,581 hard primes below 10^9 , exactly *four* admit a unique working residual ≤ 107 (8,803,369 with 107; 142,361,209 with 59; 287,567,281 with 83; 794,037,841 with 63) — the entire obstruction to a shorter covering list. A greedy cover of the full solvability table uses 18 residuals; a packing argument gives a lower bound of 12.

6.2 Calibration of the base case of Theorem 1.11

The proportion of hard primes failing both $R = 3$ and $R = 7$, multiplied by $\log x$, is constant to within 1% across three decades of the 10^9 data (Table 4).

That is, the joint exceptional density is $\approx 5.17/\log x$: the decay matches the order predicted by the base case of Theorem 1.11, and 5.17 estimates the (conjectural) Selberg–Delange constant.

The exceptional set is density zero but decays slowly: 26.5% of hard primes below 10^9 fail both residuals; all are settled by deeper residuals of the list. The fitted per-residual decay exponents are consistent with the sieve prediction $\kappa = \frac{1}{2}$ where exact criteria exist: $\hat{\kappa} = 0.539$ for $R = 7$, and $\hat{\kappa} = 0.387$ for $R = 3$, where the short fitting range and Selberg–Delange lower-order terms bias the exponent downward.

6.3 Failure taxonomy

Sampling the complete solvability table at 10^9 : failures at $R = 7$ are 100% Type I (as Theorem 1.4 requires); at $R = 11$, 89.7% Type I and 10.3% the budget pattern of Theorem 1.5(b); the Type-I share broadly decreases with $\varphi(R)$ (e.g. 51% at $R = 103$, though not monotonically), as hitting one specific class among $\varphi(R)$ grows harder while the character dichotomy stays binary. Taken together, the data of this section matches the framework of Sections 2–4: each distributional feature — the head, the geometric tail, the once-empty band, the static record — has a structural or heuristic explanation (the exact criteria for the head; the calibrated model, built on the sieve dimensions, for the tail and the band; Theorem 1.8 for the record).

7 Open problems

1. **Finite covering.** Prove that some finite admissible set S covers all hard primes (Theorem 1.13 gives density $(\log x)^{-A}$ for every A ; emptiness is beyond the sieve).
2. **Beyond fixed lists.** The support bound is proved unconditionally for every residual, prime and composite (Theorem 1.10), Theorem 4.5 carries list growth to length $\varepsilon \log \log x$, and the branch classification that governs everything further is now formulated in Section 4.2: the kernel case is proved (Theorem 4.8), the count form refuted (Theorem 4.7), and what remains open is exactly Conjecture 4.9 — the aperiodic window-container statement, whose difficulty is now localized to the control of non-critical deletion steps (the punctured-chain mechanism is empty by Proposition 4.12); its resolution — for Hajós orders alone — gives Theorem 4.11. The fixed-list format itself cannot reach emptiness (Remark 4.3). The proof of Theorem 1.10 shows the extremal failing supports are *subgroup-trapped* (the stabilizer H is nontrivial and S concentrates on H up to at most $d/|H| - 2$ outliers); exploiting this structure *uniformly in R* — fewer, p -independent branches — is now a well-posed additive-combinatorics question, and the structurally right target remains a slowly growing bound on $R_{\min}(p)$.
3. **The constant 5.17.** The joint $\{3, 7\}$ -exceptional density is measured as $\approx 5.17/\log x$, stable to 1% across three decades. Conjecturally 5.17 is a Selberg–Delange-type constant: the count should satisfy $E(x) \sim Cx/(\log x)^2$ with $C = c_0 \prod_q \delta_q$, the local factors δ_q recording, for each prime q , the probability that q divides neither form in a forbidden class. Proving the asymptotic (not just the upper bound) requires a two-form Selberg–Delange or shifted-convolution analysis — open, and a natural test problem for such methods (see also Remark 1.14 for the density context).
4. **Structural reformulations and complexity.** Theorem 1.8 invites a graph formulation: on the bipartite graph joining primes q to residuals R by $q \mid (p + R)/4$, the edge characters $(\frac{q}{R})$ collapse to vertex signs $(\frac{p}{q})$, and joint Type-I failure is an all-plus vertex assignment — a percolation-style object whose structure could organize the growth law. Separately, the

meta-theorem’s enumeration complexity is unresolved: the DP state count grows empirically by $\approx 2.8\times$ per successive prime residual (271 at $R = 19$ to 1.95×10^7 at $R = 107$); is the equivalence-state complexity genuinely exponential in R ? (For the support-bound property itself the question is settled: Theorem 1.10 is a constant-size certificate. The open version concerns the full exact criteria.)

5. **The record and the covering hypothesis.** Theorem 1.8 explains the record prime quantitatively (see Corollary 2.3) and yields a conditional resolution of the covering question: under Dickson’s prime-tuple conjecture, for every B there are infinitely many hard primes failing every residual $\leq B$ — fix the forced smooth parts of the shifted values (character-neutral by Corollary 2.3(i)), make the cofactors prime in prescribed residue classes, and apply Proposition 2.1. So no fixed finite list suffices, conditionally, and the record grows without bound (at the pure-Type-I level, like $\log x / \log \log x$). Unconditional versions appear out of reach: already “infinitely many hard p fail $R = 3$ ” contains a two-form prime-tuple statement. What remains open is an unconditional upper bound on $R_{\min}(p)$ — and this is not an incremental item: by Proposition 1.1, $R_{\min}(p) \leq 2p$ is *equivalent* to the conjecture at p , so an unconditional bound $R_{\min}(p) \leq f(p)$ for any f whatsoever is the full problem.

A Verification details

This appendix documents every computer-assisted claim of the paper (their proof status is stated in Remark 1.7): the finite enumerations that are *proof components* (Appendix A.1), the large-scale certificate computations behind Section 6 and how each dataset is verified (Appendix A.2), and the commands that reproduce all of it (Appendix A.3). All code is in the accompanying repository; every check is deterministic and uses exact integer arithmetic only.

A.1 Proof components

Each component below is an exhaustive enumeration of a finite space whose sufficiency is established in the text (Theorem 1.3 for the exact criteria; for Lemma 1.9, monotonicity of the reachability of the target in support and multiplicities, checked at minimal exponents). Each run completes in seconds on commodity hardware, except the largest support-bound check ($R = 107$: 331 s); the entry points named are in the module `erdos_straus.theory`.

Trusted computing base. The certificate data does not rest on the solver: every certificate claim is re-checked by evaluating the integer identity $4abc = p(bc+ca+ab)$ directly, so the trusted base for the data is a bignum multiply-and-compare (the host language’s built-in integers), not the search pipeline. The finite case analyses trust more — the enumeration code itself, roughly sixty lines per check. Four mitigations apply. First, the enumerated spaces are defined by statements proved in the text (Theorem 1.3, Lemma 4.1), and the code mirrors those definitions one-to-one. Second, independent implementations agree where they overlap: the subset dynamic program and the direct binomial enumeration verify identical claims at $R = 19, 23$, and the exact criteria are cross-checked against ground-truth divisor search on up to 1,587,581 primes with zero discrepancies. Third, every check is deterministic, uses exact integer arithmetic only (no floating point, randomness, or external libraries), and runs from the repository in seconds (the

largest, the $R = 107$ dynamic program, in 331s) — an independent re-implementation is an afternoon’s work, and we would welcome one.

Fourth, much of the development is formalized in the accompanying repository (`lean/`) in Lean 4 [3] with its mathematical library `mathlib` [13]. For readers unfamiliar with proof assistants: Lean is a language in which theorem statements and proofs are written formally and every proof is checked by a small trusted *kernel*, so a theorem accepted by the kernel is correct relative to its formal statement, independently of who — or what — wrote the proof. Two Lean-specific mechanisms matter below: `native_decide` discharges a decidable proposition by running compiled code, which extends the trusted base from the kernel to Lean’s compiler and evaluator; and `#print axioms` reports, for any theorem, exactly which axioms (including any such evaluator appeal) its proof ultimately uses. The formalization comprises:

- *The elementary layer*, as ordinary kernel-checked proofs: certificate soundness and integrality, both directions of Theorem 1.2, both families of Proposition 1.12, Theorem 1.8, and the ladder’s elementary layer — Lemmas 5.1 and 5.2 in full, and the selected-residual identity of Section 5 (the formalization shows the coprimality hypothesis of Lemma 5.1 is not needed: for $q \mid R$ both symbols vanish).
- *The enumeration layer*: the reachability model underlying Theorem 1.3 is defined in the multiplicative structure of $\mathbb{Z}/R$ directly, and both halves of Lemma 4.1 are proved as theorems, with no computational content to trust. The finite checks behind Theorem 1.4 (1,536 configurations, checked as the 384 that remain after the neutral class collapses), Theorem 1.5 (497,664 capped configurations, including the exact two-case failure classification), and Lemma 1.9 at $R = 19, 23$ (437,580 resp. 7,759,752 checks), and Theorem 1.6 (349,920 configurations, via a product-index bridge for the non-cyclic unit group) are discharged by `native_decide`.
- *The coordinate bridges*: the larger enumerations are *evaluated* in discrete-log coordinates for speed, but the translation to the multiplicative model is itself a proved theorem, so the coordinates add nothing to the trusted base; the $R = 11, 19, 23$ statements hold in the multiplicative model, each reporting exactly its enumeration’s single evaluator axiom.
- *The reduction of Theorem 1.3*: membership of the target class in the reachability model is proved equivalent to the existence of a divisor $k \mid m^2$ in that class (via the fundamental theorem of arithmetic), and produces the explicit positive certificate (k, b, c) solving (1.1) — closing the loop between the abstract model the enumerations check and the integer certificates the conjecture asks for.
- *Two capstones*: a composed corollary (if $(p + 19)/4$ has nine prime factors in pairwise-distinct unit classes other than 1 modulo 19, an explicit solution of (1.1) at p exists, machine-checked end to end), and Lemma 1.9 at $R = 31$ by *certified dynamic programming* — the $\binom{29}{15} = 77,558,760$ supports are never enumerated; a machine-checked induction shows the dynamic program covers every support, and the compiled evaluator runs the dynamic program itself together with the final check that every state realized by ≥ 15 classes hits all 30 targets (3,001 states in the run) — the program’s *correctness* is proved; its evaluation is what `native_decide` trusts.

Every main theorem is audited with `#print axioms`, and helper lemmas are covered transitively through them. The elementary layer is now formalized outright, including Proposition 1.1 in

both directions (a solution is a certificate at its own residual, and conversely — so the equivalence organizing this paper is machine-checked) and Propositions 2.1 and 2.2. Theorem 1.6’s finite check is formalized by a product-index mask bridge over the non-cyclic $(\mathbb{Z}/15)^* \cong C_2 \times C_4$, and Theorem 1.10 fully symbolically: Kneser’s addition theorem vendored from its existing sorry-free formalization, an iterated many-summand form proved on top, and the support bound derived for an *arbitrary* finite abelian group with an explicit involution count — part (iii)’s general form, from which parts (i) and the odd-order variant follow as instances. So too the branch-classification layer: Theorem 4.8’s single kernel step in all three parts, with the quotient *constructed* rather than assumed and the container size bound $\leq d/2 - 1$ derived from the support bounds downstairs; Proposition 4.12 in full; and Theorem 4.7 end to end — the complete-partition machinery, the maximality of the constructed supports, their injectivity, and an explicit $2^{c\sqrt{d}}$ lower bound. (Not formalized: the transfer of Theorem 1.10(ii) to the multiplicative divisor-class model, the identity $t = 2^{\omega(R)} - 1$ for the involution count of $(\mathbb{Z}/R)^*$ — so the general bound is quoted with that t supplied by hand — the *iterated* kernel reduction, and the analytic arguments of Sections 4 and 5, which rest on the hand proofs alone.) For these enumerations the sixty-line trust base above is thus replaced by the Lean toolchain evaluating a formally stated proposition — and the Python and Lean implementations, written independently against the same statements, agree. (The development notes and Lean sources use working letter names for several results; the repository’s documentation includes the dictionary.)

In addition (as sanity checks, not proof components), the criteria of Theorems 1.2, 1.4, 1.5 and 1.6 were compared against ground-truth divisor search on the complete solvability data below 10^9 : agreement in 1,587,581/1,587,581 cases for $R = 3$, 41,421/41,421 sampled failure cases for $R = 7$ (all Type I, as Theorem 1.4 requires), 158,759/158,759 sampled cases for $R = 15$, and 79,380 sampled cases for $R = 11$, with zero discrepancies; the general composite- R engine (`solvable_exact_general`) likewise agrees on 7,938 sampled primes for every composite admissible residual, and the strong form of Theorem 1.10(ii) is confirmed numerically at every residual checked (`verify_support_bound_strong`). One check named in this paragraph’s vicinity is *not* a sanity check: `kneser_support_general` is a proof component of the $R \leq 107$ clause of Theorem 1.10(iii), and is listed as such in Table 5. The identity of Theorem 1.8 and its forced-prime corollary were likewise checked on 135,588 resp. 61,438 instances from the same data, with zero violations (Corollary 2.3).

A.2 Certificate datasets and their verification

The computational claims of Section 6 rest on the following datasets, all archived in the repository under `data/`. A certificate (R, a, b, c) for p is *accepted* only if $4abc = p(bc + ac + ab)$ holds in exact integer arithmetic and $R = 4a - p$; for the compact minimal- R maps, the verifier *reconstructs* the triple from (p, R) by the divisor search of (1.2) and then applies the same exact check, so a map entry is verified as a genuine certificate, not taken on faith.

The masks dataset of Table 6 is the basis for the covering-set statements (the 18-residual cover, the lower bound 12, the four critical primes), the failure taxonomy, and the calibration of Section 6; the minimal- R maps are the basis for Table 3, the record, and the gap analysis.

A.3 Reproduction

With the repository installed (`pip install -e .`), the commands of Table 7 regenerate and re-verify everything above. Generation of the full 10^{10} map takes about 11 minutes on four

cores; every other command runs in seconds to minutes.

The ladder census of Section 5 is reproduced by `python -m erdos_straus.burgess_scan` (the 10^9 population) and by the scaled drivers in the same module for the window samples and the deep tail; the archives are the five `burgess_*.json` files under `data/analysis/`.

The full test suite (`python -m pytest`; 56 tests) exercises the proof-component checks, the exact criteria against ground truth, and sampled dataset validation on every run. Table 8 summarizes the computational environment.

Data and code availability

All code, all datasets of Table 6, and the Lean development are openly available in the repository <https://github.com/jeffmhopkins/Erd-s-Straus-attack>.

result	proof inputs	proved in	machine-checked?
Prop. 1.1	elementary	§1.1	Lean, both directions
Props. 2.1, 2.2	elementary (characters; explicit k)	§2	Lean
Thm. 1.2 ($R=3$)	elementary	§3	Lean, both directions
Thm. 1.3	elementary	§3	model + reduction in Lean
Thm. 1.4 ($R=7$)	Prop. 2.1 + Thm. 1.3 + enumeration (1,536)	§3	Lean (<code>native_decide</code>)
Thm. 1.5 ($R=11$)	Thm. 1.3 + enumeration (497,664)	§3	Lean (<code>native_decide</code>)
Thm. 1.6 ($R=15$)	Thm. 1.3, composite form + enumeration (349,920)	§3	Lean (<code>native_decide</code>)
Thm. 1.8	quadratic reciprocity	§2	Lean
Lem. 1.9	Lem. 4.1 + enumerations / DP	§4	Lean at $R = 19, 23, 31$
Thm. 1.10	Kneser’s theorem + Lem. 4.1; (iii) upgrade at $R \leq 107$: enumeration	§4	Lean, parts (i) and (iii); (iii) upgrade: Python
Thm. 1.11	Thms. 1.2, 1.4–1.6, 1.10 + Fundamental Lemma	§4	no
Prop. 1.12	elementary + Thm. 1.11	§4	families in Lean
Thm. 1.13	Prop. 2.2 + Fundamental Lemma	§4	no
Thm. 4.5	Thm. 1.10(iii), general tier + Brun’s sieve, uniform constants	§4	no
Thm. 4.7	explicit construction (complete partitions)	§4	Lean
Thm. 4.8	Kneser + Lem. 4.1 + quotient induction	§4	Lean, single step
Thm. 4.10	Thm. 4.8 + Conjecture 4.9	§4	no (conditional)
Prop. 4.12	elementary (isolation of punctures)	§4	Lean
Thm. 4.11	Thm. 4.10 + the audit of Thm. 4.5	§4	no (conditional)
Lems. 5.1–5.2	Jacobi reciprocity, elementary	§5	Lean
Thms. 5.3–5.5	Thm. 1.10 / Prop. 2.2 + sieve	§5	no
Thm. 5.7	norm-form bridge + [11] + Fundamental Lemma	§5	no
Thm. 5.8	Prop. 2.1 + [11] + [9]	§5	no
Thm. 5.11	Hypothesis 5.10 (measured) + large sieve	§5	no (conditional)

Table 1: Audit map: what each main result depends on, where it is proved, and its formal-verification status. “—” means the result is elementary enough that no machine check was warranted; “no” means the proof is by hand only. Theorems 5.5 and 5.8 are established in outline (proof sketches).

population	first-rung success	ladder ($R \leq 400$)
all hard $p < 10^9$ (exhaustive)	94.78%	100% (two q 's)
10^9 – 3.2×10^9 (sampled)	95.4%	100%
3.2×10^9 – 10^{10}	95.0%	100%
10^{10} – 3.2×10^{10}	95.8%	100%
3.2×10^{10} – 10^{11}	95.7%	100%
10^{11} deep tail ($R_{\min} \geq 43$, complete)	75.1%	99.99% (100% at $R \leq 435$)

Table 2: The selected-residual census: success rates at the first rung and along the q -ladder, capped at rungs $R_j \leq 400$. “(two q 's)”: below 10^9 , exactly two primes needed the ladder of the *second*-least non-residue to resolve within the cap. The deep-tail row conditions on $R_{\min} \geq 43$, the adversarial 0.016% of the 10^{11} population; there, raising the cap to $R \leq 435$ restores 100%. The per-rung failure rate is flat in p across two decades, including on the deep-tail population.

R	count	R	count	R	count
3	69,951,190	39	25,258	75	29
7	31,335,473	43	6,850	79	50
11	19,439,771	47	8,811	83	23
15	3,974,286	51	1,450	87	8
19	2,441,687	55	1,369	91	3
23	1,070,808	59	954	95	5
27	187,963	63	283	99	1
31	193,848	67	104	103	1
35	30,784	71	209	107	1

Table 3: Distribution of the minimal residual $R_{\min}$ over the 128,671,219 hard primes below 10^{11} (plotted in Figure 1).

bin ($p \sim$)	10^6	10^7	6×10^7	2.5×10^8	10^9
density $\times \log x$	5.14	5.19	5.18	5.17	5.16

Table 4: The proportion of hard primes near x that fail both $R = 3$ and $R = 7$, multiplied by $\log x$, in five bins spanning three decades of the 10^9 data (x is the bin's representative value). Constancy of the product is the decay rate predicted by the base case of Theorem 1.11.

component	configuration space	count	entry point
Thm. 1.4 ($R = 7$)	class multisets in $(\mathbb{Z}/7)^*$, multiplicities ≤ 3 , with the forced even factor, $p \bmod 7 \in \{1, 2, 4\}$, and the consistency relation	1,536	<code>verify_R7_finite</code>
Thm. 1.5 ($R = 11$)	dynamic program over $(\mathbb{Z}/11)^*$ with the forced factor 3; collapses to 25 equivalence states (16 succeed, 6 fail by character, 3 by budget)	25	<code>finite_criterion_dp</code>
Thm. 1.6 ($R = 15$)	class multisets in $(\mathbb{Z}/15)^*$, multiplicities ≤ 5 , with the forced even factor, $p \bmod 15 \in \{1, 4\}$, and the consistency relation (349,380 succeed, 540 fail all-residue, zero budget failures)	349,920	<code>verify_R15_finite</code>
Lem. 1.9 ($R = 19$)	all supports of size 9 over the nonzero logs $\times$ all classes of p , at minimal multiplicities	437,580	<code>verify_support_bound</code>
Lem. 1.9 ($R = 23$)	all supports of size 11, likewise	7,759,752	<code>verify_support_bound</code>
Lem. 1.9 ($R = 19$ –107, all twelve primes)	subset dynamic program over reachability-mask states; state counts 271 ($R=19$), 3,001 ($R=31$), 2,414,145 ($R=83$), 19,530,971 ($R=107$); zero violations throughout	$\approx 19.5\text{M}$	<code>verify_support_bound_dp</code>
Thm. 1.10(iii), $R \leq 107$ clause	maximal support of a proper product-set in $(\mathbb{Z}/R)^*$ equals $\varphi(R)/2 - 1$ for every composite admissible $R \leq 107$ (e.g. $R = 95$: max 35 over 583,455 realizable product-sets)	12 moduli	<code>kneser_support_general</code>

Table 5: Machine-verified proof components: enumerated configuration spaces and the entry points in `erdos_straus.theory`.

dataset	entries	verification performed
explicit triples, $p < 1.2 \times 10^8$	213,131	every triple checked exactly
minimal- R map, $p < 10^9$	1,587,581	<i>exhaustive</i> : every entry reconstructed and checked exactly
minimal- R map, $p < 10^{10}$	14,215,707	143,000 systematically sampled entries reconstructed and checked; additionally <i>all</i> 842 tail entries ($R \geq 51$) checked with full minimality (no smaller admissible R succeeds)
solvability masks, $p < 10^9$	$27 \times 1,587,581$	per-prime solvability of every admissible $R \leq 107$ by full divisor-class computation; consistency with the minimal- R map
R -sequence dataset, $p < 10^{11}$	128,671,219	257,343 systematically sampled entries reconstructed and checked; <i>all</i> 20,151 tail entries ($R \geq 43$) checked with full minimality; prime array pinned by sha256; pipeline validated byte-for-byte against the exhaustively verified 10^9 map

Table 6: Certificate datasets and the verification applied to each.

claim	command
Thm. 1.4 finite check	<code>python -c "from erdos_straus.theory \ import verify_R7_finite; \ print(verify_R7_finite())"</code>
Thm. 1.5 enumeration	<code>finite_criterion_dp(11, forced_logs=[_dlog_table(11)[1][3]]) from erdos_straus.theory (as in the test suite, tests/test_solver.py)</code>
Lem. 1.9 at $R = 19, 23$	<code>verify_support_bound(19), (23)</code>
Thm. 1.6 finite check	<code>python -c "from erdos_straus.theory \ import verify_R15_finite; \ print(verify_R15_finite())"</code>
Thm. 1.10 numerical confirmation	<code>verify_support_bound_strong(31), kneser_support_general(15)</code>
generate the 10^{10} map	<code>python -m erdos_straus.bulk_generate \ -max 10000000000 -segmented -store rmap</code>
verify a dataset exhaustively	<code>python -m erdos_straus.verify data/<file></code>
distribution / covering / taxonomy	<code>python -m erdos_straus.analyze dist \ -rmap data/hard_primes_1e9_minimalR.json.gz (also cover, tail); python -m erdos_straus.theory</code>
Lean formalization (elementary + enumeration layers)	<code>cd lean/ErdosStraus lake exe cache get && lake build</code>

Table 7: Reproduction commands. Shell commands are given verbatim, with `\` marking an explicit line continuation within a single command; the finite-check rows are Python calls from the module `erdos_straus.theory`.

implementation	Python 3, exact integer arithmetic only; no external libraries in any proof-component check (<code>sympy</code> appears only in the reference solver)
factorization	trial division by sieved primes up to $\sqrt{n}$ (numpy-sieved prime tables for bulk generation)
wall-clock times	10^{11} generation: 2.85 h on 4 cores; 10^{10} : 656 s; 10^9 : 49 s; largest proof-component check ($R = 107$ dynamic program): 331 s
verification	every certificate re-checked via $4abc = p(bc+ca+ab)$; exhaustive below 10^9 ; 257,343 samples plus all 20,151 tail entries at 10^{11} (Table 6)
integrity	datasets archived in the repository; the 10^{11} prime array and R -sequence are pinned by SHA-256 in <code>hard_primes_1e11_minimalR.meta.json</code> (prefixes 53cf17c4..., 2a69f3aa...)
formal layer	Lean 4.32.2 + mathlib (pinned toolchain); six <code>native_decide</code> checks; Theorem 1.10(i) fully symbolic (vendored Kneser); every main theorem audited by <code>#print axioms</code>

Table 8: Computational environment and integrity summary.

References

- [1] N. C. Ankeny, *The least quadratic non residue*, Ann. of Math. (2) **55** (1952), 65–72.
- [2] D. A. Burgess, *The distribution of quadratic residues and non-residues*, Mathematika **4** (1957), 106–112.
- [3] L. de Moura and S. Ullrich, *The Lean 4 theorem prover and programming language*, in: Automated Deduction – CADE 28, Lecture Notes in Computer Science, vol. 12699, Springer, Cham, 2021, pp. 625–635.
- [4] C. Elsholtz and T. Tao, *Counting the number of solutions to the Erdős–Straus equation on unit fractions*, J. Aust. Math. Soc. **94** (2013), 50–105.
- [5] P. Erdős, *Az $1/x_1 + 1/x_2 + \cdots + 1/x_n = a/b$ egyenlet egész számú megoldásairól* [On the integer solutions of the equation $1/x_1 + 1/x_2 + \cdots + 1/x_n = a/b$] (Hungarian), Mat. Lapok **1** (1950), 192–210.
- [6] J. Friedlander and H. Iwaniec, *Opera de Cribro*, American Mathematical Society Colloquium Publications, vol. 57, American Mathematical Society, Providence, RI, 2010.
- [7] H. Halberstam and H.-E. Richert, *Sieve Methods*, Academic Press, London, 1974.
- [8] D. A. Cox, *Primes of the Form $x^2 + ny^2$: Fermat, Class Field Theory, and Complex Multiplication*, 2nd ed., Wiley, Hoboken, NJ, 2013.
- [9] E. Fuchs, K. Hsu, J. Rickards, D. Schindler and K. Stange, *Primitive representations by binary quadratic forms*, preprint, arXiv:2504.20289 [math.NT], 2025.
- [10] H. Iwaniec, *The half dimensional sieve*, Acta Arith. **29** (1976), 69–95.
- [11] H. Iwaniec, *Primes of the type $\varphi(x, y) + A$ where φ is a quadratic form*, Acta Arith. **21** (1972), 203–234.
- [12] M. Kneser, *Abschätzung der asymptotischen Dichte von Summenmengen*, Math. Z. **58** (1953), 459–484.
- [13] The mathlib Community, *The Lean mathematical library*, in: Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2020), ACM, New York, 2020, pp. 367–381.
- [14] S. Mihnea and D. C. Bogdan, *Further verification and empirical evidence for the Erdős–Straus conjecture*, preprint, arXiv:2509.00128 [math.NT], 2025.
- [15] L. J. Mordell, *Diophantine Equations*, Academic Press, London, 1969.
- [16] S. E. Salez, *The Erdős–Straus conjecture: new modular equations and checking up to $N = 10^{17}$* , preprint, arXiv:1406.6307 [math.NT], 2014.
- [17] J. W. Sander, *On $4/n = 1/x + 1/y + 1/z$ and Rosser’s sieve*, Acta Arith. **59** (1991), 183–204.
- [18] A. Swett, *The Erdős–Straus conjecture*, online verification record to 10^{14} , 1999; formerly at <http://math.uindy.edu/swett/esc.htm> (archived snapshots available via the Internet Archive; accessed August 2026).

- [19] T. Tao and V. H. Vu, *Additive Combinatorics*, Cambridge Studies in Advanced Mathematics, vol. 105, Cambridge University Press, Cambridge, 2006.
- [20] R. C. Vaughan, *On a problem of Erdős, Straus and Schinzel*, *Mathematika* **17** (1970), 193–198.
- [21] W. A. Webb, *On $4/n = 1/x + 1/y + 1/z$* , *Proc. Amer. Math. Soc.* **25** (1970), 578–584.
- [22] K. Yamamoto, *On the Diophantine equation $4/n = 1/x + 1/y + 1/z$* , *Mem. Fac. Sci. Kyushu Univ. Ser. A* **19** (1965), 37–47.